
Polynomial Links as Parameter-Tied Rank Lifts for Tucker Tensors

Anonymous Author(s)

Affiliation

Address

email

Abstract

Tucker decomposition is a compact model for multiway data, but its map from a latent tensor to measurements is linear. Nonlinear measurement responses can therefore make a low-rank latent tensor appear higher-rank in the measured tensor. We propose Nonlinear Tucker Decomposition with Polynomial Links (NTD-PL), a Tucker enhancement that learns a polynomial link jointly with a Tucker backbone. NTD-PL keeps the backbone ranks fixed and adds only scalar link coefficients. Its polynomial powers induce a parameter-tied rank lift: each degree forms Veronese-lifted Tucker factors tied to the same latent tensor, rather than an untied high-rank correction.

We analyze this structure through Veronese rank expansion, separation from matched-rank Tucker, and approximation results for nonlinear responses. We optimize a masked reconstruction and completion objective by alternating Link Refresh, a small ridge update for the scalar link, with Linked-Backbone Descent and degree continuation. Experiments on controlled tensors, hyperspectral reconstruction and completion, rank-lift and Link Refresh diagnostics, and external cross-domain cases show scoped gains: NTD-PL helps when the residual left by matched-rank Tucker is coherent with a nonlinear response.

1 Introduction

Tensor decompositions provide compact models for multiway data in scientific data analysis, signal processing, and masked multiway learning [Kolda and Bader, 2009, Sidiropoulos et al., 2017, Liu et al., 2013, Liu and Moitra, 2020]. Tucker decomposition is a standard choice because it separates a small core from mode-wise factors. It uses a linear map from the Tucker latent tensor to the measured tensor: the reconstructed Tucker value is fitted directly to the measurement. This assumption is natural for many data sets, but it can be mismatched to measurement or data-generation pipelines with response curves, saturation, calibration effects, or material-dependent nonlinearities [Debevec and Malik, 1997, Grossberg and Nayar, 2004, Bioucas-Dias et al., 2012, Heylen et al., 2014, Luo et al., 2024]. In such settings, the latent signal may be well organized by a low-rank Tucker tensor while the measured tensor appears to have larger multilinear rank after a nonlinear response.

A direct remedy is to increase Tucker rank, but this changes the multilinear budget and expands the core and factors throughout the model. More flexible nonlinear tensor decoders based on kernels, Gaussian processes, or neural networks can also be useful, but they address a broader problem and often make the relation to a Tucker backbone less explicit [Xu et al., 2011, Zhe et al., 2016, Liu et al., 2017, Saragadam et al., 2024, Ahn et al., 2024]. We ask a narrower question: can a Tucker model keep its chosen backbone ranks fixed while learning a response that explains coherent nonlinear residuals left by matched-rank Tucker?

We instantiate this idea with a polynomial link. Let

$$\mathcal{S} = \llbracket \mathcal{X}; \mathbf{A}^{(1)}, \dots, \mathbf{A}^{(N)} \rrbracket$$

37 be a Tucker latent tensor. Instead of fitting the measured tensor directly by $\mathcal{S}$, we model

$$\hat{\mathcal{Y}} = f_{\beta}(\mathcal{S}) = \sum_{p=0}^P \beta_p \mathcal{S}^{\odot p},$$

38 where $\odot p$ denotes the Hadamard power. We call the model Nonlinear Tucker Decomposition
 39 with Polynomial Links (NTD-PL). It contains Tucker as the degree-one case, keeps the Tucker
 40 backbone and multilinear ranks fixed, and adds only $P + 1$ scalar link coefficients. The higher powers
 41 $\mathcal{S}^{\odot p}$ induce high-order interactions among the same core and factors, so the added expressivity is
 42 parameter-tied to one learned latent tensor.

43 This viewpoint is central to the paper. NTD-PL is best understood as a Tucker enhancement and
 44 shared-backbone rank lift, not as a generic nonlinear decoder. A degree- p power of the latent tensor
 45 behaves like an ordinary Tucker tensor whose factors are polynomial feature lifts of the original
 46 factors, but those lifted features remain tied to one learned backbone. The new degrees therefore
 47 reshape the low-rank geometry through the polynomial link, rather than adding an untied high-rank
 48 tensor block. This is also why the link is not a post-hoc calibration layer.

49 Our claims are intentionally scoped. Increasing Tucker rank remains the right comparison when the
 50 residual is ordinary multilinear structure, and we do not claim to dominate neural, kernel, or other
 51 nonlinear tensor factorization methods. The intended regime is more specific: the residual left by
 52 matched-rank Tucker should be coherent with a nonlinear response.

53 This scoped regime is also supported empirically. In controlled tensors, the gain grows with the
 54 injected nonlinear residual energy and with the degree needed to approximate the response. On
 55 real hyperspectral tensors, NTD-PL improves matched-rank Tucker in both full reconstruction and
 56 held-out completion. More importantly, a one-step Link Refresh diagnostic predicts when NTD-PL
 57 will help: across HSI scenes and non-HSI tensor units, larger diagnostic reduction is associated with
 58 larger NTD-PL gains, while low-yield units show little or no benefit. These results suggest that
 59 nonlinear responses are identifiable structure in several real tensor families.

60 **Contributions.** We make the following contributions.

- 61 • **Model.** We propose NTD-PL, a Tucker enhancement with a polynomial link on a Tucker latent ten-
 62 sor. It keeps the Tucker backbone and multilinear ranks fixed while adding only scalar coefficients
 63 for a parameter-tied response.
- 64 • **Theory.** We characterize the induced model class through Veronese rank expansion, separation
 65 from matched-rank Tucker, and approximation results for nonlinear responses.
- 66 • **Optimization.** We derive a masked objective for reconstruction and completion, and optimize it
 67 with Link Refresh, Linked-Backbone Descent, and degree continuation.
- 68 • **Experiments.** We evaluate controlled nonlinear residuals, hyperspectral reconstruction and com-
 69 pletion, rank-lift and Link Refresh diagnostics, and external cross-domain boundary cases.

70 2 Related Work

71 **Low-rank tensor models and Tucker backbones.** CP, Tucker, tensor train, and tensor ring
 72 decompositions provide compact representations for multiway data through different rank notions
 73 and factor topologies [Kolda and Bader, 2009, Sidiropoulos et al., 2017, De Lathauwer et al., 2000,
 74 Oseledets, 2011, Zhao et al., 2016]. Recent completion work has improved scalability, sampling
 75 guarantees, and visual-data priors while still mostly changing the tensor rank model, tensor-network
 76 topology, or low-rank regularizer [Liu and Moitra, 2020, Wang et al., 2021, Zheng et al., 2021].
 77 Tucker is the closest baseline for NTD-PL because it uses the same core-factor multilinear backbone.
 78 NTD-PL does not primarily compete by changing the rank tuple, replacing the factor topology, or
 79 moving to another tensor-network format. Instead, it keeps the Tucker backbone and multilinear
 80 ranks fixed and changes the measurement map on top of the Tucker latent tensor.

81 **Nonlinear tensor factorization and neural/kernel decoders.** Nonlinear tensor models include
 82 generalized likelihood and link formulations, Bayesian and kernelized tensor factorizations, Gaussian-
 83 process tensor models, and neural tensor decoders [Hong et al., 2020, Schein et al., 2016, Xu et al.,

2011, Zhe et al., 2016, Liu et al., 2017, Pan et al., 2020, Saragadam et al., 2024, Varolgunes et al., 2023, Ahn et al., 2024, Luo et al., 2024]. Related imaging methods also learn linear or nonlinear transforms under low-rank tensor priors [Luo et al., 2022, Wang et al., 2024, Xu et al., 2024]. These methods can be more flexible than a scalar polynomial link, but this flexibility often comes with heavier inference, additional hyperparameters, or a less transparent finite-dimensional relation to Tucker. NTD-PL is complementary: it deliberately uses a low-dimensional scalar response rather than a general kernel, Gaussian-process, neural decoder, or learned transform, so that the induced effect can be analyzed relative to a fixed Tucker backbone.

Polynomial feature lifts and high-order interactions. Polynomial and Volterra-style expansions are classical tools for representing nonlinear responses [Volterra, 1959, Boyd and Chua, 1985]. Related matrix and tensor methods use polynomial feature maps, quadratic factorizations, or explicit high-order interactions to enlarge a linear model [Zhai et al., 2024, Ruegamer, 2024, Liu et al., 2015, Deng et al., 2016]. Such constructions can make nonlinear interactions explicit, but untied high-order parameterizations can increase parameter dimension and contraction cost. NTD-PL does not introduce an independent high-order tensor parameter. Its polynomial powers are powers of one learned Tucker latent tensor, inducing parameter-tied high-order interactions and Veronese-style backbone-tied rank lifts.

Positioning. Existing nonlinear tensor models either use broad nonlinear decoders, modify likelihoods or kernels, or explicitly parameterize high-order interactions. What is less developed is a Tucker-preserving, low-dimensional, analyzable polynomial link that yields a parameter-tied rank lift and supports masked learning. NTD-PL trades general nonlinear flexibility for structure, interpretability, and a direct finite-dimensional rank-lift view. Its expected benefit is therefore scoped to settings where matched-rank Tucker leaves a coherent nonlinear residual explained by a scalar response.

3 Model

3.1 Model and Measurement Assumption

Let $\mathcal{Y} \in \mathbb{R}^{I_1 \times \dots \times I_N}$ be a measured tensor and let

$$\mathcal{S} = [\![\mathcal{X}; \mathbf{A}^{(1)}, \dots, \mathbf{A}^{(N)}]\!], \quad \mathcal{X} \in \mathbb{R}^{R_1 \times \dots \times R_N}, \quad \mathbf{A}^{(n)} \in \mathbb{R}^{I_n \times R_n}. \quad (1)$$

Our measurement assumption is that the measured tensor is generated from this low-rank Tucker latent tensor through a scalar response, up to noise. NTD-PL models this response by a polynomial:

$$Y_{\mathbf{i}} = f_{\beta}(S_{\mathbf{i}}) + \varepsilon_{\mathbf{i}}, \quad f_{\beta}(s) = \sum_{p=0}^P \beta_p s^p, \quad \mathbf{i} = (i_1, \dots, i_N). \quad (2)$$

Equivalently, the tensor prediction is

$$\hat{\mathcal{Y}} = f_{\beta}(\mathcal{S}) = \sum_{p=0}^P \beta_p \mathcal{S}^{\odot p}, \quad \beta \in \mathbb{R}^{P+1}. \quad (3)$$

The constant term β_0 captures affine response shifts, the linear term $\beta_1 \mathcal{S}$ is the usual Tucker channel, and terms with $p \geq 2$ activate nonlinear channels on the same latent tensor. Standard Tucker is recovered by setting $\beta_1 = 1$ and $\beta_p = 0$ for $p \neq 1$. Thus NTD-PL changes the measurement map while leaving the Tucker backbone and its multilinear ranks fixed.

3.2 Backbone-Tied High-Order Interactions

Although f_{β} is scalar, the term $\mathcal{S}^{\odot p}$ is not merely an independent output calibration. We arrive at it by first defining an explicit p -fold Tucker extension: instead of mapping one latent index per mode, the mode- n projection is a higher-order tensor $\mathcal{A}^{(n,p)} \in \mathbb{R}^{I_n \times R_n^p}$ that couples p latent indices. The case $p = 1$ is ordinary Tucker, while $p = 2$ gives a quadratic Tucker-style interaction between two latent Tucker tuples. Appendix A.1 gives the formal construction and tensor-network view.

The explicit construction is a reference model, not the parameterization we train. Its projection tensors grow as $I_n R_n^p$, and each degree would add a new set of mode projections. NTD-PL keeps the

125 same high-order interaction pattern by tying the p -fold projection to the original Tucker factor:

$$\mathcal{A}_{i_n, r_n^{(1)}, \dots, r_n^{(p)}}^{(n,p)} = \prod_{q=1}^p A_{i_n, r_n^{(q)}}^{(n)}.$$

126 Under this tying, the degree- p channel is exactly S_i^p . Thus it still couples p latent Tucker index tuples,
 127 but all copies share the same mode factors. Across degrees, the only new free parameters are the
 128 scalar coefficients $\{\beta_p\}$, which mix these tied channels.

129 The latent tensor $\mathcal{S}$ and the coefficients β are fitted jointly, so the low-rank backbone is optimized
 130 together with the polynomial channels. Section 4 analyzes the induced rank expansion and the
 131 resulting separation from standard Tucker.

132 3.3 Masked Least-Squares Objective

133 Let $\Omega \in \{0, 1\}^{I_1 \times \dots \times I_N}$ denote the observed-entry mask and let $|\Omega| = \sum_i \Omega_i$. We estimate the core,
 134 factors, and link coefficients by the regularized masked objective

$$\min_{\mathcal{X}, \{\mathbf{A}^{(n)}\}, \beta} \frac{1}{2|\Omega|} \|\Omega \odot (\mathcal{Y} - f_{\beta}(\mathcal{S}))\|_F^2 + \frac{\lambda_X}{2} \|\mathcal{X}\|_F^2 + \sum_{n=1}^N \frac{\lambda_n}{2} \|\mathbf{A}^{(n)}\|_F^2 + \frac{\lambda_{\beta}}{2} \|\beta\|_2^2. \quad (4)$$

135 The same formulation covers full reconstruction by taking $\Omega \equiv \mathbf{1}$ and tensor completion by restricting
 136 the loss to the observed entries. The quadratic penalties regularize the Tucker backbone and the
 137 low-dimensional polynomial link.

138 4 Theory

139 The theory addresses the rank-lift view raised by the model definition. A scalar link has few free
 140 parameters, but its powers can induce ordinary Tucker tensors with larger multilinear ranks. We
 141 first make this induced structure explicit, then give a separation from matched-rank Tucker and an
 142 approximation route for nonlinear responses.

143 4.1 Rank Expansion and Separation

144 **Proposition 1** (Shared-backbone rank expansion). *Let $\mathcal{S} = \llbracket \mathcal{X}; \mathbf{A}^{(1)}, \dots, \mathbf{A}^{(N)} \rrbracket$ have Tucker rank*
 145 *at most $(R_1, \dots, R_N)$. For every integer $p \geq 1$, $\mathcal{S}^{\odot p}$ admits a Tucker representation with mode ranks*
 146 *at most*

$$\left(\binom{R_1 + p - 1}{p}, \dots, \binom{R_N + p - 1}{p} \right). \quad (5)$$

147 *More explicitly, the ordinary Tucker factors of $\mathcal{S}^{\odot p}$ can be chosen as Veronese lifts of the original*
 148 *factors. For $\mathbf{A} \in \mathbb{R}^{I \times R}$, define $\mathbf{A}^{(p)} \in \mathbb{R}^{I \times D}$, $D = \binom{R+p-1}{p}$, by*

$$A_{i,\alpha}^{(p)} = \prod_{r=1}^R A_{i,r}^{\alpha_r}, \quad (6)$$

149 *where $\alpha \in \mathbb{N}^R$ ranges over $|\alpha| = p$. Then $\mathcal{S}^{\odot p}$ factors through $\mathbf{A}^{(1)(p)}, \dots, \mathbf{A}^{(N)(p)}$.*

150 This proposition makes the tying in Section 3.2 explicit in standard Tucker language: a degree- p
 151 channel is an ordinary Tucker tensor whose mode factors are polynomial feature lifts of the same
 152 backbone factors. The link therefore changes rank structure through tied features, not through an
 153 independent high-order parameter block.

154 **Theorem 1** (Tensorial separation from standard Tucker). *Let $N \geq 3$, $R \geq 2$, $p \geq 2$, and suppose*
 155 *$I_n \geq D$ for every n , where $D = \binom{R+p-1}{p}$. Then there exists an order- N tensor $\mathcal{Y} \in \mathbb{R}^{I_1 \times \dots \times I_N}$ that*
 156 *is representable by a degree- p NTD-PL model with Tucker backbone rank $(R, \dots, R)$, but whose*
 157 *exact standard Tucker rank is*

$$(D, \dots, D). \quad (7)$$

158 *Equivalently, a degree- p polynomial link can generate tensors whose multilinear ranks are*
 159 *$(\binom{R+p-1}{p}, \dots, \binom{R+p-1}{p})$ from the same Tucker backbone of rank $(R, \dots, R)$.*

160 The proof in Appendix B uses a CP-diagonal Tucker backbone and generic full-rank Veronese feature
 161 matrices. The theorem is an existence separation, not a universality claim over all high-rank Tucker
 162 tensors: it identifies a structured high-rank subset generated by a low-rank backbone.

163 4.2 Approximation Under Nonlinear Measurement Responses

164 **Theorem 2** (Polynomial generation). *Suppose $\mathcal{Y} = g(\mathcal{S}^*)$, where $\mathcal{S}^*$ has Tucker rank at most
 165 $(R_1, \dots, R_N)$ and g is a polynomial of degree at most P . Then degree- P NTD-PL with the same
 166 Tucker rank can represent $\mathcal{Y}$ exactly.*

167 **Theorem 3** (Analytic link approximation). *Suppose $\mathcal{Y} = g(\mathcal{S}^*)$, where $\mathcal{S}^*$ has Tucker rank at most
 168 $(R_1, \dots, R_N)$ and $\|\mathcal{S}^*\|_\infty \leq B$. If g is analytic on the complex disk $D(0, \rho)$ with $B < \rho$, and
 169 $M_\rho = \sup_{|z|=\rho} |g(z)|$, then there exists a degree- P NTD-PL model with the same Tucker backbone
 170 satisfying*

$$\frac{1}{\sqrt{M}} \|\mathcal{Y} - \hat{\mathcal{Y}}\|_F \leq M_\rho \frac{(B/\rho)^{P+1}}{1 - B/\rho}, \quad M = \prod_{n=1}^N I_n. \quad (8)$$

171 Thus the same finite link that induces a tied rank lift also gives a direct approximation route for
 172 smooth responses on a bounded latent range. The statement uses the simple Taylor truncation bound;
 173 sharper polynomial approximation bounds can replace it under stronger assumptions.

174 Together, the results clarify the scope of NTD-PL. Polynomial links can approximate nonlinear
 175 responses while keeping the same latent Tucker rank, and their powers induce Veronese-lifted Tucker
 176 factors. This explains how a scalar link can change multilinear rank structure without introducing
 177 untied high-order tensor parameters.

178 5 Optimization

179 We optimize the masked objective in Section 3.3 by alternating between the scalar link and the Tucker
 180 backbone. The nonlinear part is low-dimensional once the latent Tucker tensor is fixed, while the
 181 backbone update remains a standard differentiable Tucker fitting step through the current link. We
 182 call the two stages *Link Refresh* and *Linked-Backbone Descent*.

183 **Link Refresh.** Fix the current latent tensor

$$S = \llbracket \mathcal{X}; \mathbf{A}^{(1)}, \dots, \mathbf{A}^{(N)} \rrbracket. \quad (9)$$

184 For active degree $Q \leq P$, updating the polynomial coefficients reduces to a small masked ridge
 185 regression,

$$\min_{\beta_{0:Q}} \frac{1}{2|\Omega|} \|y_\Omega - \Phi_{\Omega,Q}(S)\beta_{0:Q}\|_2^2 + \frac{\lambda_\beta}{2} \|\beta_{0:Q}\|_2^2, \quad (10)$$

186 where $\Phi_{\Omega,Q}(S)$ contains the observed powers $[1, s, \dots, s^Q]$. This step is cheap because it solves
 187 only a $(Q+1) \times (Q+1)$ linear system. It also gives an interpretable view of the link: once the
 188 backbone exposes scalar latent values over the tensor, the refresh estimates the best scalar response
 189 on the observed set.

190 **Linked-Backbone Descent.** With β fixed, the Tucker parameters are updated through the current
 191 nonlinear measurement map. The masked residual is first propagated to the latent tensor as

$$T = \frac{1}{|\Omega|} \Omega \odot (f_\beta(S) - Y) \odot f'_\beta(S), \quad (11)$$

192 and then backpropagated through the Tucker reconstruction to update $\mathcal{X}$ and $\{\mathbf{A}^{(n)}\}$. In implemen-
 193 tation we use Adam for these first-order steps, but the update itself is just standard Tucker reverse
 194 contractions composed with the current link. The backbone is therefore not calibrated after fitting; it
 195 is jointly reshaped under the nonlinear link.

196 **Stabilization.** Directly activating all polynomial degrees from the start is brittle, because higher
 197 powers of S amplify scale ambiguity. We therefore begin from the Tucker-compatible degree-one
 198 model and activate higher degrees gradually. When a new degree is switched on, its coefficient is
 199 initialized at zero, so the model expands from the current lower-degree solution without changing
 200 the prediction abruptly. We also normalize Tucker factor scales and absorb the corresponding
 201 scales into the core, which fixes a Tucker gauge without changing S . For the coefficient solve, the
 202 implementation may regress on a centered and scaled latent variable and then map the solution back
 203 to the original power basis. Appendix Figure 6 shows representative training curves where activating
 204 the relevant higher degrees produces systematic loss drops. Full gradients, stabilization details, costs
 205 analysis and pseudocode are given in Appendix A.2.

206 **Masked uniform convergence.** The masked objective also admits a standard sampling in-
 207 terpretation. Let $m = |\Omega|$ and let $\Theta \subset \mathbb{R}^d$ be a bounded NTD-PL parameter set, where
 208 $d = \prod_{n=1}^N R_n + \sum_{n=1}^N I_n R_n + (P + 1)$. If the per-sample losses are bounded by $0 \leq \ell_i(\theta) \leq B_\ell$
 209 and G -Lipschitz in θ , then with probability at least $1 - \delta$,

$$\sup_{\theta \in \Theta} |\mathcal{L}(\theta) - \hat{\mathcal{L}}_m(\theta)| \leq 2G\varepsilon + B_\ell \sqrt{\frac{d \log(3B_\Theta/\varepsilon) + \log(2/\delta)}{2m}}, \quad (12)$$

210 for every $0 < \varepsilon \leq B_\Theta$. This is a standard uniform convergence bound, not a sharp minimax result,
 211 but it clarifies that masked training controls the full-tensor risk with a complexity term determined by
 212 the tied NTD-PL parameter count. The proof is given in Appendix B.4.

213 6 Experiments

214 The experiments test three mechanism questions: whether NTD-PL behaves as a parameter-tied rank
 215 lift under controlled nonlinear residuals; whether the polynomial link improves matched-rank Tucker
 216 on real reconstruction and held-out completion; and whether D_{link} diagnoses beneficial tensor units
 217 across domains. Unless stated otherwise, real tensors are globally max-normalized and used without
 218 spatial downsampling or cropping. We report RMSE and SAM for reconstruction, and RMSE* and
 219 SAM* only on held-out entries for completion. NTD-PL is compared mainly against same-rank
 220 Tucker, except in the rank-lift ablation.

221 6.1 Controlled Nonlinear Residuals

222 These synthetic experiments are mechanism checks rather than generic benchmarks. We generate
 223 measured tensors as

$$Y_\alpha = \sqrt{1 - \alpha} S^* + \sqrt{\alpha} \tilde{R}^*,$$

224 where S^* is a low-rank Tucker tensor and $\tilde{R}^*$ is the normalized component of a scalar nonlinear
 225 response orthogonal to S^* . Thus α controls nonlinear residual energy, not just the raw amplitude
 226 of a nonlinearity. Figure 1 shows the expected transition: NTD-PL is Tucker-like when the residual
 227 is nearly linear and separates as nonlinear residual energy increases. Tucker can lag behind CP,
 228 but NTD-PL then pulls ahead systematically. Figure 2 shows the complementary degree effect:
 229 gains emerge when the active polynomial degree can match or approximate the response complexity.
 230 Degree continuation and coefficient-contribution details are deferred to Appendix C.1, where Table 4
 231 shows the linear consistency negative control. Taken together, these checks indicate that NTD-PL is
 232 not a blanket improvement rule; its gain depends on the data’s nonlinear structure.




Figure 1: Controlled nonlinear residual sweep. NTD-PL stays Tucker-like near the linear regime and separates as α , the nonlinear residual energy, increases. Lower RMSE is better.




Figure 2: Degree-dependent gains at $\alpha = 0.25$. Cell values are Tucker minus NTD-PL RMSE; gains appear when the active degree is sufficient for the response complexity.

233 6.2 Tucker Enhancement on Real Hyperspectral Tensors

234 This section tests whether the same fixed-rank enhancement appears on real hyperspectral tensors,
 235 both when the tensor is fully measured and when evaluation is performed on held-out missing entries.
 236 We use the CAVE multispectral dataset [Columbia CAVE Laboratory, n.d., Yasuma et al., 2010]
 237 because it offers multiple real scenes with rich spectral and spatial structure, making it a reasonable
 238 mechanism check rather than a toy-sized tensor.

239 **Representation enhancement.** Full reconstruction is a representation-capacity check, not a gener-
 240 alization test. Figure 3 shows that NTD-PL improves RMSE and SAM on CAVE at matched Tucker
 241 backbone rank and similar parameter budget. The gap shrinks as Tucker rank increases, consistent
 242 with a fixed-rank enhancement: ordinary multilinear capacity already explains more of the tensor.
 243 Table 1 gives the complementary rank-lift ablation. Tucker can match the NTD-PL RMSE only after
 244 spending substantially more ordinary rank and parameters, supporting the parameter-tied rank-lift
 245 interpretation. Even when RMSE is close, NTD-PL usually retains a clearer SAM advantage.




Figure 3: CAVE full reconstruction as a representation-capacity check. At matched Tucker backbone rank, NTD-PL reduces RMSE and SAM; annotations show paired relative gains over Tucker at each rank. The gap shrinks as Tucker rank increases, consistent with fixed-rank enhancement. Each point averages 15 scenes.

246 **Held-out enhancement.** Held-out completion is the stronger evidence because it tests whether the
 247 learned geometry transfers to unseen entries. Table 2 reports CAVE completion under random masks,

Table 1: Rank-lifted Tucker comparison on CAVE reconstruction. For each NTD-PL backbone rank, we increase Tucker’s spatial ranks while keeping the spectral rank fixed and report the first Tucker rank that matches NTD-PL RMSE, along with the extra Tucker parameters required.

NTD-PL reference				First Tucker matching RMSE				Extra params
Rank	Params	RMSE↓	SAM↓	Rank	Params	RMSE↓	SAM↓	
(18, 18, 3)	20k	0.0319	19.23	(28, 28, 3)	31k	0.0316	21.41	+59.5%
(24, 24, 4)	27k	0.0256	15.15	(35, 35, 4)	41k	0.0254	16.26	+51.3%
(33, 33, 4)	38k	0.0223	14.41	(49, 49, 4)	60k	0.0222	15.28	+56.5%

with RMSE* and SAM* computed only on missing entries. At the same rank budget, NTD-PL gives consistent held-out RMSE* gains and usually improves SAM* as well. The paired sign and Wilcoxon tests in Appendix C.2 show the same scene-level pattern.

Table 2: Held-out CAVE completion at rank (33, 33, 4) under random masks. RMSE* and SAM* are computed only on missing entries. Columns report mean $\pm$ standard deviation over scenes; gains are paired relative improvements over matched-rank Tucker.

Missing rate	RMSE*↓			SAM*↓		
	Tucker	NTD-PL	Gain	Tucker	NTD-PL	Gain
0.1	0.0263 $\pm$ 0.0179	0.0229 $\pm$ 0.0158	12.9%	11.82 $\pm$ 5.33	8.36 $\pm$ 3.38	29.3%
0.3	0.0263 $\pm$ 0.0179	0.0230 $\pm$ 0.0158	12.5%	14.75 $\pm$ 6.33	12.47 $\pm$ 5.06	15.5%
0.5	0.0264 $\pm$ 0.0179	0.0230 $\pm$ 0.0158	12.9%	15.56 $\pm$ 6.85	13.48 $\pm$ 5.75	13.4%
0.7	0.0266 $\pm$ 0.0180	0.0219 $\pm$ 0.0128	17.7%	16.02 $\pm$ 7.21	13.79 $\pm$ 5.43	13.9%

6.3 Diagnosing Nonlinear Low-Rank Structure Across Data

NTD-PL is meant to be diagnosable rather than uniformly beneficial. Its gain should appear when matched-rank Tucker leaves a coherent residual that a scalar link can reduce. We measure this with D_{link} , a cheap frozen-backbone Link Refresh probe: after fitting Tucker, freeze the backbone, run one scalar Link Refresh ridge solve, and measure residual-energy reduction on the relevant evaluation samples. Larger D_{link} indicates more link-like residual structure:

$$D_{\text{link}} = 10 \log_{10} \frac{\|Y_{\text{eval}} - \hat{Y}_T\|_F^2}{\|Y_{\text{eval}} - \hat{Y}_{\text{link}}\|_F^2},$$

where $\hat{Y}_T$ is the Tucker output and $\hat{Y}_{\text{link}}$ is the output after the one-step refresh.

For the CAVE completion, Table 3 groups scenes by D_{link} tercile. Higher D_{link} consistently corresponds to larger held-out RMSE* gains across missing rates, with Spearman correlations around 0.829–0.889. Appendix Figure 7 shows the corresponding scene-level scatter plots.

Table 3: CAVE completion diagnostic. Scenes are grouped by D_{link} terciles, and entries report mean NTD-PL RMSE* gains over matched-rank Tucker. Larger D_{link} consistently corresponds to larger held-out gains across missing rates.

Missing rate	Low D_{link}	Mid D_{link}	High D_{link}	Spearman ρ_s
0.1	0.75%	7.92%	12.63%	0.829
0.3	-1.13%	9.10%	13.12%	0.871
0.5	-1.09%	9.16%	13.18%	0.889
0.7	-0.62%	9.22%	13.07%	0.868

We then apply the same per-unit protocol beyond HSI. Each natural image is one tensor unit; each object or instance forms an object-view tensor; each KTH segment or UCF101 clip is an action-video tensor. Figure 4 plots D_{link} against NTD-PL RMSE gain across these independently fitted units. Diagnostic labels and terciles are assigned from D_{link} , not from NTD-PL gain, and the relationship is domain-dependent but positive: natural images show weaker link yield, object-view tensors clearer yield, and action-video clips an intermediate pattern.




Figure 4: Per-unit cross-domain diagnostic correlation. Each panel plots D_{link} against NTD-PL RMSE gain over independently fitted tensor units.

These results support the scoped claim that NTD-PL helps when matched-rank Tucker leaves link-like residual structure, not uniformly across all visual tensors. Appendix Table 7 gives the numeric statistics behind Figure 4, and Appendix Figure 7 gives the CAVE scene-level scatter supplement.

7 Limitations

NTD-PL assumes that a common nonlinear response explains a coherent part of the residual left by matched-rank Tucker. This scalar link may be less effective when nonlinearities vary strongly across spatial regions, spectral bands, modes, or samples. In such cases, a single polynomial response can underfit the domain variation even if the Tucker backbone is appropriate.

NTD-PL enhances a fixed Tucker rank budget by changing the measurement map on top of the latent tensor. When the remaining error is ordinary multilinear structure rather than a scalar response, increasing Tucker rank remains the more direct modeling choice.

Polynomial degree introduces a second tradeoff. Higher degrees increase the approximation capacity of the link, but they can amplify Tucker scale ambiguity and numerical conditioning issues in the power basis. This motivates the ridge-regularized Link Refresh step, factor normalization, and degree continuation used in the optimizer.

Finally, the joint problem remains nonconvex. The Link Refresh subproblem is a small ridge regression once the latent tensor is fixed, but the Tucker backbone and polynomial link are learned jointly through a nonconvex objective. We do not claim global optimality.

Our strongest empirical evidence is currently from controlled nonlinear residuals and hyperspectral reconstruction and completion. Broader validation on large sparse tensors, temporal or multimodal tensors, and domain-varying nonlinear responses is left for future work.

8 Conclusion

This paper develops polynomial links as parameter-tied rank lifts for Tucker tensors. The model keeps a Tucker latent tensor as the backbone and learns a polynomial link, adding only scalar coefficients. The resulting powers of the same latent tensor act as shared-backbone high-order interactions rather than untied high-rank corrections.

The analysis makes this rank-lift view explicit through Veronese rank expansion, separation from matched-rank Tucker, and approximation results for nonlinear links. The optimization uses a

masked objective with Link Refresh, Linked-Backbone Descent, and degree continuation, while the experiments and diagnostics identify the intended regime: the link helps when matched-rank Tucker leaves a coherent nonlinear residual explained by a scalar response.

Future work may consider spatially or spectrally adaptive links, scalable sparse implementations, and richer structured links beyond scalar polynomials.

References

- Dawon Ahn, Uday Singh Saini, Evangelos E. Papalexakis, and Ali Payani. Neural additive tensor decomposition for sparse tensors. In *Proceedings of the 33rd ACM International Conference on Information and Knowledge Management*, pages 14–23, 2024. doi: 10.1145/3627673.3679833.
- José M. Bioucas-Dias, Antonio Plaza, Nicolas Dobigeon, Mario Parente, Qian Du, Paul Gader, and Jocelyn Chanussot. Hyperspectral unmixing overview: Geometrical, statistical, and sparse regression-based approaches. *IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing*, 5(2):354–379, 2012. doi: 10.1109/JSTARS.2012.2194696.
- Stephen Boyd and Leon O. Chua. Fading memory and the problem of approximating nonlinear operators with volterra series. *IEEE Transactions on Circuits and Systems*, 32(11):1150–1161, 1985. doi: 10.1109/TCS.1985.1085649.
- Columbia CAVE Laboratory. Multispectral image database. <https://cave.cs.columbia.edu/repository/Multispectral>, n.d. Accessed: 2026-04-13.
- Lieven De Lathauwer, Bart De Moor, and Joos Vandewalle. A multilinear singular value decomposition. *SIAM Journal on Matrix Analysis and Applications*, 21(4):1253–1278, 2000. doi: 10.1137/S0895479896305696.
- Paul E. Debevec and Jitendra Malik. Recovering high dynamic range radiance maps from photographs. In *Proceedings of the 24th Annual Conference on Computer Graphics and Interactive Techniques, SIGGRAPH ’97*, pages 369–378, 1997. doi: 10.1145/258734.258884.
- Jian Deng, Haotian Liu, Kim Batselier, Yu-Kwong Kwok, and Ngai Wong. Storm: a nonlinear model order reduction method via symmetric tensor decomposition. In *2016 21st Asia and South Pacific design automation conference (ASP-DAC)*, pages 557–562. IEEE, 2016.
- Michael D. Grossberg and Shree K. Nayar. Modeling the space of camera response functions. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 26(10):1272–1282, 2004. doi: 10.1109/TPAMI.2004.88.
- Rob Heylen, Mario Parente, and Paul Gader. A review of nonlinear hyperspectral unmixing methods. *IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing*, 7(6):1844–1868, 2014. doi: 10.1109/JSTARS.2014.2320576.
- David Hong, Tamara G Kolda, and Jed A Duersch. Generalized canonical polyadic tensor decomposition. *SIAM review*, 62(1):133–163, 2020.
- Tamara G. Kolda and Brett W. Bader. Tensor decompositions and applications. *SIAM Review*, 51(3): 455–500, 2009. doi: 10.1137/07070111X.
- Alex Krizhevsky. Learning multiple layers of features from tiny images. Technical report, University of Toronto, 2009.
- Yann LeCun, Fu Jie Huang, and Léon Bottou. Learning methods for generic object recognition with invariance to pose and lighting. In *Proceedings of the 2004 IEEE Computer Society Conference on Computer Vision and Pattern Recognition*, volume 2, pages 97–104, 2004. doi: 10.1109/CVPR.2004.1315150.
- Allen Liu and Ankur Moitra. Tensor completion made practical. In *Advances in Neural Information Processing Systems*, volume 33, pages 18905–18916, 2020. URL <https://proceedings.neurips.cc/paper/2020/hash/dab1263d1e6a88c9ba5e7e294def5e8b-Abstract.html>.

- 341 Bin Liu, Lirong He, Shandian Zhe, Yingmin Li, and Zenglin Xu. Deepcp: Flexible nonlinear tensor
342 decomposition. In *NeurIPS Workshops*, volume 3, 2017.
- 343 Haotian Liu, Luca Daniel, and Ngai Wong. Model reduction and simulation of nonlinear circuits via
344 tensor decomposition. *IEEE Transactions on Computer-Aided Design of Integrated Circuits and*
345 *Systems*, 34(7):1059–1069, 2015.
- 346 Ji Liu, Przemyslaw Musialski, Peter Wonka, and Jieping Ye. Tensor completion for estimating
347 missing values in visual data. *IEEE Transactions on Pattern Analysis and Machine Intelligence*,
348 35(1):208–220, 2013. doi: 10.1109/TPAMI.2012.39.
- 349 Yi-Si Luo, Xi-Le Zhao, Tai-Xiang Jiang, Yi Chang, Michael K. Ng, and Chao Li. Self-supervised
350 nonlinear transform-based tensor nuclear norm for multi-dimensional image recovery. *IEEE*
351 *Transactions on Image Processing*, 31:3793–3808, 2022. doi: 10.1109/TIP.2022.3176220.
- 352 Yisi Luo, Xile Zhao, Zhemin Li, Michael K. Ng, and Deyu Meng. Low-rank tensor function
353 representation for multi-dimensional data recovery. *IEEE Transactions on Pattern Analysis and*
354 *Machine Intelligence*, 46(5):3351–3369, 2024. doi: 10.1109/TPAMI.2023.3341688.
- 355 David Martin, Charless Fowlkes, Doron Tal, and Jitendra Malik. A database of human segmented
356 natural images and its application to evaluating segmentation algorithms and measuring ecological
357 statistics. In *Proceedings of the IEEE International Conference on Computer Vision*, pages
358 416–423, 2001.
- 359 Sameer A. Nene, Shree K. Nayar, and Hiroshi Murase. Columbia Object Image Library (COIL-100).
360 Technical Report CUCS-006-96, Columbia University, 1996.
- 361 Ivan V. Oseledets. Tensor-train decomposition. *SIAM Journal on Scientific Computing*, 33(5):
362 2295–2317, 2011.
- 363 Zhimeng Pan, Zheng Wang, and Shandian Zhe. Streaming nonlinear bayesian tensor decomposition.
364 In *Conference on Uncertainty in Artificial Intelligence*, pages 490–499. PMLR, 2020.
- 365 David Ruegamer. Scalable higher-order tensor product spline models. In *Proceedings of The 27th*
366 *International Conference on Artificial Intelligence and Statistics*, volume 238 of *Proceedings*
367 *of Machine Learning Research*, pages 1–9. PMLR, 2024. URL [https://proceedings.mlr.](https://proceedings.mlr.press/v238/ruegamer24a.html)
368 [press/v238/ruegamer24a.html](https://proceedings.mlr.press/v238/ruegamer24a.html).
- 369 Vishwanath Saragadam, Randall Balestrieri, Ashok Veeraraghavan, and Richard G Baraniuk.
370 Deeptensor: Low-rank tensor decomposition with deep network priors. *IEEE Transactions*
371 *on Pattern Analysis and Machine Intelligence*, 46(12):10337–10348, 2024.
- 372 Aaron Schein, Mingyuan Zhou, David Blei, and Hanna Wallach. Bayesian poisson tucker decompo-
373 sition for learning the structure of international relations. In *International Conference on Machine*
374 *Learning*, pages 2810–2819. PMLR, 2016.
- 375 Christian Schuldt, Ivan Laptev, and Barbara Caputo. Recognizing human actions: A local SVM
376 approach. In *Proceedings of the 17th International Conference on Pattern Recognition*, volume 3,
377 pages 32–36, 2004. doi: 10.1109/ICPR.2004.1334462.
- 378 Nicholas D. Sidiropoulos, Lieven De Lathauwer, Xiao Fu, Kejun Huang, Evangelos E. Papalexakis,
379 and Christos Faloutsos. Tensor decomposition for signal processing and machine learning. *IEEE*
380 *Transactions on Signal Processing*, 65(13):3551–3582, 2017. doi: 10.1109/TSP.2017.2690524.
- 381 Khurram Soomro, Amir Roshan Zamir, and Mubarak Shah. UCF101: A dataset of 101 hu-
382 man action classes from videos in the wild. Technical Report CRCV-TR-12-01, University
383 of Central Florida, 2012. URL [https://www.crcv.ucf.edu/wp-content/uploads/2019/](https://www.crcv.ucf.edu/wp-content/uploads/2019/03/UCF101_CRCV-TR-12-01.pdf)
384 [03/UCF101_CRCV-TR-12-01.pdf](https://www.crcv.ucf.edu/wp-content/uploads/2019/03/UCF101_CRCV-TR-12-01.pdf).
- 385 Uras Varolgunes, Dan Zhou, Dantong Yu, and Ajim Uddin. Nmtucker: Non-linear matryoshka
386 tucker decomposition for financial time series imputation. In *Proceedings of the Fourth ACM*
387 *International Conference on AI in Finance*, pages 516–523, 2023.

- 388 Vito Volterra. *Theory of Functionals and of Integral and Integro-Differential Equations*. Dover
389 Publications, 1959.
- 390 Jian-Li Wang, Ting-Zhu Huang, Xi-Le Zhao, Tai-Xiang Jiang, and Michael K. Ng. Multi-dimensional
391 visual data completion via low-rank tensor representation under coupled transform. *IEEE Transac-*
392 *tions on Image Processing*, 30:3581–3596, 2021. doi: 10.1109/TIP.2021.3062995.
- 393 Jian-Li Wang, Ting-Zhu Huang, Xi-Le Zhao, Yi-Si Luo, and Tai-Xiang Jiang. CoNoT: Coupled
394 nonlinear transform-based low-rank tensor representation for multidimensional image completion.
395 *IEEE Transactions on Neural Networks and Learning Systems*, 35(7):8969–8983, 2024. doi:
396 10.1109/TNNLS.2022.3217198.
- 397 Shuang Xu, Jiangjun Peng, Teng-Yu Ji, Xiangyong Cao, Kai Sun, Rongrong Fei, and Deyu Meng.
398 Stacked tucker decomposition with multi-nonlinear products for remote sensing imagery inpainting.
399 *IEEE Transactions on Geoscience and Remote Sensing*, 62:1–13, 2024.
- 400 Zenglin Xu, Feng Yan, et al. Infinite tucker decomposition: Nonparametric bayesian models for
401 multiway data analysis. *arXiv preprint arXiv:1108.6296*, 2011.
- 402 Fumihito Yasuma, Tomoo Mitsunaga, Daisuke Iso, and Shree K. Nayar. Generalized assorted pixel
403 camera: Postcapture control of resolution, dynamic range, and spectrum. *IEEE Transactions on*
404 *Image Processing*, 19(9):2241–2253, 2010. doi: 10.1109/TIP.2010.2046811.
- 405 Zheng Zhai, Hengchao Chen, and Qiang Sun. Quadratic matrix factorization with applications
406 to manifold learning. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 46(9):
407 6384–6401, 2024.
- 408 Qibin Zhao, Guoxu Zhou, Shengli Xie, Liqing Zhang, and Andrzej Cichocki. Tensor ring decomposi-
409 tion. *arXiv preprint arXiv:1606.05535*, 2016.
- 410 Shandian Zhe, Kai Zhang, Pengyuan Wang, Kuang-chih Lee, Zenglin Xu, Yuan Qi, and Zoubin
411 Ghahramani. Distributed flexible nonlinear tensor factorization. *Advances in neural information*
412 *processing systems*, 29, 2016.
- 413 Yu-Bang Zheng, Ting-Zhu Huang, Xi-Le Zhao, Qibin Zhao, and Tai-Xiang Jiang. Fully-connected
414 tensor network decomposition and its application to higher-order tensor completion. In *Proceedings*
415 *of the AAAI Conference on Artificial Intelligence*, volume 35, pages 11071–11078, 2021. doi:
416 10.1609/aaai.v35i12.17321.

417 A Additional Model Details

418 A.1 Explicit p -Fold Tucker Motivation

419 This section expands the motivation behind the explicit p -fold Tucker channel in Section 3.2. The
420 construction is best viewed as the tensor analogue of adding polynomial interactions to a linear map.
421 In the matrix case, one may move from a linear model $\mathbf{y} = \mathbf{A}\mathbf{x}$ to a quadratic model

$$\mathbf{y} = \mathbf{A}\mathbf{x} + \mathcal{A}(\mathbf{x}, \mathbf{x}),$$

422 where the second term couples pairs of latent coordinates. The corresponding question for Tucker is:
423 if standard Tucker is the multilinear low-rank map for multiway data, what is the direct p -th order
424 interaction analogue?

425 For a Tucker backbone, an explicit p -fold interaction can be written as

$$\tilde{Y}_i^{(p)} = \sum_{\mathbf{r}^{(1)}, \dots, \mathbf{r}^{(p)}} \left(\prod_{q=1}^p \mathcal{X}_{r_1^{(q)}, \dots, r_N^{(q)}} \right) \prod_{n=1}^N \mathcal{A}_{i_n, r_n^{(1)}, \dots, r_n^{(p)}}^{(n,p)}. \quad (13)$$

426 Here each mode projection depends on p latent indices instead of one. For $p = 1$, this reduces to the
427 ordinary Tucker entry formula. For $p = 2$, the same core appears twice and the mode projections
428 couple two latent tuples, giving a quadratic Tucker-style channel. Figure 5 illustrates this progression.




Figure 5: Tensor-network view of a formal p -fold Tucker extension. The construction keeps the Tucker idea of mode-wise generation, but lets each mode projection couple multiple latent index tuples.

The explicit construction is useful as a reference model, but it is not the parameterization used by NTD-PL. Its mode projections $\mathcal{A}^{(n,p)}$ have size $I_n R_n^p$, so a direct implementation introduces new high-order projection tensors and rapidly increases storage and contraction cost. NTD-PL keeps the same high-order interaction pattern by tying the projection tensor to the original Tucker factor:

$$\mathcal{A}_{i_n, r_n^{(1)}, \dots, r_n^{(p)}}^{(n,p)} = \prod_{q=1}^p A_{i_n, r_n^{(q)}}^{(n)}. \quad (14)$$

Substituting (14) into (13) gives exactly $\mathcal{S}^{\odot p}$, where $\mathcal{S} = \llbracket \mathcal{X}; \mathbf{A}^{(1)}, \dots, \mathbf{A}^{(N)} \rrbracket$. Thus the degree- p NTD-PL channel is a shared-projection version of the explicit p -fold Tucker prototype.

This is the modeling role of the p -fold Tucker prototype: it defines the high-order Tucker interaction that NTD-PL compresses. The final model does not train the untied tensors $\mathcal{A}^{(n,p)}$; it keeps a single Tucker backbone and adds only the polynomial coefficients $\{\beta_p\}$.

A.2 Optimization Details

This section gives the update equations behind the solver summarized in Section 5. The implementation follows the same structure: one Tucker reconstruction per iteration, a fused evaluation of the polynomial and its derivative, first-order updates of the backbone, and periodic ridge-regression updates of the link coefficients.

Backbone gradients. For a fixed β , define

$$f'_\beta(\mathcal{S}) = \sum_{p=1}^P p \beta_p \mathcal{S}^{\odot(p-1)}. \quad (15)$$

The gradient of the normalized masked data term with respect to the latent tensor $\mathcal{S}$ is

$$\mathbf{T} = \frac{1}{|\Omega|} \Omega \odot (f_\beta(\mathcal{S}) - \mathcal{Y}) \odot f'_\beta(\mathcal{S}). \quad (16)$$

Backpropagating $\mathbf{T}$ through the Tucker reconstruction gives the core gradient

$$\nabla_{\mathcal{X}} \mathcal{L} = \mathbf{T} \times_1 \mathbf{A}^{(1)\top} \times_2 \dots \times_N \mathbf{A}^{(N)\top} + \lambda_X \mathcal{X}. \quad (17)$$

For the n -th factor, let

$$\mathbf{Z}^{(n)} = \left[\mathcal{X} \times_1 \mathbf{A}^{(1)} \dots \times_{n-1} \mathbf{A}^{(n-1)} \times_{n+1} \mathbf{A}^{(n+1)} \dots \times_N \mathbf{A}^{(N)} \right]_{(n)}. \quad (18)$$

Then

$$\nabla_{\mathbf{A}^{(n)}} \mathcal{L} = \mathbf{T}_{(n)} \mathbf{Z}^{(n)\top} + \lambda_n \mathbf{A}^{(n)}. \quad (19)$$

The experiments use Adam for these first-order steps. This choice is not part of the model definition; any comparable first-order optimizer can use the same gradients.

450 **Link update.** With S fixed and active degree $Q = P_{\text{act}}$, the coefficient subproblem is the ridge
 451 regression in (10). Its closed-form solution is

$$\beta_{0:Q}^* = (\Phi_{\Omega,Q}^\top \Phi_{\Omega,Q} + |\Omega| \lambda_\beta \mathbf{I})^{-1} \Phi_{\Omega,Q}^\top \mathbf{y}_\Omega. \quad (20)$$

452 Equivalently, one may absorb $|\Omega|$ into the ridge parameter and solve the usual unnormalized normal
 453 equations. If the constant term is disabled, the first column is removed and β_0 is fixed at zero. In
 454 the code, this small system is solved either from an explicit Vandermonde design or from moment
 455 equations built from cached powers.

456 **Stabilization and continuation.** High powers of S can be sensitive to scale. The implementation
 457 therefore normalizes factor columns after backbone updates and absorbs the corresponding scales
 458 into the core; this is a gauge fixing of the Tucker parameterization and does not change $\mathcal{S}$. For the
 459 coefficient solve, it may regress on the centered and scaled latent variable $(S - \mu)/\sigma$, then convert
 460 the solution back to the original power basis of S .

461 The active degree follows the schedule

$$t_k = \text{round} \left(\frac{kT}{P} \right), \quad k = 1, \dots, P-1.$$

462 When degree $p+1$ is activated, the previous coefficients and Tucker factors are kept and the new
 463 coefficient β_{p+1} is initialized at zero.

464 **Cost.** The dominant cost is forming the Tucker reconstruction and its reverse contractions on the
 465 observed tensor. For dense tensors this is comparable to standard Tucker optimization at the same
 466 rank, plus $O(P|\Omega|)$ work to evaluate the scalar polynomial and its derivative. The link update forms
 467 a Vandermonde design or the corresponding moment equations and solves a $(P+1) \times (P+1)$ linear
 468 system, with direct cost $O(P^2|\Omega| + P^3)$. Since P is small, this cost is negligible compared with the
 469 Tucker contractions.

470 Algorithm 1 gives the full procedural summary.

Algorithm 1 Masked NTD-PL optimization

Require: observed tensor $\mathcal{Y}$, mask Ω , rank $(R_1, \dots, R_N)$, degree P

```

1: initialize  $\mathcal{X}, \{\mathbf{A}^{(n)}\}$  from Tucker on the masked tensor
2: initialize active degree  $P_{\text{act}} \leftarrow 1$  and  $\beta = (0, 1)$ 
3: for  $t = 1, \dots, T$  do
4:   form  $\mathcal{S} = \llbracket \mathcal{X}; \mathbf{A}^{(1)}, \dots, \mathbf{A}^{(N)} \rrbracket$ 
5:   if a link update is scheduled then
6:     update active coefficients  $\beta_{0:P_{\text{act}}}$  by ridge regression on  $\{(S_i, Y_i) : \Omega_i = 1\}$ 
7:   end if
8:   compute  $f_\beta(\mathcal{S})$  and  $f'_\beta(\mathcal{S})$ 
9:   take first-order steps on  $\mathcal{X}$  and  $\{\mathbf{A}^{(n)}\}$  using the masked gradient
10:  normalize factor scales and absorb them into the core
11:  increase  $P_{\text{act}}$  according to the continuation schedule, initializing the new coefficient at zero
12: end for
13: return  $\mathcal{X}, \{\mathbf{A}^{(n)}\}, \beta$ 

```

471 B Additional Proofs

472 B.1 Polynomial and Analytic Approximation

473 *Proof of Theorem 2.* Since g has degree at most P , write

$$g(s) = \sum_{p=0}^P c_p s^p. \quad (21)$$

474 Take the NTD-PL latent tensor to be the true latent tensor, $\mathcal{S} = \mathcal{S}^*$, and set $\beta_p = c_p$. Then, entrywise,

$$\hat{\mathcal{Y}}_{i_1, \dots, i_N} = \sum_{p=0}^P \beta_p (\mathcal{S}_{i_1, \dots, i_N}^*)^p = g(\mathcal{S}_{i_1, \dots, i_N}^*) = \mathcal{Y}_{i_1, \dots, i_N}. \quad (22)$$

475 Thus $\hat{\mathcal{Y}} = \mathcal{Y}$. □

476 **Lemma 1** (Scalar-to-tensor approximation). *Let $\|\mathcal{S}^*\|_\infty \leq B$. For any scalar function g and any*
 477 *degree- P polynomial q_P ,*

$$\|g(\mathcal{S}^*) - q_P(\mathcal{S}^*)\|_F \leq \sqrt{\prod_{n=1}^N I_n} \sup_{|s| \leq B} |g(s) - q_P(s)|. \quad (23)$$

478 *Proof.* Every entry of $\mathcal{S}^*$ lies in $[-B, B]$, so every entrywise approximation error is bounded by
 479 $\sup_{|s| \leq B} |g(s) - q_P(s)|$. Squaring, summing over all entries, and taking the square root gives the
 480 claim. □

481 *Proof of Theorem 3.* Let $g(s) = \sum_{k=0}^\infty a_k s^k$ be the Taylor expansion of g around zero, and choose
 482 $q_P(s) = \sum_{k=0}^P a_k s^k$. By Cauchy's estimate,

$$|a_k| \leq \frac{M_\rho}{\rho^k}. \quad (24)$$

483 For $|s| \leq B < \rho$,

$$|g(s) - q_P(s)| \leq \sum_{k=P+1}^\infty |a_k| |s|^k \quad (25)$$

$$\leq M_\rho \sum_{k=P+1}^\infty (B/\rho)^k = M_\rho \frac{(B/\rho)^{P+1}}{1 - B/\rho}. \quad (26)$$

484 Applying Lemma 1 and taking $\mathcal{S} = \mathcal{S}^*$, $\beta_k = a_k$ for $0 \leq k \leq P$, yields the stated unnormalized
 485 bound. Dividing by $\sqrt{M}$, $M = \prod_{n=1}^N I_n$, gives the normalized form used in Theorem 3. □

486 B.2 Rank Expansion

487 *Proof of Proposition 1.* Write the Tucker tensor entrywise as

$$\mathcal{S}_{i_1, \dots, i_N} = \sum_{r_1=1}^{R_1} \cdots \sum_{r_N=1}^{R_N} \mathcal{X}_{r_1, \dots, r_N} \prod_{n=1}^N A_{i_n, r_n}^{(n)}. \quad (27)$$

488 Taking an element-wise p -th power gives

$$(\mathcal{S}_{i_1, \dots, i_N})^p = \prod_{q=1}^p \left(\sum_{r_1^{(q)}=1}^{R_1} \cdots \sum_{r_N^{(q)}=1}^{R_N} \mathcal{X}_{r_1^{(q)}, \dots, r_N^{(q)}} \prod_{n=1}^N A_{i_n, r_n^{(q)}}^{(n)} \right) \quad (28)$$

$$= \sum_{\gamma_1 \in [R_1]^p} \cdots \sum_{\gamma_N \in [R_N]^p} \left(\prod_{q=1}^p \mathcal{X}_{r_1^{(q)}, \dots, r_N^{(q)}} \right) \prod_{n=1}^N \left(\prod_{q=1}^p A_{i_n, r_n^{(q)}}^{(n)} \right), \quad (29)$$

489 where $\gamma_n = (r_n^{(1)}, \dots, r_n^{(p)})$. For each mode n , let $\alpha_n \in \mathbb{N}^{R_n}$ record the counts of the entries in γ_n ,
 490 so $|\alpha_n| = p$. Then

$$\prod_{q=1}^p A_{i_n, r_n^{(q)}}^{(n)} = \prod_{r=1}^{R_n} (A_{i_n, r}^{(n)})^{\alpha_{n,r}} = A_{i_n, \alpha_n}^{(n)\langle p \rangle}. \quad (30)$$

491 Group all ordered tuples $(\gamma_1, \dots, \gamma_N)$ that produce the same multi-indices $(\alpha_1, \dots, \alpha_N)$, and define
 492 the aggregated core

$$\mathcal{X}_{\alpha_1, \dots, \alpha_N}^{(p)} = \sum_{\substack{\gamma_1, \dots, \gamma_N: \\ \text{count}(\gamma_n) = \alpha_n \ \forall n}} \prod_{q=1}^p \mathcal{X}_{r_1^{(q)}, \dots, r_N^{(q)}}^{(q)}. \quad (31)$$

493 The preceding display is exactly the Tucker reconstruction

$$\mathcal{S}^{\odot p} = \llbracket \mathcal{X}^{(p)}; \mathbf{A}^{(1)(p)}, \dots, \mathbf{A}^{(N)(p)} \rrbracket. \quad (32)$$

494 The n -th Veronese factor has $\binom{R_n+p-1}{p}$ columns, giving the claimed Tucker rank bound. $\square$

495 B.3 Separation

496 **Lemma 2** (Full-rank Veronese evaluations). *Let $D = \binom{R+p-1}{p}$. There exists a matrix $\mathbf{A} \in \mathbb{R}^{D \times R}$
 497 whose Veronese feature matrix $\mathbf{A}^{(p)} \in \mathbb{R}^{D \times D}$ has rank D . Consequently, for $I \geq D$, a generic
 498 matrix $\mathbf{A} \in \mathbb{R}^{I \times R}$ has $\text{rank}(\mathbf{A}^{(p)}) = D$.*

499 *Proof.* The homogeneous degree- p monomials in R variables are linearly independent polynomials.
 500 Hence there exist evaluation points $x_1, \dots, x_D \in \mathbb{R}^R$ such that the matrix $(x_i^\alpha)_{i,\alpha}$, $|\alpha| = p$, is
 501 nonsingular; otherwise every $D \times D$ evaluation determinant would vanish, which would imply a
 502 nontrivial linear dependence among the monomials. Taking these points as the rows of $\mathbf{A}$ gives the
 503 first claim.

504 For $I \geq D$, choose any D rows. The determinant of the corresponding $D \times D$ Veronese submatrix
 505 is a polynomial in the entries of $\mathbf{A}$ that is not identically zero by the first claim. Its nonvanishing is
 506 therefore a generic property, so $\mathbf{A}^{(p)}$ has full column rank for generic $\mathbf{A}$. $\square$

507 *Proof of Theorem 1.* Let $D = \binom{R+p-1}{p}$. Choose generic factor matrices $\mathbf{A}^{(n)} = [a_1^{(n)}, \dots, a_R^{(n)}] \in$
 508 $\mathbb{R}^{I_n \times R}$, and define the CP-diagonal Tucker backbone

$$\mathcal{S} = \sum_{r=1}^R a_r^{(1)} \otimes \dots \otimes a_r^{(N)}. \quad (33)$$

509 Its mode- n unfolding is

$$\mathcal{S}_{(n)} = \mathbf{A}^{(n)} \left(\mathbf{A}^{(N)} \odot \dots \odot \mathbf{A}^{(n+1)} \odot \mathbf{A}^{(n-1)} \odot \dots \odot \mathbf{A}^{(1)} \right)^\top. \quad (34)$$

510 For generic $\mathbf{A}^{(n)}$, $\mathbf{A}^{(n)}$ has rank R , and the Khatri-Rao product on the right has rank R . Thus
 511 $\text{rank}(\mathcal{S}_{(n)}) = R$ for every n , so $\mathcal{S}$ has exact Tucker rank $(R, \dots, R)$.

512 Set $\beta_p = 1$ and $\beta_q = 0$ for $q \neq p$. Then $\mathcal{Y} = \mathcal{S}^{\odot p}$ is a valid degree- p NTD-PL tensor. Entrywise,

$$\mathcal{Y}_{i_1, \dots, i_N} = \left(\sum_{r=1}^R \prod_{n=1}^N A_{i_n, r}^{(n)} \right)^p \quad (35)$$

$$= \sum_{|\alpha|=p} c_\alpha \prod_{n=1}^N \prod_{r=1}^R \left(A_{i_n, r}^{(n)} \right)^{\alpha_r}, \quad (36)$$

513 where $c_\alpha = p! / (\alpha_1! \dots \alpha_R!)$. Let $\mathbf{B}^{(n)} = \mathbf{A}^{(n)(p)} \in \mathbb{R}^{I_n \times D}$. By Lemma 2, each $\mathbf{B}^{(n)}$ has full
 514 column rank D for generic $\mathbf{A}^{(n)}$. The mode- n unfolding is

$$\mathcal{Y}_{(n)} = \mathbf{B}^{(n)} \text{diag}(c) \left(\mathbf{B}^{(N)} \odot \dots \odot \mathbf{B}^{(n+1)} \odot \mathbf{B}^{(n-1)} \odot \dots \odot \mathbf{B}^{(1)} \right)^\top. \quad (37)$$

515 The diagonal matrix $\text{diag}(c)$ is nonsingular. Since a Khatri-Rao product of full-column-rank matrices
 516 is full column rank, the right factor has rank D . Therefore $\text{rank}(\mathcal{Y}_{(n)}) = D$. This holds for every
 517 mode n , and the exact Tucker rank of $\mathcal{Y}$ is $(D, \dots, D)$. $\square$

518 B.4 Masked Generalization

519 **Theorem 4** (Masked-entry uniform convergence). *Let $\Theta \subset \mathbb{R}^d$ be a bounded NTD-PL parameter set*
 520 *with*

$$d = \prod_{n=1}^N R_n + \sum_{n=1}^N I_n R_n + (P + 1).$$

521 For $\theta \in \Theta$, define the full-entry risk

$$\mathcal{L}(\theta) = \frac{1}{M} \sum_{\mathbf{i}} \ell_{\mathbf{i}}(\theta), \quad M = \prod_{n=1}^N I_n,$$

522 and the masked empirical risk

$$\hat{\mathcal{L}}_m(\theta) = \frac{1}{m} \sum_{t=1}^m \ell_{\mathbf{i}_t}(\theta),$$

523 where $\mathbf{i}_t$ are sampled uniformly from tensor entries. Assume $0 \leq \ell_{\mathbf{i}}(\theta) \leq B_\ell$ and $\ell_{\mathbf{i}}(\theta)$ is G -Lipschitz
 524 in θ , uniformly over $\mathbf{i}$. If Θ is contained in a Euclidean ball of radius B_Θ , then with probability at
 525 least $1 - \delta$,

$$\sup_{\theta \in \Theta} |\mathcal{L}(\theta) - \hat{\mathcal{L}}_m(\theta)| \leq 2G\varepsilon + B_\ell \sqrt{\frac{d \log(3B_\Theta/\varepsilon) + \log(2/\delta)}{2m}}$$

526 for every $0 < \varepsilon \leq B_\Theta$.

527 *Proof of Theorem 4.* Fix $0 < \varepsilon \leq B_\Theta$, and let $\mathcal{N}_\varepsilon$ be an ε -net of Θ in Euclidean norm. Since Θ lies
 528 in a d -dimensional ball of radius B_Θ , it admits a net with

$$|\mathcal{N}_\varepsilon| \leq \left(\frac{3B_\Theta}{\varepsilon} \right)^d. \quad (38)$$

529 For any fixed $\theta_0 \in \mathcal{N}_\varepsilon$, the random variables $\ell_{\mathbf{i}_t}(\theta_0)$ are independent and bounded in $[0, B_\ell]$.
 530 Hoeffding's inequality gives

$$\Pr \left(\left| \mathcal{L}(\theta_0) - \hat{\mathcal{L}}_m(\theta_0) \right| \geq u \right) \leq 2 \exp \left(-\frac{2mu^2}{B_\ell^2} \right). \quad (39)$$

531 Taking a union bound over $\mathcal{N}_\varepsilon$, with probability at least $1 - \delta$,

$$\max_{\theta_0 \in \mathcal{N}_\varepsilon} \left| \mathcal{L}(\theta_0) - \hat{\mathcal{L}}_m(\theta_0) \right| \leq B_\ell \sqrt{\frac{\log(2|\mathcal{N}_\varepsilon|/\delta)}{2m}}. \quad (40)$$

532 Substituting the net-size bound yields the stated stochastic term.

533 Now take any $\theta \in \Theta$, and choose $\theta_0 \in \mathcal{N}_\varepsilon$ with $\|\theta - \theta_0\|_2 \leq \varepsilon$. By the uniform Lipschitz assumption,

$$|\mathcal{L}(\theta) - \mathcal{L}(\theta_0)| \leq G\varepsilon, \quad |\hat{\mathcal{L}}_m(\theta) - \hat{\mathcal{L}}_m(\theta_0)| \leq G\varepsilon. \quad (41)$$

534 Therefore

$$\left| \mathcal{L}(\theta) - \hat{\mathcal{L}}_m(\theta) \right| \leq |\mathcal{L}(\theta) - \mathcal{L}(\theta_0)| + \left| \mathcal{L}(\theta_0) - \hat{\mathcal{L}}_m(\theta_0) \right| \quad (42)$$

$$+ |\hat{\mathcal{L}}_m(\theta_0) - \hat{\mathcal{L}}_m(\theta)| \quad (43)$$

$$\leq 2G\varepsilon + B_\ell \sqrt{\frac{d \log(3B_\Theta/\varepsilon) + \log(2/\delta)}{2m}}. \quad (44)$$

535 Taking the supremum over $\theta \in \Theta$ completes the proof. $\square$

536 **Remark 1.** *The Lipschitz and bounded-loss assumptions are standard for compact parametric*
 537 *classes. For NTD-PL, they follow when the observed values, core, factor matrices, and polynomial*
 538 *coefficients are bounded and the polynomial degree P is finite. The resulting Lipschitz constant may*
 539 *depend polynomially on these bounds and on P , but the covering dimension remains the number*
 540 *of tied NTD-PL parameters rather than the number of parameters in an untied high-order Tucker*
 541 *expansion.*

C Additional Experiments

C.1 Controlled Synthetic Mechanism Checks

The controlled experiments in the main text are designed to isolate the role of the scalar link. Table 4 first checks a negative control: when the source is linear Tucker, NTD-PL does not improve by inventing nonlinear terms. The only exception is an affine offset, where the constant link term is the intended correction.

For the nonlinear controlled tensors, the nonlinearity is introduced as an orthogonal residual rather than by directly setting $Y = g(S^*)$. Let S^* be the linear low-rank Tucker backbone and let g be one of the candidate scalar responses. We form

$$R^* = g(S^*) - \frac{\langle g(S^*), S^* \rangle}{\|S^*\|_F^2} S^*, \quad (45)$$

and, when $\|R^*\|_F > 0$, normalize it as

$$\tilde{R}^* = \frac{\|S^*\|_F}{\|R^*\|_F} R^*. \quad (46)$$

The observed tensor is then

$$Y_\alpha = \sqrt{1-\alpha} S^* + \sqrt{\alpha} \tilde{R}^*, \quad 0 \leq \alpha \leq 1. \quad (47)$$

By construction, $\langle S^*, \tilde{R}^* \rangle = 0$ and $\|\tilde{R}^*\|_F = \|S^*\|_F$, so α is exactly the energy fraction assigned to the nonlinear residual. This keeps the total energy fixed while moving the observation away from the linear Tucker backbone. The candidate residual families are $g(s) = s^2 + s^3$, $g(s) = (\exp(\kappa s) - 1)/\kappa$, and $g(s) = \tanh(\kappa s)$, with κ set from the input range.

The main text reports the nonlinear alpha sweep and degree sweep. Figure 6 shows that degree continuation activates higher powers with systematic loss drops at relevant degrees. Table 5 complements those figures with a small but consistent term-mass concentration effect: the dominant contribution usually lands on the degree that matches the injected response, while smooth odd links spread mass across nearby odd orders.

Table 4: Linear consistency on synthetic Tucker tensors. In the linear negative control, effective $p > 1$ contributions remain small except for the intended affine offset case.

Method	bias = 0		bias = 0.5	
	RMSE↓	$p > 1$	RMSE↓	$p > 1$
Tucker	0.0114	–	0.0647	–
NTD-PL, $\beta_0 = 0$	0.0116	4.19%	0.0522	39.17%
NTD-PL, β_0 free	0.0117	3.19%	0.0118	2.98%




Figure 6: NTD-PL training curves for controlled nonlinear mechanism checks at $\alpha = 0.25$ and $P = 5$. Dashed vertical lines mark degree activations, which produce systematic drops when the activated degree matches the response structure.

Table 5: Effective NTD-PL term contributions under moderate nonlinearity, reported as percentages over ten seeds. Polynomial targets concentrate mass at matching nonlinear degrees, while odd smooth links place most nonlinear mass on odd powers.

Link	Effective contribution (%)						
	$p = 0$	$p = 1$	$p = 2$	$p = 3$	$p = 4$	$p = 5$	$p = 6$
$s^2 + s^3$	0.1	24.1	9.8	38.1	7.3	11.5	9.1
$(e^{\kappa s} - 1)/\kappa$	0.3	39.8	16.8	13.2	10.1	7.7	12.1
$\tanh(\kappa s)$	0.2	37.9	1.3	30.9	5.4	17.3	7.0

562 C.2 CAVE Completion Significance and Robustness

563 The main text reports the full-resolution random-missing CAVE completion table at rank (33, 33, 4).
 564 We do not repeat that table here; instead, Table 6 gives the paired sign and Wilcoxon tests for the
 565 same protocol.

Table 6: Paired significance tests for full-resolution CAVE completion at rank (33, 33, 4). Metrics are computed only on missing entries.

ρ	RMSE				SAM			
	W/L	Med. gain	Sign p	Wilcoxon p	W/L	Med. gain	Sign p	Wilcoxon p
0.1	14/1	0.0027	9.77×10^{-4}	1.22×10^{-4}	14/1	2.9742	9.77×10^{-4}	8.54×10^{-4}
0.3	14/1	0.0027	9.77×10^{-4}	1.22×10^{-4}	13/2	1.9194	0.007	0.004
0.5	14/1	0.0026	9.77×10^{-4}	1.22×10^{-4}	13/2	1.6566	0.007	0.012
0.7	14/1	0.0027	9.77×10^{-4}	1.22×10^{-4}	12/3	1.3394	0.035	0.048

566 C.3 Link Refresh Residual Diagnostic

567 The diagnostic used in the main text reuses the Link Refresh step from Section 5 as a one-round
 568 residual probe. Let $\hat{\mathcal{Y}}_T$ denote the Tucker output and $\mathcal{Y}$ the target tensor. Starting from this fitted
 569 Tucker latent tensor, we update only the scalar polynomial coefficients by fitting

$$h_\gamma(x) = \sum_{p=0}^4 \gamma_p x^p$$

570 by ridge regression on pairs $(\hat{Y}_{T,i}, Y_i)$. In full reconstruction, all entries are used for this fit. In
 571 completion, the fit uses only observed entries and is then evaluated on held-out entries. The fitted
 572 map is applied entrywise to $\hat{\mathcal{Y}}_T$. We summarize this cheap refresh by the Link Refresh diagnostic

$$D_{\text{link}} = 10 \log_{10} \left(\frac{\|\mathcal{Y}_{\text{eval}} - \hat{\mathcal{Y}}_T\|_F^2}{\|\mathcal{Y}_{\text{eval}} - \hat{\mathcal{Y}}_{\text{link}}\|_F^2} \right),$$

573 which measures, on a log scale, how much Tucker residual energy is removed on the task-specific
 574 evaluation entries. In other words, it measures how much shared scalar-link yield is already present
 575 before the full NTD-PL alternating solver starts reshaping the Tucker parameters.

576 C.4 CAVE Completion Scene Diagnostics

577 Figure 7 supplements the main CAVE diagnostic table with the scene-level scatter plots used to
 578 visualize the link yield across missing rates. The plots show the same diagnostic relationship between
 579 D_{link} and held-out RMSE* gain, now at the scene level.

580 C.5 Non-HSI Link Refresh Diagnostics

581 The main diagnostic figure uses per-unit non-HSI experiments rather than collection-level tensors.
 582 CBSD68 [Martin et al., 2001] uses per-image units resized to $96 \times 96 \times 3$ with rank (32, 32, 3).
 583 CIFAR-10 [Krizhevsky, 2009] uses 1000 class-balanced test images of shape $32 \times 32 \times 3$ with rank




Figure 7: Scene-level supplement for the CAVE completion diagnostic table. Missing rates 0.3 and 0.5 are shown; D_{link} from frozen- backbone Link Refresh predicts larger held-out RMSE* gains.

(20, 20, 3). COIL-100 [Nene et al., 1996] uses one object at a time as a $36 \times 48 \times 48 \times 3$ view tensor with rank (8, 12, 12, 2). smallNORB [LeCun et al., 2004] uses one object instance at a time as an $18 \times 64 \times 64 \times 2$ stereo view tensor with rank (12, 24, 24, 2). KTH-Action [Schuld et al., 2004] and UCF101 [Soomro et al., 2012] provide the action-video tensor units used in the main diagnostic scatter. All units are max-normalized independently, and diagnostic labels are assigned only from D_{link} , not from NTD-PL gain.

Table 7: Dataset-level statistics behind the non-HSI diagnostic scatter in Figure 4. The gain columns report mean NTD-PL RMSE improvement over independently fitted matched-rank Tucker within D_{link} terciles.

Domain	Dataset	Units	Median D_{link}	Spearman ρ_s	Low gain	Mid gain	High gain
Natural images	CBSD68	68	0.091	0.85	4.17%	8.40%	14.60%
Natural images	CIFAR-10	1000	0.008	0.93	-0.01%	2.59%	7.85%
Object-view	COIL-100	100	0.208	0.93	2.77%	5.46%	9.14%
Object-view	smallNORB	25	0.264	0.95	6.92%	9.91%	14.16%
Action video	KTH-Action	240	0.107	0.88	2.74%	4.27%	10.87%
Action video	UCF101	240	0.065	0.89	1.44%	2.97%	5.89%

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes].

Justification: The abstract and introduction state the model, theory, optimizer, and scoped empirical claim. They explicitly frame NTD-PL as a Tucker enhancement for coherent nonlinear residuals, not as a universal tensor decoder.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes].

Justification: Section 7 discusses the scalar-link assumption, rank-budget scope, polynomial-degree stability, nonconvexity, and empirical scope.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes].

Justification: The main theory statements list their assumptions in Section 4; Appendix B gives full proofs and the auxiliary lemma used by the separation result.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes].

Justification: Sections 6 and Appendix C.1–C.5 describe datasets, ranks, mask protocols, metrics, diagnostics, and optimization details needed to reproduce the reported comparisons.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [No].

Justification: The submission describes the algorithms, public datasets, and experimental protocols, but the anonymous submission PDF does not include an anonymized code-release URL or exact command manifest for all experiments.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes].

Justification: The experiment section and appendix specify the main ranks, mask rates, held-out metrics, datasets, and optimizer structure; Appendix A.2 gives the update equations and algorithm.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes].

Justification: The CAVE completion table reports mean and standard deviation over scenes, and Appendix C.2 reports paired sign and Wilcoxon tests for the same protocol.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [No].

Justification: The paper reports method runtimes in several tables and gives cost analysis, but it does not specify the hardware, memory, or total compute budget used for each experiment.

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes].

Justification: The work uses mathematical analysis and public benchmark data for tensor modeling, and no human-subjects study, deceptive deployment, or high-risk data release is introduced.

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [N/A].

Justification: The contribution is a foundational tensor-modeling method with experiments on standard public benchmarks; it does not introduce a deployed system, user-facing model, or application-specific decision pipeline.

11. Safeguards

659 Question: Does the paper describe safeguards that have been put in place for responsible
660 release of data or models that have a high risk for misuse (e.g., pre-trained language models,
661 image generators, or scraped datasets)?

662 Answer: [N/A].

663 Justification: The paper does not release a high-risk pretrained model, generative model, or
664 newly scraped dataset requiring special safeguards.

665 **12. Licenses for existing assets**

666 Question: Are the creators or original owners of assets (e.g., code, data, models), used in
667 the paper, properly credited and are the license and terms of use explicitly mentioned and
668 properly respected?

669 Answer: [No].

670 Justification: The datasets used in the experiments are cited, but the manuscript does not list
671 license or terms-of-use details for every dataset.

672 **13. New assets**

673 Question: Are new assets introduced in the paper well documented and is the documentation
674 provided alongside the assets?

675 Answer: [N/A].

676 Justification: The paper introduces a method and experiments, not a new dataset, model
677 checkpoint, or benchmark asset.

678 **14. Crowdsourcing and research with human subjects**

679 Question: For crowdsourcing experiments and research with human subjects, does the paper
680 include the full text of instructions given to participants and screenshots, if applicable, as
681 well as details about compensation (if any)?

682 Answer: [N/A].

683 Justification: The paper does not involve crowdsourcing or human-subjects experiments.

684 **15. Institutional review board (IRB) approvals or equivalent for research with human**
685 **subjects**

686 Question: Does the paper describe potential risks incurred by study participants, whether
687 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)
688 approvals (or an equivalent approval/review based on the requirements of your country or
689 institution) were obtained?

690 Answer: [N/A].

691 Justification: The paper does not involve human-subjects research.

692 **16. Declaration of LLM usage**

693 Question: Does the paper describe the usage of LLMs if it is an important, original, or
694 non-standard component of the core methods in this research? Note that if the LLM is used
695 only for writing, editing, or formatting purposes and does *not* impact the core methodology,
696 scientific rigor, or originality of the research, declaration is not required.

697 Answer: [N/A].

698 Justification: The core method, theory, and experiments do not use LLMs as an important,
699 original, or non-standard research component.