
Polynomial Links as Parameter-Tied Rank Lifts for Tucker Tensors

Anonymous Author(s)

Affiliation

Address

email

Abstract

Tucker decomposition is a compact model for multiway data, but its map from a latent tensor to measurements is linear. Nonlinear measurement responses can therefore make a low-rank latent tensor appear higher-rank in the measured tensor. We propose Nonlinear Tucker Decomposition with Polynomial Links (NTD-PL), a Tucker enhancement that learns a polynomial link jointly with a Tucker backbone. NTD-PL keeps the backbone ranks fixed and adds only scalar link coefficients. Its polynomial powers induce a parameter-tied rank lift: each degree forms Veronese-lifted Tucker factors tied to the same latent tensor, rather than an untied high-rank correction.

We analyze this structure through Veronese rank expansion, separation from matched-rank Tucker, and approximation results for nonlinear responses. We optimize a masked reconstruction and completion objective by alternating Link Refresh, a small ridge update for the scalar link, with Linked-Backbone Descent and degree continuation. Experiments on controlled tensors, hyperspectral reconstruction and completion, rank-lift and Link Refresh diagnostics, and external cross-domain cases show scoped gains: NTD-PL helps when the residual left by matched-rank Tucker is coherent with a nonlinear response.

1 Introduction

Tensor decompositions provide compact models for multiway data in scientific data analysis and signal processing [Kolda and Bader, 2009, Sidiropoulos et al., 2017, Liu et al., 2013]. Tucker decomposition is a standard choice because it separates a small core from mode-wise factors. It uses a linear map from the Tucker latent tensor to the measured tensor: the reconstructed Tucker value is fitted directly to the measurement. This assumption is natural for many data sets, but it can be mismatched to measurement or data-generation pipelines with response curves, saturation, calibration effects, or material-dependent nonlinearities [Debevec and Malik, 1997, Grossberg and Nayar, 2004, Bioucas-Dias et al., 2012, Heylen et al., 2014]. In such settings, the latent signal may be well organized by a low-rank Tucker tensor while the measured tensor appears to have larger multilinear rank after a nonlinear response.

A direct remedy is to increase Tucker rank, but this changes the multilinear budget and expands the core and factors throughout the model. More flexible nonlinear tensor decoders based on kernels, Gaussian processes, or neural networks can also be useful, but they address a broader problem and often make the relation to a Tucker backbone less explicit [Xu et al., 2011, Zhe et al., 2016, Liu et al., 2017, Saragadam et al., 2024]. We ask a narrower question: can a Tucker model keep its chosen backbone ranks fixed while learning a response that explains coherent nonlinear residuals left by matched-rank Tucker?

We instantiate this idea with a polynomial link. Let

$$\mathcal{S} = \llbracket \mathcal{X}; \mathbf{A}^{(1)}, \dots, \mathbf{A}^{(N)} \rrbracket$$

37 be a Tucker latent tensor. Instead of fitting the measured tensor directly by $\mathcal{S}$, we model

$$\hat{\mathcal{Y}} = f_{\beta}(\mathcal{S}) = \sum_{p=0}^P \beta_p \mathcal{S}^{\odot p},$$

38 where $\odot p$ denotes the Hadamard power. We call the model Nonlinear Tucker Decomposition
 39 with Polynomial Links (NTD-PL). It contains Tucker as the degree-one case, keeps the Tucker
 40 backbone and multilinear ranks fixed, and adds only $P + 1$ scalar link coefficients. The higher powers
 41 $\mathcal{S}^{\odot p}$ induce high-order interactions among the same core and factors, so the added expressivity is
 42 parameter-tied to one learned latent tensor.

43 This viewpoint is central to the paper. NTD-PL is best understood as a Tucker enhancement and
 44 shared-backbone rank lift, not as a generic nonlinear decoder. A degree- p power of the latent tensor
 45 behaves like an ordinary Tucker tensor whose factors are polynomial feature lifts of the original
 46 factors, but those lifted features remain tied to one learned backbone. The new degrees therefore
 47 reshape the low-rank geometry through the polynomial link, rather than adding an untied high-rank
 48 tensor block. This is also why the link is not a post-hoc calibration layer.

49 Our claims are intentionally scoped. Increasing Tucker rank remains the right comparison when the
 50 residual is ordinary multilinear structure, and we do not claim to dominate neural, kernel, or other
 51 nonlinear tensor factorization methods. The intended regime is more specific: the residual left by
 52 matched-rank Tucker should be coherent with a nonlinear response.

53 This scoped regime is also supported empirically. In controlled tensors, the gain grows with the
 54 injected nonlinear residual energy and with the degree needed to approximate the response. On
 55 real hyperspectral tensors, NTD-PL improves matched-rank Tucker in both full reconstruction and
 56 held-out completion. More importantly, a one-step Link Refresh diagnostic predicts when NTD-PL
 57 will help: across HSI scenes and non-HSI tensor units, larger diagnostic reduction is associated with
 58 larger NTD-PL gains, while low-yield units show little or no benefit. These results suggest that
 59 nonlinear responses are identifiable structure in several real tensor families.

60 **Contributions.** We make the following contributions.

- 61 • **Model.** We propose NTD-PL, a Tucker enhancement with a polynomial link on a Tucker latent ten-
 62 sor. It keeps the Tucker backbone and multilinear ranks fixed while adding only scalar coefficients
 63 for a parameter-tied response.
- 64 • **Theory.** We characterize the induced model class through Veronese rank expansion, separation
 65 from matched-rank Tucker, and approximation results for nonlinear responses.
- 66 • **Optimization.** We derive a masked objective for reconstruction and completion, and optimize it
 67 with Link Refresh, Linked-Backbone Descent, and degree continuation.
- 68 • **Experiments.** We evaluate controlled nonlinear residuals, hyperspectral reconstruction and com-
 69 pletion, rank-lift and Link Refresh diagnostics, and external cross-domain boundary cases.

70 2 Related Work

71 **Low-rank tensor models and Tucker backbones.** CP, Tucker, tensor train, and tensor ring
 72 decompositions provide compact representations for multiway data through different rank notions
 73 and factor topologies [Kolda and Bader, 2009, Sidiropoulos et al., 2017, De Lathauwer et al., 2000,
 74 Oseledets, 2011, Zhao et al., 2016]. Tucker is the closest baseline for NTD-PL because it uses the
 75 same core-factor multilinear backbone. NTD-PL does not primarily compete by changing the rank
 76 tuple, replacing the factor topology, or moving to another tensor-network format. Instead, it keeps the
 77 Tucker backbone and multilinear ranks fixed and changes the measurement map on top of the Tucker
 78 latent tensor.

79 **Nonlinear tensor factorization and neural/kernel decoders.** Nonlinear tensor models include
 80 generalized likelihood and link formulations, Bayesian and kernelized tensor factorizations, Gaussian-
 81 process tensor models, and neural tensor decoders [Hong et al., 2020, Schein et al., 2016, Xu et al.,
 82 2011, Zhe et al., 2016, Liu et al., 2017, Pan et al., 2020, Saragadam et al., 2024, Varolgunes et al.,
 83 2023]. These methods can be more flexible than a scalar polynomial link, but this flexibility often

comes with heavier inference, additional hyperparameters, or a less transparent finite-dimensional relation to Tucker. NTD-PL is complementary: it deliberately uses a low-dimensional scalar response rather than a general kernel, Gaussian-process, or neural decoder, so that the induced effect can be analyzed relative to a fixed Tucker backbone.

Polynomial feature lifts and high-order interactions. Polynomial and Volterra-style expansions are classical tools for representing nonlinear responses [Volterra, 1959, Boyd and Chua, 1985]. Related matrix and tensor methods use polynomial feature maps, quadratic factorizations, or explicit high-order interactions to enlarge a linear model [Zhai et al., 2024, Liu et al., 2015, Deng et al., 2016]. Such constructions can make nonlinear interactions explicit, but untied high-order parameterizations can increase parameter dimension and contraction cost. NTD-PL does not introduce an independent high-order tensor parameter. Its polynomial powers are powers of one learned Tucker latent tensor, inducing parameter-tied high-order interactions and Veronese-style backbone-tied rank lifts.

Positioning. Existing nonlinear tensor models either use broad nonlinear decoders, modify likelihoods or kernels, or explicitly parameterize high-order interactions. What is less developed is a Tucker-preserving, low-dimensional, analyzable polynomial link that yields a parameter-tied rank lift and supports masked learning. NTD-PL trades general nonlinear flexibility for structure, interpretability, and a direct finite-dimensional rank-lift view. Its expected benefit is therefore scoped to settings where matched-rank Tucker leaves a coherent nonlinear residual explained by a scalar response.

3 Model

3.1 Model and Measurement Assumption

Let $\mathcal{Y} \in \mathbb{R}^{I_1 \times \dots \times I_N}$ be a measured tensor and let

$$\mathcal{S} = \llbracket \mathcal{X}; \mathbf{A}^{(1)}, \dots, \mathbf{A}^{(N)} \rrbracket, \quad \mathcal{X} \in \mathbb{R}^{R_1 \times \dots \times R_N}, \quad \mathbf{A}^{(n)} \in \mathbb{R}^{I_n \times R_n}. \quad (1)$$

Our measurement assumption is that the measured tensor is generated from this low-rank Tucker latent tensor through a scalar response, up to noise. NTD-PL models this response by a polynomial:

$$Y_{\mathbf{i}} = f_{\beta}(S_{\mathbf{i}}) + \varepsilon_{\mathbf{i}}, \quad f_{\beta}(s) = \sum_{p=0}^P \beta_p s^p, \quad \mathbf{i} = (i_1, \dots, i_N). \quad (2)$$

Equivalently, the tensor prediction is

$$\hat{\mathcal{Y}} = f_{\beta}(\mathcal{S}) = \sum_{p=0}^P \beta_p \mathcal{S}^{\odot p}, \quad \beta \in \mathbb{R}^{P+1}. \quad (3)$$

The constant term β_0 captures affine response shifts, the linear term $\beta_1 \mathcal{S}$ is the usual Tucker channel, and terms with $p \geq 2$ activate nonlinear channels on the same latent tensor. Standard Tucker is recovered by setting $\beta_1 = 1$ and $\beta_p = 0$ for $p \neq 1$. Thus NTD-PL changes the measurement map while leaving the Tucker backbone and its multilinear ranks fixed.

3.2 Backbone-Tied High-Order Interactions

Although f_{β} is scalar, the term $\mathcal{S}^{\odot p}$ is not merely an independent output calibration. We arrive at it by first defining an explicit p -fold Tucker extension: instead of mapping one latent index per mode, the mode- n projection is a higher-order tensor $\mathcal{A}^{(n,p)} \in \mathbb{R}^{I_n \times R_n^p}$ that couples p latent indices. The case $p = 1$ is ordinary Tucker, while $p = 2$ gives a quadratic Tucker-style interaction between two latent Tucker tuples. Appendix A.1 gives the formal construction and tensor-network view.

The explicit construction is a reference model, not the parameterization we train. Its projection tensors grow as $I_n R_n^p$, and each degree would add a new set of mode projections. NTD-PL keeps the same high-order interaction pattern by tying the p -fold projection to the original Tucker factor:

$$\mathcal{A}_{i_n, r_n^{(1)}, \dots, r_n^{(p)}}^{(n,p)} = \prod_{q=1}^p A_{i_n, r_n^{(q)}}^{(n)}.$$

Under this tying, the degree- p channel is exactly S_1^p . Thus it still couples p latent Tucker index tuples, but all copies share the same mode factors. Across degrees, the only new free parameters are the scalar coefficients $\{\beta_p\}$, which mix these tied channels.

The latent tensor $\mathcal{S}$ and the coefficients β are fitted jointly, so the low-rank backbone is optimized together with the polynomial channels. Section 4 analyzes the induced rank expansion and the resulting separation from standard Tucker.

3.3 Masked Least-Squares Objective

Let $\Omega \in \{0, 1\}^{I_1 \times \dots \times I_N}$ denote the observed-entry mask and let $|\Omega| = \sum_i \Omega_i$. We estimate the core, factors, and link coefficients by the regularized masked objective

$$\min_{\mathcal{X}, \{\mathbf{A}^{(n)}\}, \beta} \frac{1}{2|\Omega|} \|\Omega \odot (\mathcal{Y} - f_\beta(\mathcal{S}))\|_F^2 + \frac{\lambda_X}{2} \|\mathcal{X}\|_F^2 + \sum_{n=1}^N \frac{\lambda_n}{2} \|\mathbf{A}^{(n)}\|_F^2 + \frac{\lambda_\beta}{2} \|\beta\|_2^2. \quad (4)$$

The same formulation covers full reconstruction by taking $\Omega \equiv \mathbf{1}$ and tensor completion by restricting the loss to the observed entries. The quadratic penalties regularize the Tucker backbone and the low-dimensional polynomial link.

4 Theory

The theory addresses the rank-lift view raised by the model definition. A scalar link has few free parameters, but its powers can induce ordinary Tucker tensors with larger multilinear ranks. We first make this induced structure explicit, then give a separation from matched-rank Tucker and an approximation route for nonlinear responses.

4.1 Rank Expansion and Separation

Proposition 1 (Shared-backbone rank expansion). *Let $\mathcal{S} = \llbracket \mathcal{X}; \mathbf{A}^{(1)}, \dots, \mathbf{A}^{(N)} \rrbracket$ have Tucker rank at most $(R_1, \dots, R_N)$. For every integer $p \geq 1$, $\mathcal{S}^{\odot p}$ admits a Tucker representation with mode ranks at most*

$$\left(\binom{R_1 + p - 1}{p}, \dots, \binom{R_N + p - 1}{p} \right). \quad (5)$$

More explicitly, the ordinary Tucker factors of $\mathcal{S}^{\odot p}$ can be chosen as Veronese lifts of the original factors. For $\mathbf{A} \in \mathbb{R}^{I \times R}$, define $\mathbf{A}^{(p)} \in \mathbb{R}^{I \times D}$, $D = \binom{R+p-1}{p}$, by

$$A_{i,\alpha}^{(p)} = \prod_{r=1}^R A_{i,r}^{\alpha_r}, \quad (6)$$

where $\alpha \in \mathbb{N}^R$ ranges over $|\alpha| = p$. Then $\mathcal{S}^{\odot p}$ factors through $\mathbf{A}^{(1)(p)}, \dots, \mathbf{A}^{(N)(p)}$.

This proposition makes the tying in Section 3.2 explicit in standard Tucker language: a degree- p channel is an ordinary Tucker tensor whose mode factors are polynomial feature lifts of the same backbone factors. The link therefore changes rank structure through tied features, not through an independent high-order parameter block.

Theorem 1 (Tensorial separation from standard Tucker). *Let $N \geq 3$, $R \geq 2$, $p \geq 2$, and suppose $I_n \geq D$ for every n , where $D = \binom{R+p-1}{p}$. Then there exists an order- N tensor $\mathcal{Y} \in \mathbb{R}^{I_1 \times \dots \times I_N}$ that is representable by a degree- p NTD-PL model with Tucker backbone rank $(R, \dots, R)$, but whose exact standard Tucker rank is*

$$(D, \dots, D). \quad (7)$$

Equivalently, a degree- p polynomial link can generate tensors whose multilinear ranks are $((\binom{R+p-1}{p}), \dots, (\binom{R+p-1}{p}))$ from the same Tucker backbone of rank $(R, \dots, R)$.

The proof in Appendix B uses a CP-diagonal Tucker backbone and generic full-rank Veronese feature matrices. The theorem is an existence separation, not a universality claim over all high-rank Tucker tensors: it identifies a structured high-rank subset generated by a low-rank backbone.

158 4.2 Approximation Under Nonlinear Measurement Responses

159 **Theorem 2** (Polynomial generation). Suppose $\mathcal{Y} = g(\mathcal{S}^*)$, where $\mathcal{S}^*$ has Tucker rank at most
 160 $(R_1, \dots, R_N)$ and g is a polynomial of degree at most P . Then degree- P NTD-PL with the same
 161 Tucker rank can represent $\mathcal{Y}$ exactly.

162 **Theorem 3** (Analytic link approximation). Suppose $\mathcal{Y} = g(\mathcal{S}^*)$, where $\mathcal{S}^*$ has Tucker rank at most
 163 $(R_1, \dots, R_N)$ and $\|\mathcal{S}^*\|_\infty \leq B$. If g is analytic on the complex disk $D(0, \rho)$ with $B < \rho$, and
 164 $M_\rho = \sup_{|z|=\rho} |g(z)|$, then there exists a degree- P NTD-PL model with the same Tucker backbone
 165 satisfying

$$\frac{1}{\sqrt{M}} \|\mathcal{Y} - \hat{\mathcal{Y}}\|_F \leq M_\rho \frac{(B/\rho)^{P+1}}{1 - B/\rho}, \quad M = \prod_{n=1}^N I_n. \quad (8)$$

166 Thus the same finite link that induces a tied rank lift also gives a direct approximation route for
 167 smooth responses on a bounded latent range. The statement uses the simple Taylor truncation bound;
 168 sharper polynomial approximation bounds can replace it under stronger assumptions.

169 Together, the results clarify the scope of NTD-PL. Polynomial links can approximate nonlinear
 170 responses while keeping the same latent Tucker rank, and their powers induce Veronese-lifted Tucker
 171 factors. This explains how a scalar link can change multilinear rank structure without introducing
 172 untied high-order tensor parameters.

173 5 Optimization

174 We optimize the masked objective in Section 3.3 by alternating between the scalar link and the Tucker
 175 backbone. The nonlinear part is low-dimensional once the latent Tucker tensor is fixed, while the
 176 backbone update remains a standard differentiable Tucker fitting step through the current link. We
 177 call the two stages *Link Refresh* and *Linked-Backbone Descent*.

178 **Link Refresh.** Fix the current latent tensor

$$S = \llbracket \mathcal{X}; \mathbf{A}^{(1)}, \dots, \mathbf{A}^{(N)} \rrbracket. \quad (9)$$

179 For active degree $Q \leq P$, updating the polynomial coefficients reduces to a small masked ridge
 180 regression,

$$\min_{\beta_{0:Q}} \frac{1}{2|\Omega|} \|y_\Omega - \Phi_{\Omega,Q}(S)\beta_{0:Q}\|_2^2 + \frac{\lambda_\beta}{2} \|\beta_{0:Q}\|_2^2, \quad (10)$$

181 where $\Phi_{\Omega,Q}(S)$ contains the observed powers $[1, s, \dots, s^Q]$. This step is cheap because it solves
 182 only a $(Q+1) \times (Q+1)$ linear system. It also gives an interpretable view of the link: once the
 183 backbone exposes scalar latent values over the tensor, the refresh estimates the best scalar response
 184 on the observed set.

185 **Linked-Backbone Descent.** With β fixed, the Tucker parameters are updated through the current
 186 nonlinear measurement map. The masked residual is first propagated to the latent tensor as

$$T = \frac{1}{|\Omega|} \Omega \odot (f_\beta(S) - Y) \odot f'_\beta(S), \quad (11)$$

187 and then backpropagated through the Tucker reconstruction to update $\mathcal{X}$ and $\{\mathbf{A}^{(n)}\}$. In implemen-
 188 tation we use Adam for these first-order steps, but the update itself is just standard Tucker reverse
 189 contractions composed with the current link. The backbone is therefore not calibrated after fitting; it
 190 is jointly reshaped under the nonlinear link.

191 **Stabilization.** Directly activating all polynomial degrees from the start is brittle, because higher
 192 powers of S amplify scale ambiguity. We therefore begin from the Tucker-compatible degree-one
 193 model and activate higher degrees gradually. When a new degree is switched on, its coefficient is
 194 initialized at zero, so the model expands from the current lower-degree solution without changing
 195 the prediction abruptly. We also normalize Tucker factor scales and absorb the corresponding
 196 scales into the core, which fixes a Tucker gauge without changing S . For the coefficient solve, the

implementation may regress on a centered and scaled latent variable and then map the solution back to the original power basis. Appendix Figure 6 shows representative training curves where activating the relevant higher degrees produces systematic loss drops. Full gradients, stabilization details, costs analysis and pseudocode are given in Appendix A.2.

Masked uniform convergence. The masked objective also admits a standard sampling interpretation. Let $m = |\Omega|$ and let $\Theta \subset \mathbb{R}^d$ be a bounded NTD-PL parameter set, where $d = \prod_{n=1}^N R_n + \sum_{n=1}^N I_n R_n + (P + 1)$. If the per-sample losses are bounded by $0 \leq \ell_i(\theta) \leq B_\ell$ and G -Lipschitz in θ , then with probability at least $1 - \delta$,

$$\sup_{\theta \in \Theta} |\mathcal{L}(\theta) - \hat{\mathcal{L}}_m(\theta)| \leq 2G\varepsilon + B_\ell \sqrt{\frac{d \log(3B_\Theta/\varepsilon) + \log(2/\delta)}{2m}}, \quad (12)$$

for every $0 < \varepsilon \leq B_\Theta$. This is a standard uniform convergence bound, not a sharp minimax result, but it clarifies that masked training controls the full-tensor risk with a complexity term determined by the tied NTD-PL parameter count. The proof is given in Appendix B.4.

6 Experiments

The experiments test three mechanism questions: whether NTD-PL behaves as a parameter-tied rank lift under controlled nonlinear residuals; whether the polynomial link improves matched-rank Tucker on real reconstruction and held-out completion; and whether D_{link} diagnoses beneficial tensor units across domains. Unless stated otherwise, real tensors are globally max-normalized and used without spatial downsampling or cropping. We report RMSE and SAM for reconstruction, and RMSE* and SAM* only on held-out entries for completion. NTD-PL is compared mainly against same-rank Tucker, except in the rank-lift ablation.

6.1 Controlled Nonlinear Residuals

These synthetic experiments are mechanism checks rather than generic benchmarks. We generate measured tensors as

$$Y_\alpha = \sqrt{1 - \alpha} S^* + \sqrt{\alpha} \tilde{R}^*,$$

where S^* is a low-rank Tucker tensor and $\tilde{R}^*$ is the normalized component of a scalar nonlinear response orthogonal to S^* . Thus α controls nonlinear residual energy, not just the raw amplitude of a nonlinearity. Figure 1 shows the expected transition: NTD-PL is Tucker-like when the residual is nearly linear and separates as nonlinear residual energy increases. Tucker can lag behind CP, but NTD-PL then pulls ahead systematically. Figure 2 shows the complementary degree effect: gains emerge when the active polynomial degree can match or approximate the response complexity. Degree continuation and coefficient-contribution details are deferred to Appendix C.1, where Table 4 shows the linear consistency negative control. Taken together, these checks indicate that NTD-PL is not a blanket improvement rule; its gain depends on the data’s nonlinear structure.




Figure 1: Controlled nonlinear residual sweep. NTD-PL stays Tucker-like near the linear regime and separates as α , the nonlinear residual energy, increases. Lower RMSE is better.




Figure 2: Degree-dependent gains at $\alpha = 0.25$. Cell values are Tucker minus NTD-PL RMSE; gains appear when the active degree is sufficient for the response complexity.

6.2 Tucker Enhancement on Real Hyperspectral Tensors

This section tests whether the same fixed-rank enhancement appears on real hyperspectral tensors, both when the tensor is fully measured and when evaluation is performed on held-out missing entries. We use CAVE dataset because it offers multiple real scenes with rich spectral and spatial structure, making it a reasonable mechanism check rather than a toy-sized tensor.

Representation enhancement. Full reconstruction is a representation-capacity check, not a generalization test. Figure 3 shows that NTD-PL improves RMSE and SAM on CAVE at matched Tucker backbone rank and similar parameter budget. The gap shrinks as Tucker rank increases, consistent with a fixed-rank enhancement: ordinary multilinear capacity already explains more of the tensor. Table 1 gives the complementary rank-lift ablation. Tucker can match the NTD-PL RMSE only after spending substantially more ordinary rank and parameters, supporting the parameter-tied rank-lift interpretation. Even when RMSE is close, NTD-PL usually retains a clearer SAM advantage.




Figure 3: CAVE full reconstruction as a representation-capacity check. At matched Tucker backbone rank, NTD-PL reduces RMSE and SAM; annotations show paired relative gains over Tucker at each rank. The gap shrinks as Tucker rank increases, consistent with fixed-rank enhancement. Each point averages 15 scenes.

Table 1: Rank-lifted Tucker comparison on CAVE reconstruction. For each NTD-PL backbone rank, we increase Tucker’s spatial ranks while keeping the spectral rank fixed and report the first Tucker rank that matches NTD-PL RMSE, along with the extra Tucker parameters required.

NTD-PL reference				First Tucker matching RMSE				Extra params
Rank	Params	RMSE↓	SAM↓	Rank	Params	RMSE↓	SAM↓	
(18, 18, 3)	20k	0.0319	19.23	(28, 28, 3)	31k	0.0316	21.41	+59.5%
(24, 24, 4)	27k	0.0256	15.15	(35, 35, 4)	41k	0.0254	16.26	+51.3%
(33, 33, 4)	38k	0.0223	14.41	(49, 49, 4)	60k	0.0222	15.28	+56.5%

Held-out enhancement. Held-out completion is the stronger evidence because it tests whether the learned geometry transfers to unseen entries. Table 2 reports CAVE completion under random masks, with RMSE* and SAM* computed only on missing entries. At the same rank budget, NTD-PL gives consistent held-out RMSE* gains and usually improves SAM* as well. The paired sign and Wilcoxon tests in Appendix C.2 show the same scene-level pattern.

Table 2: Held-out CAVE completion at rank (33, 33, 4) under random masks. RMSE* and SAM* are computed only on missing entries. Columns report mean $\pm$ standard deviation over scenes; gains are paired relative improvements over matched-rank Tucker.

Missing rate	RMSE* $\downarrow$			SAM* $\downarrow$		
	Tucker	NTD-PL	Gain	Tucker	NTD-PL	Gain
0.1	0.0263 $\pm$ 0.0179	0.0229 $\pm$ 0.0158	12.9%	11.82 $\pm$ 5.33	8.36 $\pm$ 3.38	29.3%
0.3	0.0263 $\pm$ 0.0179	0.0230 $\pm$ 0.0158	12.5%	14.75 $\pm$ 6.33	12.47 $\pm$ 5.06	15.5%
0.5	0.0264 $\pm$ 0.0179	0.0230 $\pm$ 0.0158	12.9%	15.56 $\pm$ 6.85	13.48 $\pm$ 5.75	13.4%
0.7	0.0266 $\pm$ 0.0180	0.0219 $\pm$ 0.0128	17.7%	16.02 $\pm$ 7.21	13.79 $\pm$ 5.43	13.9%

6.3 Diagnosing Nonlinear Low-Rank Structure Across Data

NTD-PL is meant to be diagnosable rather than uniformly beneficial. Its gain should appear when matched-rank Tucker leaves a coherent residual that a scalar link can reduce. We measure this with D_{link} , a cheap frozen-backbone Link Refresh probe: after fitting Tucker, freeze the backbone, run one scalar Link Refresh ridge solve, and measure residual-energy reduction on the relevant evaluation samples. Larger D_{link} indicates more link-like residual structure:

$$D_{\text{link}} = 10 \log_{10} \frac{\|Y_{\text{eval}} - \hat{Y}_T\|_F^2}{\|Y_{\text{eval}} - \hat{Y}_{\text{link}}\|_F^2},$$

where $\hat{Y}_T$ is the Tucker output and $\hat{Y}_{\text{link}}$ is the output after the one-step refresh.

For the CAVE completion, Table 3 groups scenes by D_{link} tercile. Higher D_{link} consistently corresponds to larger held-out RMSE* gains across missing rates, with Spearman correlations around 0.829–0.889. Appendix Figure 7 shows the corresponding scene-level scatter plots.

Table 3: CAVE completion diagnostic. Scenes are grouped by D_{link} terciles, and entries report mean NTD-PL RMSE* gains over matched-rank Tucker. Larger D_{link} consistently corresponds to larger held-out gains across missing rates.

Missing rate	Low D_{link}	Mid D_{link}	High D_{link}	Spearman ρ_s
0.1	0.75%	7.92%	12.63%	0.829
0.3	-1.13%	9.10%	13.12%	0.871
0.5	-1.09%	9.16%	13.18%	0.889
0.7	-0.62%	9.22%	13.07%	0.868

We then apply the same per-unit protocol beyond HSI. Each natural image is one tensor unit; each object or instance forms an object-view tensor; each KTH segment or UCF101 clip is an action-video tensor. Figure 4 plots D_{link} against NTD-PL RMSE gain across these independently fitted units. Diagnostic labels and terciles are assigned from D_{link} , not from NTD-PL gain, and the relationship is domain-dependent but positive: natural images show weaker link yield, object-view tensors clearer yield, and action-video clips an intermediate pattern.

These results support the scoped claim that NTD-PL helps when matched-rank Tucker leaves link-like residual structure, not uniformly across all visual tensors. Appendix Figure 7 gives the CAVE scene-level scatter supplement, and Appendix Table 7 gives the exact non-HSI numeric statistics.




Figure 4: Per-unit non-HSI diagnostic correlation. Each panel plots D_{link} against NTD-PL RMSE gain over independently fitted tensor units. D_{link} is computed by frozen-backbone Link Refresh, and diagnostic labels are assigned from D_{link} , not from NTD-PL gain. The figure shows domain-dependent but positive diagnostic structure.

7 Limitations

NTD-PL assumes that a common nonlinear response explains a coherent part of the residual left by matched-rank Tucker. This scalar link may be less effective when nonlinearities vary strongly across spatial regions, spectral bands, modes, or samples. In such cases, a single polynomial response can underfit the domain variation even if the Tucker backbone is appropriate.

NTD-PL enhances a fixed Tucker rank budget by changing the measurement map on top of the latent tensor. When the remaining error is ordinary multilinear structure rather than a scalar response, increasing Tucker rank remains the more direct modeling choice.

Polynomial degree introduces a second tradeoff. Higher degrees increase the approximation capacity of the link, but they can amplify Tucker scale ambiguity and numerical conditioning issues in the power basis. This motivates the ridge-regularized Link Refresh step, factor normalization, and degree continuation used in the optimizer.

Finally, the joint problem remains nonconvex. The Link Refresh subproblem is a small ridge regression once the latent tensor is fixed, but the Tucker backbone and polynomial link are learned jointly through a nonconvex objective. We do not claim global optimality.

Our strongest empirical evidence is currently from controlled nonlinear residuals and hyperspectral reconstruction and completion. Broader validation on large sparse tensors, temporal or multimodal tensors, and domain-varying nonlinear responses is left for future work.

8 Conclusion

This paper develops polynomial links as parameter-tied rank lifts for Tucker tensors. The model keeps a Tucker latent tensor as the backbone and learns a polynomial link, adding only scalar coefficients. The resulting powers of the same latent tensor act as shared-backbone high-order interactions rather than untied high-rank corrections.

The analysis makes this rank-lift view explicit through Veronese rank expansion, separation from matched-rank Tucker, and approximation results for nonlinear links. The optimization uses a masked objective with Link Refresh, Linked-Backbone Descent, and degree continuation, while the

experiments and diagnostics identify the intended regime: the link helps when matched-rank Tucker leaves a coherent nonlinear residual explained by a scalar response.

Future work may consider spatially or spectrally adaptive links, scalable sparse implementations, and richer structured links beyond scalar polynomials.

References

- José M. Bioucas-Dias, Antonio Plaza, Nicolas Dobigeon, Mario Parente, Qian Du, Paul Gader, and Jocelyn Chanussot. Hyperspectral unmixing overview: Geometrical, statistical, and sparse regression-based approaches. *IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing*, 5(2):354–379, 2012. doi: 10.1109/JSTARS.2012.2194696.
- Stephen Boyd and Leon O. Chua. Fading memory and the problem of approximating nonlinear operators with volterra series. *IEEE Transactions on Circuits and Systems*, 32(11):1150–1161, 1985. doi: 10.1109/TCS.1985.1085649.
- Lieven De Lathauwer, Bart De Moor, and Joos Vandewalle. A multilinear singular value decomposition. *SIAM Journal on Matrix Analysis and Applications*, 21(4):1253–1278, 2000. doi: 10.1137/S0895479896305696.
- Paul E. Debevec and Jitendra Malik. Recovering high dynamic range radiance maps from photographs. In *Proceedings of the 24th Annual Conference on Computer Graphics and Interactive Techniques, SIGGRAPH ’97*, pages 369–378, 1997. doi: 10.1145/258734.258884.
- Jian Deng, Haotian Liu, Kim Batselier, Yu-Kwong Kwok, and Ngai Wong. Storm: a nonlinear model order reduction method via symmetric tensor decomposition. In *2016 21st Asia and South Pacific design automation conference (ASP-DAC)*, pages 557–562. IEEE, 2016.
- Michael D. Grossberg and Shree K. Nayar. Modeling the space of camera response functions. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 26(10):1272–1282, 2004. doi: 10.1109/TPAMI.2004.88.
- Rob Heylen, Mario Parente, and Paul Gader. A review of nonlinear hyperspectral unmixing methods. *IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing*, 7(6):1844–1868, 2014. doi: 10.1109/JSTARS.2014.2320576.
- David Hong, Tamara G Kolda, and Jed A Duersch. Generalized canonical polyadic tensor decomposition. *SIAM review*, 62(1):133–163, 2020.
- Tamara G. Kolda and Brett W. Bader. Tensor decompositions and applications. *SIAM Review*, 51(3): 455–500, 2009. doi: 10.1137/07070111X.
- Bin Liu, Lirong He, Shandian Zhe, Yingmin Li, and Zenglin Xu. Deepcp: Flexible nonlinear tensor decomposition. In *NeurIPS Workshops*, volume 3, 2017.
- Haotian Liu, Luca Daniel, and Ngai Wong. Model reduction and simulation of nonlinear circuits via tensor decomposition. *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, 34(7):1059–1069, 2015.
- Ji Liu, Przemyslaw Musialski, Peter Wonka, and Jieping Ye. Tensor completion for estimating missing values in visual data. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 35(1):208–220, 2013. doi: 10.1109/TPAMI.2012.39.
- Ivan V. Oseledets. Tensor-train decomposition. *SIAM Journal on Scientific Computing*, 33(5): 2295–2317, 2011.
- Zhimeng Pan, Zheng Wang, and Shandian Zhe. Streaming nonlinear bayesian tensor decomposition. In *Conference on Uncertainty in Artificial Intelligence*, pages 490–499. PMLR, 2020.
- Vishwanath Saragadam, Randall Balestrieri, Ashok Veeraraghavan, and Richard G Baraniuk. Deeptensor: Low-rank tensor decomposition with deep network priors. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 46(12):10337–10348, 2024.

- 336 Aaron Schein, Mingyuan Zhou, David Blei, and Hanna Wallach. Bayesian poisson tucker decompo-
 337 sition for learning the structure of international relations. In *International Conference on Machine*
 338 *Learning*, pages 2810–2819. PMLR, 2016.
- 339 Nicholas D. Sidiropoulos, Lieven De Lathauwer, Xiao Fu, Kejun Huang, Evangelos E. Papalexakis,
 340 and Christos Faloutsos. Tensor decomposition for signal processing and machine learning. *IEEE*
 341 *Transactions on Signal Processing*, 65(13):3551–3582, 2017. doi: 10.1109/TSP.2017.2690524.
- 342 Uras Varolgunes, Dan Zhou, Dantong Yu, and Ajim Uddin. Nmtucker: Non-linear matryoshka
 343 tucker decomposition for financial time series imputation. In *Proceedings of the Fourth ACM*
 344 *International Conference on AI in Finance*, pages 516–523, 2023.
- 345 Vito Volterra. *Theory of Functionals and of Integral and Integro-Differential Equations*. Dover
 346 Publications, 1959.
- 347 Zenglin Xu, Feng Yan, et al. Infinite tucker decomposition: Nonparametric bayesian models for
 348 multiway data analysis. *arXiv preprint arXiv:1108.6296*, 2011.
- 349 Zheng Zhai, Hengchao Chen, and Qiang Sun. Quadratic matrix factorization with applications
 350 to manifold learning. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 46(9):
 351 6384–6401, 2024.
- 352 Qibin Zhao, Guoxu Zhou, Shengli Xie, Liqing Zhang, and Andrzej Cichocki. Tensor ring decomposi-
 353 tion. *arXiv preprint arXiv:1606.05535*, 2016.
- 354 Shandian Zhe, Kai Zhang, Pengyuan Wang, Kuang-chih Lee, Zenglin Xu, Yuan Qi, and Zoubin
 355 Ghahramani. Distributed flexible nonlinear tensor factorization. *Advances in neural information*
 356 *processing systems*, 29, 2016.

357 A Additional Model Details

358 A.1 Explicit p -Fold Tucker Motivation

359 This section expands the motivation behind the explicit p -fold Tucker channel in Section 3.2. The
 360 construction is best viewed as the tensor analogue of adding polynomial interactions to a linear map.
 361 In the matrix case, one may move from a linear model $\mathbf{y} = \mathbf{A}\mathbf{x}$ to a quadratic model

$$\mathbf{y} = \mathbf{A}\mathbf{x} + \mathcal{A}(\mathbf{x}, \mathbf{x}),$$

362 where the second term couples pairs of latent coordinates. The corresponding question for Tucker is:
 363 if standard Tucker is the multilinear low-rank map for multiway data, what is the direct p -th order
 364 interaction analogue?

365 For a Tucker backbone, an explicit p -fold interaction can be written as

$$\tilde{Y}_i^{(p)} = \sum_{\mathbf{r}^{(1)}, \dots, \mathbf{r}^{(p)}} \left(\prod_{q=1}^p \mathcal{X}_{r_1^{(q)}, \dots, r_N^{(q)}} \right) \prod_{n=1}^N \mathcal{A}_{i_n, r_n^{(1)}, \dots, r_n^{(p)}}^{(n,p)}. \quad (13)$$

366 Here each mode projection depends on p latent indices instead of one. For $p = 1$, this reduces to the
 367 ordinary Tucker entry formula. For $p = 2$, the same core appears twice and the mode projections
 368 couple two latent tuples, giving a quadratic Tucker-style channel. Figure 5 illustrates this progression.

369 The explicit construction is useful as a reference model, but it is not the parameterization used by
 370 NTD-PL. Its mode projections $\mathcal{A}^{(n,p)}$ have size $I_n R_n^p$, so a direct implementation introduces new
 371 high-order projection tensors and rapidly increases storage and contraction cost. NTD-PL keeps the
 372 same high-order interaction pattern by tying the projection tensor to the original Tucker factor:

$$\mathcal{A}_{i_n, r_n^{(1)}, \dots, r_n^{(p)}}^{(n,p)} = \prod_{q=1}^p \mathcal{A}_{i_n, r_n^{(q)}}^{(n)}. \quad (14)$$

373 Substituting (14) into (13) gives exactly $\mathcal{S}^{\odot p}$, where $\mathcal{S} = \llbracket \mathcal{X}; \mathbf{A}^{(1)}, \dots, \mathbf{A}^{(N)} \rrbracket$. Thus the degree- p
 374 NTD-PL channel is a shared-projection version of the explicit p -fold Tucker prototype.

375 This is the modeling role of the p -fold Tucker prototype: it defines the high-order Tucker interaction
 376 that NTD-PL compresses. The final model does not train the untied tensors $\mathcal{A}^{(n,p)}$; it keeps a single
 377 Tucker backbone and adds only the polynomial coefficients $\{\beta_p\}$.




Figure 5: Tensor-network view of a formal p -fold Tucker extension. The construction keeps the Tucker idea of mode-wise generation, but lets each mode projection couple multiple latent index tuples.

378 A.2 Optimization Details

379 This section gives the update equations behind the solver summarized in Section 5. The implementa-
 380 tion follows the same structure: one Tucker reconstruction per iteration, a fused evaluation of the
 381 polynomial and its derivative, first-order updates of the backbone, and periodic ridge-regression
 382 updates of the link coefficients.

383 **Backbone gradients.** For a fixed β , define

$$f'_\beta(\mathcal{S}) = \sum_{p=1}^P p\beta_p \mathcal{S}^{\odot(p-1)}. \quad (15)$$

384 The gradient of the normalized masked data term with respect to the latent tensor $\mathcal{S}$ is

$$\mathbf{T} = \frac{1}{|\Omega|} \Omega \odot (f_\beta(\mathcal{S}) - \mathcal{Y}) \odot f'_\beta(\mathcal{S}). \quad (16)$$

385 Backpropagating $\mathbf{T}$ through the Tucker reconstruction gives the core gradient

$$\nabla_{\mathcal{X}} \mathcal{L} = \mathbf{T} \times_1 \mathbf{A}^{(1)\top} \times_2 \cdots \times_N \mathbf{A}^{(N)\top} + \lambda_{\mathcal{X}} \mathcal{X}. \quad (17)$$

386 For the n -th factor, let

$$\mathbf{Z}^{(n)} = \left[\mathcal{X} \times_1 \mathbf{A}^{(1)} \cdots \times_{n-1} \mathbf{A}^{(n-1)} \times_{n+1} \mathbf{A}^{(n+1)} \cdots \times_N \mathbf{A}^{(N)} \right]_{(n)}. \quad (18)$$

387 Then

$$\nabla_{\mathbf{A}^{(n)}} \mathcal{L} = \mathbf{T}_{(n)} \mathbf{Z}^{(n)\top} + \lambda_n \mathbf{A}^{(n)}. \quad (19)$$

388 The experiments use Adam for these first-order steps. This choice is not part of the model definition;
 389 any comparable first-order optimizer can use the same gradients.

390 **Link update.** With $\mathcal{S}$ fixed and active degree $Q = P_{\text{act}}$, the coefficient subproblem is the ridge
 391 regression in (10). Its closed-form solution is

$$\beta_{0:Q}^* = (\Phi_{\Omega,Q}^\top \Phi_{\Omega,Q} + |\Omega| \lambda_\beta \mathbf{I})^{-1} \Phi_{\Omega,Q}^\top \mathbf{y}_\Omega. \quad (20)$$

392 Equivalently, one may absorb $|\Omega|$ into the ridge parameter and solve the usual unnormalized normal
 393 equations. If the constant term is disabled, the first column is removed and β_0 is fixed at zero. In
 394 the code, this small system is solved either from an explicit Vandermonde design or from moment
 395 equations built from cached powers.

396 **Stabilization and continuation.** High powers of $\mathcal{S}$ can be sensitive to scale. The implementation
 397 therefore normalizes factor columns after backbone updates and absorbs the corresponding scales
 398 into the core; this is a gauge fixing of the Tucker parameterization and does not change $\mathcal{S}$. For the
 399 coefficient solve, it may regress on the centered and scaled latent variable $(S - \mu)/\sigma$, then convert
 400 the solution back to the original power basis of $\mathcal{S}$.

401 The active degree follows the schedule

$$t_k = \text{round} \left(\frac{kT}{P} \right), \quad k = 1, \dots, P-1.$$

402 When degree $p+1$ is activated, the previous coefficients and Tucker factors are kept and the new
403 coefficient β_{p+1} is initialized at zero.

404 **Cost.** The dominant cost is forming the Tucker reconstruction and its reverse contractions on the
405 observed tensor. For dense tensors this is comparable to standard Tucker optimization at the same
406 rank, plus $O(P|\Omega|)$ work to evaluate the scalar polynomial and its derivative. The link update forms
407 a Vandermonde design or the corresponding moment equations and solves a $(P+1) \times (P+1)$ linear
408 system, with direct cost $O(P^2|\Omega| + P^3)$. Since P is small, this cost is negligible compared with the
409 Tucker contractions.

410 Algorithm 1 gives the full procedural summary.

Algorithm 1 Masked NTD-PL optimization

Require: observed tensor $\mathcal{Y}$, mask Ω , rank $(R_1, \dots, R_N)$, degree P

```

1: initialize  $\mathcal{X}, \{\mathbf{A}^{(n)}\}$  from Tucker on the masked tensor
2: initialize active degree  $P_{\text{act}} \leftarrow 1$  and  $\beta = (0, 1)$ 
3: for  $t = 1, \dots, T$  do
4:   form  $\mathcal{S} = \llbracket \mathcal{X}; \mathbf{A}^{(1)}, \dots, \mathbf{A}^{(N)} \rrbracket$ 
5:   if a link update is scheduled then
6:     update active coefficients  $\beta_{0:P_{\text{act}}}$  by ridge regression on  $\{(S_i, Y_i) : \Omega_i = 1\}$ 
7:   end if
8:   compute  $f_\beta(\mathcal{S})$  and  $f'_\beta(\mathcal{S})$ 
9:   take first-order steps on  $\mathcal{X}$  and  $\{\mathbf{A}^{(n)}\}$  using the masked gradient
10:  normalize factor scales and absorb them into the core
11:  increase  $P_{\text{act}}$  according to the continuation schedule, initializing the new coefficient at zero
12: end for
13: return  $\mathcal{X}, \{\mathbf{A}^{(n)}\}, \beta$ 

```

411 B Additional Proofs

412 B.1 Polynomial and Analytic Approximation

413 *Proof of Theorem 2.* Since g has degree at most P , write

$$g(s) = \sum_{p=0}^P c_p s^p. \quad (21)$$

414 Take the NTD-PL latent tensor to be the true latent tensor, $\mathcal{S} = \mathcal{S}^*$, and set $\beta_p = c_p$. Then, entrywise,

$$\hat{\mathcal{Y}}_{i_1, \dots, i_N} = \sum_{p=0}^P \beta_p (\mathcal{S}_{i_1, \dots, i_N}^*)^p = g(\mathcal{S}_{i_1, \dots, i_N}^*) = \mathcal{Y}_{i_1, \dots, i_N}. \quad (22)$$

415 Thus $\hat{\mathcal{Y}} = \mathcal{Y}$. □

416 **Lemma 1** (Scalar-to-tensor approximation). *Let $\|\mathcal{S}^*\|_\infty \leq B$. For any scalar function g and any*
417 *degree- P polynomial q_P ,*

$$\|g(\mathcal{S}^*) - q_P(\mathcal{S}^*)\|_F \leq \sqrt{\prod_{n=1}^N I_n} \sup_{|s| \leq B} |g(s) - q_P(s)|. \quad (23)$$

418 *Proof.* Every entry of $\mathcal{S}^*$ lies in $[-B, B]$, so every entrywise approximation error is bounded by
419 $\sup_{|s| \leq B} |g(s) - q_P(s)|$. Squaring, summing over all entries, and taking the square root gives the
420 claim. □

421 *Proof of Theorem 3.* Let $g(s) = \sum_{k=0}^{\infty} a_k s^k$ be the Taylor expansion of g around zero, and choose
 422 $q_P(s) = \sum_{k=0}^P a_k s^k$. By Cauchy's estimate,

$$|a_k| \leq \frac{M_\rho}{\rho^k}. \quad (24)$$

423 For $|s| \leq B < \rho$,

$$|g(s) - q_P(s)| \leq \sum_{k=P+1}^{\infty} |a_k| |s|^k \quad (25)$$

$$\leq M_\rho \sum_{k=P+1}^{\infty} (B/\rho)^k = M_\rho \frac{(B/\rho)^{P+1}}{1 - B/\rho}. \quad (26)$$

424 Applying Lemma 1 and taking $\mathcal{S} = \mathcal{S}^*$, $\beta_k = a_k$ for $0 \leq k \leq P$, yields the stated unnormalized
 425 bound. Dividing by $\sqrt{M}$, $M = \prod_{n=1}^N I_n$, gives the normalized form used in Theorem 3. $\square$

426 B.2 Rank Expansion

427 *Proof of Proposition 1.* Write the Tucker tensor entrywise as

$$\mathcal{S}_{i_1, \dots, i_N} = \sum_{r_1=1}^{R_1} \cdots \sum_{r_N=1}^{R_N} \mathcal{X}_{r_1, \dots, r_N} \prod_{n=1}^N A_{i_n, r_n}^{(n)}. \quad (27)$$

428 Taking an element-wise p -th power gives

$$(\mathcal{S}_{i_1, \dots, i_N})^p = \prod_{q=1}^p \left(\sum_{r_1^{(q)}=1}^{R_1} \cdots \sum_{r_N^{(q)}=1}^{R_N} \mathcal{X}_{r_1^{(q)}, \dots, r_N^{(q)}} \prod_{n=1}^N A_{i_n, r_n^{(q)}}^{(n)} \right) \quad (28)$$

$$= \sum_{\gamma_1 \in [R_1]^p} \cdots \sum_{\gamma_N \in [R_N]^p} \left(\prod_{q=1}^p \mathcal{X}_{r_1^{(q)}, \dots, r_N^{(q)}} \right) \prod_{n=1}^N \left(\prod_{q=1}^p A_{i_n, r_n^{(q)}}^{(n)} \right), \quad (29)$$

429 where $\gamma_n = (r_n^{(1)}, \dots, r_n^{(p)})$. For each mode n , let $\alpha_n \in \mathbb{N}^{R_n}$ record the counts of the entries in γ_n ,
 430 so $|\alpha_n| = p$. Then

$$\prod_{q=1}^p A_{i_n, r_n^{(q)}}^{(n)} = \prod_{r=1}^{R_n} \left(A_{i_n, r}^{(n)} \right)^{\alpha_{n,r}} = A_{i_n, \alpha_n}^{(n)\langle p \rangle}. \quad (30)$$

431 Group all ordered tuples $(\gamma_1, \dots, \gamma_N)$ that produce the same multi-indices $(\alpha_1, \dots, \alpha_N)$, and define
 432 the aggregated core

$$\mathcal{X}_{\alpha_1, \dots, \alpha_N}^{\langle p \rangle} = \sum_{\substack{\gamma_1, \dots, \gamma_N: \\ \text{count}(\gamma_n) = \alpha_n \ \forall n}} \prod_{q=1}^p \mathcal{X}_{r_1^{(q)}, \dots, r_N^{(q)}}. \quad (31)$$

433 The preceding display is exactly the Tucker reconstruction

$$\mathcal{S}^{\odot p} = \llbracket \mathcal{X}^{\langle p \rangle}; \mathbf{A}^{(1)\langle p \rangle}, \dots, \mathbf{A}^{(N)\langle p \rangle} \rrbracket. \quad (32)$$

434 The n -th Veronese factor has $\binom{R_n + p - 1}{p}$ columns, giving the claimed Tucker rank bound. $\square$

435 B.3 Separation

436 **Lemma 2** (Full-rank Veronese evaluations). *Let $D = \binom{R+p-1}{p}$. There exists a matrix $\mathbf{A} \in \mathbb{R}^{D \times R}$*
 437 *whose Veronese feature matrix $\mathbf{A}^{\langle p \rangle} \in \mathbb{R}^{D \times D}$ has rank D . Consequently, for $I \geq D$, a generic*
 438 *matrix $\mathbf{A} \in \mathbb{R}^{I \times R}$ has $\text{rank}(\mathbf{A}^{\langle p \rangle}) = D$.*

439 *Proof.* The homogeneous degree- p monomials in R variables are linearly independent polynomials.
 440 Hence there exist evaluation points $x_1, \dots, x_D \in \mathbb{R}^R$ such that the matrix $(x_i^\alpha)_{i,\alpha}$, $|\alpha| = p$, is
 441 nonsingular; otherwise every $D \times D$ evaluation determinant would vanish, which would imply a
 442 nontrivial linear dependence among the monomials. Taking these points as the rows of $\mathbf{A}$ gives the
 443 first claim.

444 For $I \geq D$, choose any D rows. The determinant of the corresponding $D \times D$ Veronese submatrix
 445 is a polynomial in the entries of $\mathbf{A}$ that is not identically zero by the first claim. Its nonvanishing is
 446 therefore a generic property, so $\mathbf{A}^{(p)}$ has full column rank for generic $\mathbf{A}$. $\square$

447 *Proof of Theorem 1.* Let $D = \binom{R+p-1}{p}$. Choose generic factor matrices $\mathbf{A}^{(n)} = [a_1^{(n)}, \dots, a_R^{(n)}] \in$
 448 $\mathbb{R}^{I_n \times R}$, and define the CP-diagonal Tucker backbone

$$\mathcal{S} = \sum_{r=1}^R a_r^{(1)} \otimes \dots \otimes a_r^{(N)}. \quad (33)$$

449 Its mode- n unfolding is

$$\mathcal{S}_{(n)} = \mathbf{A}^{(n)} \left(\mathbf{A}^{(N)} \odot \dots \odot \mathbf{A}^{(n+1)} \odot \mathbf{A}^{(n-1)} \odot \dots \odot \mathbf{A}^{(1)} \right)^\top. \quad (34)$$

450 For generic $\mathbf{A}^{(n)}$, $\mathbf{A}^{(n)}$ has rank R , and the Khatri-Rao product on the right has rank R . Thus
 451 $\text{rank}(\mathcal{S}_{(n)}) = R$ for every n , so $\mathcal{S}$ has exact Tucker rank $(R, \dots, R)$.

452 Set $\beta_p = 1$ and $\beta_q = 0$ for $q \neq p$. Then $\mathcal{Y} = \mathcal{S}^{\odot p}$ is a valid degree- p NTD-PL tensor. Entrywise,

$$\mathcal{Y}_{i_1, \dots, i_N} = \left(\sum_{r=1}^R \prod_{n=1}^N A_{i_n, r}^{(n)} \right)^p \quad (35)$$

$$= \sum_{|\alpha|=p} c_\alpha \prod_{n=1}^N \prod_{r=1}^R \left(A_{i_n, r}^{(n)} \right)^{\alpha_r}, \quad (36)$$

453 where $c_\alpha = p! / (\alpha_1! \dots \alpha_R!) > 0$. Let $\mathbf{B}^{(n)} = \mathbf{A}^{(n)\langle p \rangle} \in \mathbb{R}^{I_n \times D}$. By Lemma 2, each $\mathbf{B}^{(n)}$ has full
 454 column rank D for generic $\mathbf{A}^{(n)}$. The mode- n unfolding is

$$\mathcal{Y}_{(n)} = \mathbf{B}^{(n)} \text{diag}(c) \left(\mathbf{B}^{(N)} \odot \dots \odot \mathbf{B}^{(n+1)} \odot \mathbf{B}^{(n-1)} \odot \dots \odot \mathbf{B}^{(1)} \right)^\top. \quad (37)$$

455 The diagonal matrix $\text{diag}(c)$ is nonsingular. Since a Khatri-Rao product of full-column-rank matrices
 456 is full column rank, the right factor has rank D . Therefore $\text{rank}(\mathcal{Y}_{(n)}) = D$. This holds for every
 457 mode n , and the exact Tucker rank of $\mathcal{Y}$ is $(D, \dots, D)$. $\square$

458 B.4 Masked Generalization

459 **Theorem 4** (Masked-entry uniform convergence). *Let $\Theta \subset \mathbb{R}^d$ be a bounded NTD-PL parameter set*
 460 *with*

$$d = \prod_{n=1}^N R_n + \sum_{n=1}^N I_n R_n + (P + 1).$$

461 *For $\theta \in \Theta$, define the full-entry risk*

$$\mathcal{L}(\theta) = \frac{1}{M} \sum_{\mathbf{i}} \ell_{\mathbf{i}}(\theta), \quad M = \prod_{n=1}^N I_n,$$

462 *and the masked empirical risk*

$$\hat{\mathcal{L}}_m(\theta) = \frac{1}{m} \sum_{t=1}^m \ell_{\mathbf{i}_t}(\theta),$$

where $\mathbf{i}_t$ are sampled uniformly from tensor entries. Assume $0 \leq \ell_{\mathbf{i}}(\theta) \leq B_\ell$ and $\ell_{\mathbf{i}}(\theta)$ is G -Lipschitz in θ , uniformly over $\mathbf{i}$. If Θ is contained in a Euclidean ball of radius B_Θ , then with probability at least $1 - \delta$,

$$\sup_{\theta \in \Theta} |\mathcal{L}(\theta) - \widehat{\mathcal{L}}_m(\theta)| \leq 2G\varepsilon + B_\ell \sqrt{\frac{d \log(3B_\Theta/\varepsilon) + \log(2/\delta)}{2m}}$$

for every $0 < \varepsilon \leq B_\Theta$.

Proof of Theorem 4. Fix $0 < \varepsilon \leq B_\Theta$, and let $\mathcal{N}_\varepsilon$ be an ε -net of Θ in Euclidean norm. Since Θ lies in a d -dimensional ball of radius B_Θ , it admits a net with

$$|\mathcal{N}_\varepsilon| \leq \left(\frac{3B_\Theta}{\varepsilon} \right)^d. \quad (38)$$

For any fixed $\theta_0 \in \mathcal{N}_\varepsilon$, the random variables $\ell_{\mathbf{i}_t}(\theta_0)$ are independent and bounded in $[0, B_\ell]$. Hoeffding's inequality gives

$$\Pr \left(|\mathcal{L}(\theta_0) - \widehat{\mathcal{L}}_m(\theta_0)| \geq u \right) \leq 2 \exp \left(-\frac{2mu^2}{B_\ell^2} \right). \quad (39)$$

Taking a union bound over $\mathcal{N}_\varepsilon$, with probability at least $1 - \delta$,

$$\max_{\theta_0 \in \mathcal{N}_\varepsilon} |\mathcal{L}(\theta_0) - \widehat{\mathcal{L}}_m(\theta_0)| \leq B_\ell \sqrt{\frac{\log(2|\mathcal{N}_\varepsilon|/\delta)}{2m}}. \quad (40)$$

Substituting the net-size bound yields the stated stochastic term.

Now take any $\theta \in \Theta$, and choose $\theta_0 \in \mathcal{N}_\varepsilon$ with $\|\theta - \theta_0\|_2 \leq \varepsilon$. By the uniform Lipschitz assumption,

$$|\mathcal{L}(\theta) - \mathcal{L}(\theta_0)| \leq G\varepsilon, \quad |\widehat{\mathcal{L}}_m(\theta) - \widehat{\mathcal{L}}_m(\theta_0)| \leq G\varepsilon. \quad (41)$$

Therefore

$$|\mathcal{L}(\theta) - \widehat{\mathcal{L}}_m(\theta)| \leq |\mathcal{L}(\theta) - \mathcal{L}(\theta_0)| + |\mathcal{L}(\theta_0) - \widehat{\mathcal{L}}_m(\theta_0)| \quad (42)$$

$$+ |\widehat{\mathcal{L}}_m(\theta_0) - \widehat{\mathcal{L}}_m(\theta)| \quad (43)$$

$$\leq 2G\varepsilon + B_\ell \sqrt{\frac{d \log(3B_\Theta/\varepsilon) + \log(2/\delta)}{2m}}. \quad (44)$$

Taking the supremum over $\theta \in \Theta$ completes the proof. $\square$

Remark 1. The Lipschitz and bounded-loss assumptions are standard for compact parametric classes. For NTD-PL, they follow when the observed values, core, factor matrices, and polynomial coefficients are bounded and the polynomial degree P is finite. The resulting Lipschitz constant may depend polynomially on these bounds and on P , but the covering dimension remains the number of tied NTD-PL parameters rather than the number of parameters in an untied high-order Tucker expansion.

C Additional Experiments

C.1 Controlled Synthetic Mechanism Checks

The controlled experiments in the main text are designed to isolate the role of the scalar link. Table 4 first checks a negative control: when the source is linear Tucker, NTD-PL does not improve by inventing nonlinear terms. The only exception is an affine offset, where the constant link term is the intended correction.

For the nonlinear controlled tensors, the nonlinearity is introduced as an orthogonal residual rather than by directly setting $Y = g(S^*)$. Let S^* be the linear low-rank Tucker backbone and let g be one of the candidate scalar responses. We form

$$R^* = g(S^*) - \frac{\langle g(S^*), S^* \rangle}{\|S^*\|_F^2} S^*, \quad (45)$$

491 and, when $\|R^*\|_F > 0$, normalize it as

$$\tilde{R}^* = \frac{\|S^*\|_F}{\|R^*\|_F} R^*. \quad (46)$$

492 The observed tensor is then

$$Y_\alpha = \sqrt{1-\alpha} S^* + \sqrt{\alpha} \tilde{R}^*, \quad 0 \leq \alpha \leq 1. \quad (47)$$

493 By construction, $\langle S^*, \tilde{R}^* \rangle = 0$ and $\|\tilde{R}^*\|_F = \|S^*\|_F$, so α is exactly the energy fraction assigned to
 494 the nonlinear residual. This keeps the total energy fixed while moving the observation away from the
 495 linear Tucker backbone. The candidate residual families are $g(s) = s^2 + s^3$, $g(s) = (\exp(\kappa s) - 1)/\kappa$,
 496 and $g(s) = \tanh(\kappa s)$, with κ set from the input range.

497 The main text reports the nonlinear alpha sweep and degree sweep. Figure 6 shows that degree contin-
 498 uation activates higher powers with systematic loss drops at relevant degrees. Table 5 complements
 499 those figures with a small but consistent term-mass concentration effect: the dominant contribution
 500 usually lands on the degree that matches the injected response, while smooth odd links spread mass
 501 across nearby odd orders.

Table 4: Linear consistency on synthetic Tucker tensors. In the linear negative control, effective $p > 1$ contributions remain small except for the intended affine offset case.

Method	bias = 0		bias = 0.5	
	RMSE↓	$p > 1$	RMSE↓	$p > 1$
Tucker	0.0114	–	0.0647	–
NTD-PL, $\beta_0 = 0$	0.0116	4.19%	0.0522	39.17%
NTD-PL, β_0 free	0.0117	3.19%	0.0118	2.98%




Figure 6: NTD-PL training curves for controlled nonlinear mechanism checks at $\alpha = 0.25$ and $P = 5$. Dashed vertical lines mark degree activations, which produce systematic drops when the activated degree matches the response structure.

Table 5: Effective NTD-PL term contributions under moderate nonlinearity, reported as percentages over ten seeds. Polynomial targets concentrate mass at matching nonlinear degrees, while odd smooth links place most nonlinear mass on odd powers.

Link	Effective contribution (%)						
	$p = 0$	$p = 1$	$p = 2$	$p = 3$	$p = 4$	$p = 5$	$p = 6$
$s^2 + s^3$	0.1	24.1	9.8	38.1	7.3	11.5	9.1
$(e^{\kappa s} - 1)/\kappa$	0.3	39.8	16.8	13.2	10.1	7.7	12.1
$\tanh(\kappa s)$	0.2	37.9	1.3	30.9	5.4	17.3	7.0

502 C.2 CAVE Completion Significance and Robustness

503 The main text reports the full-resolution random-missing CAVE completion table at rank $(33, 33, 4)$.
 504 We do not repeat that table here; instead, Table 6 gives the paired sign and Wilcoxon tests for the
 505 same protocol.

Table 6: Paired significance tests for full-resolution CAVE completion at rank (33, 33, 4). Metrics are computed only on missing entries.

ρ	RMSE				SAM			
	W/L	Med. gain	Sign p	Wilcoxon p	W/L	Med. gain	Sign p	Wilcoxon p
0.1	14/1	0.0027	9.77×10^{-4}	1.22×10^{-4}	14/1	2.9742	9.77×10^{-4}	8.54×10^{-4}
0.3	14/1	0.0027	9.77×10^{-4}	1.22×10^{-4}	13/2	1.9194	0.007	0.004
0.5	14/1	0.0026	9.77×10^{-4}	1.22×10^{-4}	13/2	1.6566	0.007	0.012
0.7	14/1	0.0027	9.77×10^{-4}	1.22×10^{-4}	12/3	1.3394	0.035	0.048

C.3 Link Refresh Residual Diagnostic

The diagnostic used in the main text reuses the Link Refresh step from Section 5 as a one-round residual probe. Let $\hat{\mathcal{Y}}_T$ denote the Tucker output and $\mathcal{Y}$ the target tensor. Starting from this fitted Tucker latent tensor, we update only the scalar polynomial coefficients by fitting

$$h_\gamma(x) = \sum_{p=0}^4 \gamma_p x^p$$

by ridge regression on pairs $(\hat{Y}_{T,i}, Y_i)$. In full reconstruction, all entries are used for this fit. In completion, the fit uses only observed entries and is then evaluated on held-out entries. The fitted map is applied entrywise to $\hat{\mathcal{Y}}_T$. We summarize this cheap refresh by the Link Refresh diagnostic

$$D_{\text{link}} = 10 \log_{10} \left(\frac{\|\mathcal{Y}_{\text{eval}} - \hat{\mathcal{Y}}_T\|_F^2}{\|\mathcal{Y}_{\text{eval}} - \hat{\mathcal{Y}}_{\text{link}}\|_F^2} \right),$$

which measures, on a log scale, how much Tucker residual energy is removed on the task-specific evaluation entries. In other words, it measures how much shared scalar-link yield is already present before the full NTD-PL alternating solver starts reshaping the Tucker parameters.

C.4 CAVE Completion Scene Diagnostics

Figure 7 supplements the main CAVE diagnostic table with the scene-level scatter plots used to visualize the link yield across missing rates. The plots show the same diagnostic relationship between D_{link} and held-out RMSE* gain, now at the scene level.




Figure 7: Scene-level supplement for the CAVE completion diagnostic table. Missing rates 0.3 and 0.5 are shown; D_{link} from frozen- backbone Link Refresh predicts larger held-out RMSE* gains.

C.5 Non-HSI Link Refresh Diagnostics

The main diagnostic table uses per-unit non-HSI experiments rather than collection-level tensors. CBSD68 uses per-image units resized to $96 \times 96 \times 3$ with rank (32, 32, 3); the available diagnostic CSV contains 16 CBSD68 units. CIFAR-10 uses 1000 class-balanced test images of shape $32 \times 32 \times 3$ with rank (20, 20, 3). COIL-100 uses one object at a time as a $36 \times 48 \times 48 \times 3$ view tensor with

rank (8, 12, 12, 2). smallNORB uses one object instance at a time as an $18 \times 64 \times 64 \times 2$ stereo view tensor with rank (12, 24, 24, 2). All units are max-normalized independently, and diagnostic labels are assigned only from D_{link} , not from NTD-PL gain.

Table 7: Exact per-dataset statistics behind the main non-HSI scatter figure. Each unit is independently fitted by Tucker and NTD-PL at matched backbone rank. D_{link} is computed by frozen-backbone Link Refresh; diagnostic tiers are assigned from D_{link} , not from NTD-PL gain.

Missing rate	Low D_{link}	Mid D_{link}	High D_{link}	Spearman ρ_s
0.1	0.75%	7.92%	12.63%	0.829
0.3	-1.13%	9.10%	13.12%	0.871
0.5	-1.09%	9.16%	13.18%	0.889
0.7	-0.62%	9.22%	13.07%	0.868

NeurIPS Paper Checklist

This draft uses a working checklist. It should be replaced by the official NeurIPS checklist distributed with the final 2026 submission template.

- Claims.** The main claims are stated in the abstract, introduction, theory section, and experiment captions: NTD-PL is a Tucker enhancement and shared-backbone rank lift for Tucker tensors; its gains are tied to jointly reshaping the Tucker backbone under a shared scalar link rather than to modest rank increases alone.
- Limitations.** Section 7 discusses the shared scalar-link assumption, rank-selection scope, polynomial-degree stability, nonconvex optimization, and the current experimental scope.
- Theory assumptions.** The approximation results assume a bounded latent tensor and polynomial or analytic scalar links. The separation results use generic full-rank feature constructions. The masked generalization bound assumes bounded parameters, bounded observations, and random observed entries.
- Reproducibility.** The experiment section reports datasets, ranks, mask types, missing rates, and baseline definitions. The codebase includes scripts for rank sweeps, MLPCal, structured missingness, and the cross-domain real-tensor reconstruction study; the final version should add hardware details and exact command manifests.