

Dictionary-Linked Tucker Models for Nonlinear Low-Rank Tensor Signal Recovery

First Author, Second Author, and Third Author

Abstract—Many multidimensional signals are compactly represented by low-rank tensor models, but nonlinear measurement responses can make an underlying low-rank signal appear higher-rank in the observed domain. This paper studies a dictionary-linked Tucker model for nonlinear low-rank tensor signal recovery. The model represents the underlying signal with a Tucker core and factors, and learns an entrywise response function with only a few extra coefficients. With a power dictionary, this response can increase the apparent Tucker rank of the measured tensor without introducing a separate high-rank Tucker model. The finite-dictionary formulation also allows other response bases when prior knowledge is available. We analyze the induced rank inflation, response approximation, stationarity of the alternating estimator, and computational scaling. We also develop a response diagnostic that predicts when the nonlinear response is likely to help. Experiments on controlled nonlinear tensors, hyperspectral reconstruction and completion, noisy reconstruction, and cross-domain diagnostics show that the learned response improves recovery when the residual follows a common entrywise response pattern. Low-diagnostic regimes clarify the operating conditions of the model.

Index Terms—Tensor decomposition, Tucker decomposition, low-rank tensor recovery, nonlinear measurements, tensor completion, hyperspectral imaging.

I. INTRODUCTION

MULTIDIMENSIONAL signals often contain strong structure across space, spectrum, time, view, or acquisition mode. Low-rank tensor models exploit this structure by representing a signal through a small number of coupled low-dimensional factors. They have become standard tools for signal representation, denoising, completion, and recovery [1]–[4]. Tucker decomposition is especially useful in this setting because it separates a compact core from mode-wise factors and provides an explicit multilinear rank budget.

In many sensing problems, however, the structured quantity of interest is not recorded directly. Camera response curves, radiometric processing, saturation, reflectance effects, and post-nonlinear hyperspectral mixing are standard examples in which a physical signal is transformed before it is stored as data [5]–[9]. Related ideas also appear in generalized low-rank and single-index matrix models, where low-dimensional structure is observed through a non-Gaussian or nonlinear entrywise map [10], [11]. These observation models share a simple lesson: a compact pre-measurement signal can look more complicated after the measurement process.

For Tucker recovery, this distinction matters because an entrywise nonlinear response can inflate the apparent multilinear rank of the measured tensor. An underlying low-rank signal

can therefore appear to require a substantially larger Tucker rank after measurement. From a signal recovery perspective, this creates a model mismatch: the rank budget selected for the underlying signal may be appropriate, while the linear Tucker observation model is not.

One way to reduce this mismatch is to increase the Tucker rank. This is often effective, but it changes the entire multilinear representation: the core grows rapidly, each mode factor gains additional degrees of freedom, and the interpretation of a fixed rank budget for the underlying signal is lost. A second approach is to replace the multilinear model with a more flexible nonlinear tensor model, such as a kernel, Gaussian-process, generalized-link, or neural tensor decoder [12]–[14]. Such methods are powerful, but their flexibility can obscure the relation between the recovered signal and a chosen Tucker representation, and may introduce additional training and model-selection burdens.

This paper studies a narrower and more structured alternative. We assume that the underlying signal is well represented by a Tucker tensor, but that the measurements are generated from that underlying signal through an unknown entrywise response. We approximate this response in a finite dictionary of scalar functions and learn the response jointly with the Tucker core and factors. If

$$\mathcal{S} = [\mathcal{X}; \mathbf{A}^{(1)}, \dots, \mathbf{A}^{(N)}] \quad (1)$$

is the underlying Tucker signal, the measured tensor is modeled as

$$\hat{\mathcal{Y}} = f_{\beta}(\mathcal{S}), \quad (2)$$

where $f_{\beta}(s) = \sum_{q=0}^Q \beta_q \psi_q(s)$ for a chosen scalar dictionary $\{\psi_q\}_{q=0}^Q$. We call this a dictionary-linked Tucker model. It keeps the Tucker ranks fixed and adds only the response coefficients $\beta = (\beta_0, \dots, \beta_Q)$. Different finite dictionaries can be used when response priors are available. In this paper, we use the simplest power dictionary $\psi_q(s) = s^q$, which gives the polynomial-linked model and makes the rank-inflation mechanism explicit. The model is therefore not an unrestricted nonlinear decoder and not a post-hoc calibration of a fixed Tucker fit. The underlying signal and the response are estimated together, so the Tucker core and factors are shaped by the nonlinear measurement model during recovery.

The key structural point is most explicit for the polynomial dictionary, where the response creates rank inflation tied to the same Tucker factors. The degree- p term $\mathcal{S}^{\odot p}$ can be viewed as a Tucker tensor whose factors are polynomial feature lifts of the original factors. Thus higher-degree channels can increase the effective multilinear rank of the measured tensor, but they remain tied to one low-rank Tucker signal. More general

dictionaries keep the same modeling principle: the additional expressivity must be explainable by one shared entrywise response acting on the underlying signal. This constraint is what makes the model useful in the intended regime and limited outside it. If the remaining residual is ordinary multilinear structure, increasing Tucker rank is the appropriate remedy; if the residual is coherent with a measurement response, the learned response can provide a compact correction.

We develop this idea for masked tensor signal recovery, covering both full reconstruction and tensor completion. The estimation procedure alternates between a small ridge-regression update for the response coefficients and gradient-based updates of the Tucker core and factors. To improve numerical stability, the implementation uses normalized features and continuation when the selected dictionary is nested. We also introduce a response diagnostic that measures how much residual energy can be removed by refitting only the response after a linear Tucker fit. This diagnostic is designed to answer a practical question before full nonlinear training: does this signal contain residual structure that is likely to benefit from a shared entrywise response?

The resulting study is organized around the full signal-recovery pipeline: modeling, estimation, diagnosis, and validation. On the rank side, we characterize when an entrywise response creates efficient rank inflation and when increasing ordinary Tucker rank is the more direct explanation. On the estimation side, we use the special structure of the response update to keep the nonlinear part low-dimensional. On the experimental side, we evaluate both high- and low-diagnostic regimes to clarify the operating conditions of the model.

A. Contributions

The main contributions are summarized as follows.

- We formulate nonlinear low-rank tensor recovery under an unknown entrywise measurement response, separating the rank of the underlying signal from the rank inflation observed after measurement.
- We propose a dictionary-linked Tucker model that keeps the Tucker ranks fixed while learning an entrywise response with a small number of coefficients. The estimator is instantiated with a power dictionary, yielding a compact polynomial response tied to the same Tucker core and factors.
- We analyze the resulting model class and estimation procedure, including induced rank inflation, parameter economy relative to explicit higher-rank Tucker models, response approximation, descent properties, diagnostic interpretation, and computational scaling.
- We evaluate the model on controlled nonlinear tensor signals and hyperspectral recovery tasks, with rank-inflation comparisons, noisy-observation robustness tests, and diagnostic stratification that identifies when the nonlinear response is useful.

The remainder of the paper develops the model, analysis, estimation algorithm, diagnostic score, and experimental validation in this order.

II. RELATED WORK

A. Low-Rank Tensor Signal Models

Low-rank tensor decompositions are a standard language for multidimensional signal modeling. CP and Tucker models have been used for compression, denoising, source separation, completion, and data fusion across signal processing and machine learning [1], [2], [15]. Tucker decomposition, in particular, represents a tensor by a small core and mode-wise factors [16], making the multilinear rank explicit. Other structured formats, including tensor-train and tensor-ring decompositions [17], [18], control complexity through different network topologies. The present paper uses Tucker because its mode-wise ranks give a direct way to distinguish the rank of the underlying signal from the effective rank of the measured tensor.

Tensor completion methods exploit these low-dimensional structures to recover missing entries. Convex and nonconvex approaches include nuclear-norm inspired formulations, factorized Tucker or CP objectives, and practical algorithms for incomplete tensors [3], [4], [19]. These methods usually model the measured tensor itself as low-rank. Our setting is different: the underlying signal is low-rank, but an entrywise response can lift the measured tensor outside the same rank budget. Increasing Tucker rank can reduce the mismatch, but doing so changes every mode factor and grows the core. NTD-PL instead keeps the original Tucker ranks and adds a response function with only a few extra coefficients.

B. Generalized and Nonlinear Tensor Models

Several tensor models replace the Gaussian squared-loss view with generalized or nonlinear observation models. Generalized CP and probabilistic tensor factorizations model non-Gaussian observations through a chosen likelihood or link function [20], [21], while nonparametric Bayesian tensor models increase flexibility through latent function priors [22]. Kernel, functional, and neural tensor methods further expand the model class by learning nonlinear feature maps or deep decoders [12]–[14], [23]. These approaches are valuable when the data require broad nonlinear expressivity, but they can make it harder to interpret how a fixed multilinear rank budget is being used.

There is also a matrix literature in which low-rank structure is observed through non-Gaussian losses or an unknown entrywise transfer function [10], [11]. This viewpoint is close in spirit to our setting, but the present work focuses on multilinear Tucker structure, rank inflation under polynomial responses, and tensor recovery diagnostics.

The dictionary-linked Tucker model uses a constrained entrywise response. It assumes that the dominant residual left by a fixed Tucker fit can be explained by a shared entrywise map, while preserving the interpretability of the Tucker ranks. This places the model between classical low-rank recovery and general nonlinear tensor factorization. The power-dictionary instance is expressive enough to approximate common smooth responses on a bounded signal range, while remaining structured enough to admit a rank-inflation analysis and a closed-form response update. Other finite scalar dictionaries could

be substituted when calibration constraints suggest a different response basis, without replacing the Tucker core and factors.

C. Imaging Recovery and Nonlinear Responses

Nonlinear responses arise naturally in imaging and sensing pipelines. Camera response curves, radiometric processing, saturation, and high-dynamic-range imaging motivate explicit response modeling in computational photography [5], [6]. Hyperspectral imaging adds additional sources of nonlinear behavior, including nonlinear mixing, reflectance effects, and sensor-dependent preprocessing [7]–[9]. In particular, post-nonlinear mixing models use a nonlinear map after a lower-dimensional mixing step, which is an important precedent for modeling a structured signal before nonlinear observation. At the same time, hyperspectral cubes are strongly correlated across spatial and spectral modes, making them natural testbeds for low-rank tensor recovery.

Hyperspectral reconstruction and completion provide a signal recovery setting where fixed-rank Tucker mismatch is observable and measurable. The proposed response diagnostic estimates, before full nonlinear training, whether a fixed Tucker residual contains structure that the entrywise response can remove. This makes the model selection criterion explicit and testable.

III. DICTIONARY-LINKED TUCKER REPRESENTATION

We model the underlying signal with Tucker factors because multidimensional signals often have mode-specific correlations across space, spectrum, time, or view, and Tucker ranks provide an explicit budget for this multilinear structure. The nonlinear observation model then describes how this low-rank Tucker signal is transformed into the measured tensor by the sensor or acquisition pipeline.

A. Nonlinear Low-Rank Signal Model

Let $\mathcal{Y} \in \mathbb{R}^{I_1 \times \dots \times I_N}$ denote the measured tensor and let $\Omega \in \{0, 1\}^{I_1 \times \dots \times I_N}$ be the observation mask. We assume that the measured entries are generated from an underlying signal $\mathcal{S}$ through an entrywise response,

$$Y_i = f(S_i) + \varepsilon_i, \quad \mathbf{i} = (i_1, \dots, i_N), \quad (3)$$

where ε_i denotes measurement noise. The underlying signal is represented by a Tucker model,

$$\mathcal{S} = \llbracket \mathcal{X}; \mathbf{A}^{(1)}, \dots, \mathbf{A}^{(N)} \rrbracket, \quad (4)$$

where $\mathcal{X} \in \mathbb{R}^{R_1 \times \dots \times R_N}$ is the core and $\mathbf{A}^{(n)} \in \mathbb{R}^{I_n \times R_n}$ are mode factors.

This model separates two notions of complexity. The underlying signal is controlled by the Tucker rank $(R_1, \dots, R_N)$, while the measured tensor may have larger effective multilinear rank because of the response f . The recovery problem is therefore to estimate the low-rank Tucker structure and the response from the observed entries of $\mathcal{Y}$.

B. Scalar Link Dictionaries

We approximate the unknown response in (3) by a finite scalar dictionary,

$$\hat{\mathcal{Y}} = f_{\beta}(\mathcal{S}) = \sum_{q=0}^Q \beta_q \psi_q(\mathcal{S}), \quad (5)$$

where $\psi_q(\mathcal{S})$ applies ψ_q entrywise. Standard Tucker is recovered whenever the dictionary contains the identity channel and the fitted response is $f_{\beta}(s) = s$. The default choice is the power dictionary $\psi_q(s) = s^q$, which gives the polynomial-linked Tucker model

$$f_{\beta}(\mathcal{S}) = \sum_{p=0}^P \beta_p \mathcal{S}^{\odot p}. \quad (6)$$

For $p \geq 2$, the degree- p channel $\mathcal{S}^{\odot p}$ uses products of entries from the same Tucker signal. Thus the model can increase effective multilinear rank without adding an independent high-rank tensor block. Appendix A-A gives the explicit p -fold Tucker construction from which this tied channel is obtained.

The same estimator can also support non-power dictionaries that remain linear in the coefficients β , for example conditioned polynomial bases, radial basis functions, or splines. These choices change the entrywise response family but not the Tucker core and factors, the masked objective, or the response-refresh subproblem. Unless otherwise stated, the remainder of the paper uses the power dictionary.

C. Masked Recovery Objective

The same formulation covers full reconstruction and completion. We estimate the core, factors, and response coefficients by minimizing

$$\min_{\mathcal{X}, \{\mathbf{A}^{(n)}\}, \beta} \frac{1}{2|\Omega|} \|\Omega \odot (\mathcal{Y} - f_{\beta}(\mathcal{S}))\|_F^2 + \mathcal{R}(\mathcal{X}, \{\mathbf{A}^{(n)}\}, \beta), \quad (7)$$

where $\mathcal{R}$ contains regularization terms for the Tucker parameters and the response coefficients. When Ω is the all-one tensor, (7) gives full reconstruction. When Ω selects a subset of entries, the loss is evaluated only on observed measurements and prediction is assessed on the held-out entries.

IV. ESTIMATION ALGORITHM AND RESPONSE DIAGNOSTIC

Let $m = |\Omega|$ and let $\theta = (\mathcal{X}, \{\mathbf{A}^{(n)}\}_{n=1}^N, \beta)$ collect all trainable parameters. For an active response dictionary of size $Q + 1$, we minimize

$$\mathcal{F}_Q(\theta) = \frac{1}{2|\Omega|} \|\Omega \odot (\mathcal{Y} - f_{\beta}(\mathcal{S}))\|_F^2 + \frac{\lambda_X}{2} \|\mathcal{X}\|_F^2 + \sum_{n=1}^N \frac{\lambda_n}{2} \|\mathbf{A}^{(n)}\|_F^2 + \frac{\lambda_{\beta}}{2} \|\beta\|_2^2, \quad (8)$$

where $f_{\beta}(s) = \sum_{q=0}^Q \beta_q \psi_q(s)$ and $\{\psi_q\}_{q=0}^Q$ is the selected response dictionary. The monomial basis $\psi_q(s) = s^q$ gives (6); in implementation, the same polynomial space may be represented by centered, Chebyshev, or otherwise conditioned

bases. Non-polynomial choices such as RBF or spline dictionaries use the same objective as long as the response remains linear in β . The experiments use the monomial power dictionary. The estimation alternates between an exact small response-coefficient update and a descent step for the Tucker core and factors.

A. Response Refresh as a Ridge Subproblem

Fix the current Tucker signal $\mathcal{S}$ and list the observed entries as $y_\Omega \in \mathbb{R}^m$ and $s_\Omega \in \mathbb{R}^m$. Define the feature matrix $\Phi_{\Omega,Q}(\mathcal{S}) \in \mathbb{R}^{m \times (Q+1)}$ by

$$[\Phi_{\Omega,Q}(\mathcal{S})]_{j,q} = \psi_q([s_\Omega]_j), \quad q = 0, \dots, Q. \quad (9)$$

The response update is the ridge regression

$$\beta^+ = \arg \min_{\beta \in \mathbb{R}^{Q+1}} \frac{1}{2m} \|y_\Omega - \Phi_{\Omega,Q}(\mathcal{S})\beta\|_2^2 + \frac{\lambda_\beta}{2} \|\beta\|_2^2, \quad (10)$$

with closed-form solution

$$\beta^+ = \left(\frac{1}{m} \Phi_{\Omega,Q}^T \Phi_{\Omega,Q} + \lambda_\beta \mathbf{I} \right)^{-1} \frac{1}{m} \Phi_{\Omega,Q}^T y_\Omega. \quad (11)$$

This step solves only a $(Q+1) \times (Q+1)$ system. It is therefore cheap relative to reconstructing or differentiating through the Tucker tensor, and because it solves the β -subproblem exactly, it never increases $\mathcal{F}_Q$ for the current Tucker signal.

B. Descent for the Tucker Parameters

With β fixed, the residual is backpropagated through the entrywise response and then through the Tucker map. The gradient with respect to the Tucker signal is

$$\mathbf{G}_\mathcal{S} = \frac{1}{m} \Omega \odot (f_\beta(\mathcal{S}) - \mathcal{Y}) \odot f'_\beta(\mathcal{S}). \quad (12)$$

The Tucker gradients are standard reverse contractions:

$$\nabla_{\mathcal{X}} \mathcal{F}_Q = \mathbf{G}_\mathcal{S} \times_1 (\mathbf{A}^{(1)})^T \cdots \times_N (\mathbf{A}^{(N)})^T + \lambda_{\mathcal{X}} \mathcal{X}, \quad (13)$$

$$\nabla_{\mathbf{A}^{(n)}} \mathcal{F}_Q = [\mathbf{G}_\mathcal{S}]_{(n)} \mathbf{Z}_n^T + \lambda_n \mathbf{A}^{(n)}, \quad (14)$$

where $\mathbf{Z}_n = (\mathbf{A}^{(N)} \otimes \cdots \otimes \mathbf{A}^{(n+1)} \otimes \mathbf{A}^{(n-1)} \otimes \cdots \otimes \mathbf{A}^{(1)}) \mathcal{X}_{(n)}^T$. The Tucker update can use a line-search gradient step, a block coordinate gradient step, or an adaptive first-order implementation, provided the step satisfies the sufficient decrease condition used in Section V.

C. Dictionary Continuation and Normalization

Large or high-order response dictionaries can amplify Tucker scale ambiguity. For nested dictionaries, such as the power or Chebyshev bases, we use continuation: training starts with the Tucker-compatible identity response and activates additional channels gradually. When a new channel is introduced, its coefficient is initialized at zero, so the prediction is unchanged at the moment of activation. Non-nested dictionaries such as RBF and spline bases are trained with the full selected dictionary from the start. After Tucker updates, factor columns are rescaled and the inverse scales are absorbed into the core. This fixes a numerical gauge without changing the represented signal $\mathcal{S}$ or the objective value.

Algorithm 1 Alternating recovery for dictionary-linked Tucker models.

Require: measured tensor $\mathcal{Y}$, mask Ω , Tucker ranks $(R_1, \dots, R_N)$, response dictionary $\{\psi_q\}$, regularization weights, continuation schedule

- 1: Initialize $\mathcal{X}$ and $\{\mathbf{A}^{(n)}\}$ from a Tucker fit or a random Tucker factorization.
- 2: Set the active dictionary to the identity response and initialize $\beta = (0, 1, 0, \dots, 0)$.
- 3: **for** outer iterations $k = 0, 1, \dots$ **do**
- 4: Form the current Tucker signal $\mathcal{S} = \llbracket \mathcal{X}; \mathbf{A}^{(1)}, \dots, \mathbf{A}^{(N)} \rrbracket$.
- 5: Build $\Phi_{\Omega,Q}(\mathcal{S})$ on the observed entries.
- 6: Refresh β by the ridge solution in (11).
- 7: Compute $\mathbf{G}_\mathcal{S}$ from (12).
- 8: Update $\mathcal{X}$ and $\{\mathbf{A}^{(n)}\}$ by a descent step using (12)–(14).
- 9: Normalize factor columns and absorb inverse scales into $\mathcal{X}$.
- 10: **if** the continuation schedule activates a new dictionary channel **then**
- 11: Add the new response coefficient and initialize it at zero.
- 12: **end if**
- 13: **if** the stopping criterion is met **then**
- 14: Stop.
- 15: **end if**
- 16: **end for**

Ensure: fitted Tucker signal $\mathcal{S}$, response coefficients β , and prediction $f_\beta(\mathcal{S})$

It does not make the Tucker factors identifiable in a strict statistical sense; it only keeps the Tucker parameterization and response coefficients on stable numerical scales during training.

Algorithm 1 treats the response as part of the fitted representation: the Tucker update is computed through the current nonlinear measurement map, rather than by post-processing a fixed Tucker fit.

D. Response Diagnostic

The diagnostic measures whether the residual of a linear Tucker fit is coherent with an entrywise response. Let $\mathcal{S}_T$ be a fitted Tucker reconstruction and let $\beta_T = (0, 1, 0, \dots, 0)$ denote the identity response. We define the diagnostic as

$$E_{\text{lin}} = \|y_\Omega - \Phi_{\Omega,1}(\mathcal{S}_T)\beta_T\|_2^2, \quad (15)$$

$$E_{\text{link}} = \min_{\beta} \left[\|y_\Omega - \Phi_{\Omega,Q}(\mathcal{S}_T)\beta\|_2^2 + m\lambda_\beta \|\beta\|_2^2 \right], \quad (16)$$

$$D_{\text{link}} = \frac{E_{\text{lin}} - E_{\text{link}}}{E_{\text{lin}}}. \quad (17)$$

Large D_{link} means that refitting only the response coefficients can already remove substantial residual energy without changing the Tucker fit. We use this score as a model-selection and interpretation tool: high diagnostic yield suggests that the nonlinear response is aligned with the residual structure, whereas low diagnostic yield suggests that increasing multilinear rank or changing the tensor prior may be more appropriate.

V. ANALYSIS

This section gives rank, optimization, and computational results for dictionary-linked Tucker models. It addresses three questions: how an entrywise response changes the Tucker rank seen in the measured tensor, what response functions can be generated or approximated by the power dictionary, and what descent and computational guarantees hold for the alternating estimator. The rank results are stated for the power dictionary, where response-induced rank inflation has an explicit algebraic

form. The optimization results apply to any differentiable finite dictionary whose response is linear in the response coefficients.

A. Rank Inflation Under Nonlinear Observation

We begin with the rank effect of the response. These results compare the Tucker rank of the underlying signal with the Tucker rank that may appear after entrywise nonlinear measurement. They do not claim that the proposed model is a general substitute for arbitrary higher-rank Tucker models; rather, they identify the special rank inflation that can be explained by one shared response function.

Proposition 1 (Response-induced Tucker rank inflation). *Let $\mathcal{S}$ have Tucker rank at most $(R_1, \dots, R_N)$. For every $p \geq 1$, $\mathcal{S}^{\odot p}$ admits a Tucker representation with mode ranks at most*

$$\left(\binom{R_1 + p - 1}{p}, \dots, \binom{R_N + p - 1}{p} \right). \quad (18)$$

Thus each polynomial channel can be represented as an ordinary Tucker tensor with lifted factors, but those lifted factors remain tied to the same original factors. This is the structural difference between the proposed model and an untied higher-rank correction.

Corollary 1 (Parameter cost of response-induced rank inflation). *Let $D_n(p) = \binom{R_n + p - 1}{p}$. A Tucker model that explicitly represents the degree- p channel in Proposition 1 may require*

$$d_{\text{lift}}(p) = \prod_{n=1}^N D_n(p) + \sum_{n=1}^N I_n D_n(p) \quad (19)$$

parameters for that channel alone. The power-dictionary linked model represents all degrees $0, \dots, Q$ with the original Tucker parameters plus $Q+1$ response coefficients. Hence the extra parameter cost of the linked model is independent of the ambient dimensions $\{I_n\}$ and the lifted ranks $\{D_n(p)\}$.

This corollary does not imply that an explicit higher-rank Tucker model is a poor model; it explains why the linked model is economical when the high-rank structure is specifically generated by a shared entrywise response.

Theorem 1 (Separation from matched-rank Tucker). *Let $N \geq 3$, $R \geq 2$, and $p \geq 2$. Define $D = \binom{R+p-1}{p}$ and suppose $I_n \geq D$ for every mode. Then there exists an order- N tensor representable by a degree- p polynomial-linked Tucker model with Tucker rank $(R, \dots, R)$, but whose exact standard Tucker rank is $(D, \dots, D)$.*

The theorem gives a strict algebraic separation: the power response does more than post-calibrate an already fitted tensor. It can generate tensors whose ordinary multilinear ranks match the Veronese-lifted ranks in Proposition 1, while the trainable Tucker signal remains at rank $(R, \dots, R)$. This is an existence result for response-generated high-rank structure, not a claim that the linked model spans all tensors of the larger Tucker rank.

B. Response Generation and Approximation

The previous results explain how a polynomial response can inflate the Tucker rank observed after measurement. The next two results ask what kinds of entrywise responses the power dictionary can represent on top of the same low-rank Tucker signal.

Proposition 2 (Exact generation for polynomial responses). *If $\mathcal{Y} = g(\mathcal{S}^*)$, where $\mathcal{S}^*$ has Tucker rank at most $(R_1, \dots, R_N)$ and g is a polynomial of degree at most P , then a degree- P polynomial-linked Tucker model with the same Tucker rank can represent $\mathcal{Y}$ exactly.*

Theorem 2 (Continuous response approximation). *Let $\mathcal{S}^*$ have Tucker rank at most $(R_1, \dots, R_N)$ and entries in a compact interval $[a, b]$. If the measurement response f is continuous on $[a, b]$, then for every $\epsilon > 0$ there exists a polynomial degree P and coefficients β such that*

$$\|f(\mathcal{S}^*) - f_\beta(\mathcal{S}^*)\|_F \leq \epsilon \sqrt{M}, \quad M = \prod_{n=1}^N I_n. \quad (20)$$

The exact polynomial result is a useful special case. The approximation result is broader: the power response is not restricted to exactly polynomial sensors. On a bounded signal range, it can approximate any smooth or continuous entrywise response to the desired uniform accuracy, while keeping the Tucker rank of the underlying signal fixed. The degree required for a given accuracy depends on the regularity of the response and the width of the signal range. This approximation result is the reason the power dictionary is used as the default response basis in the experiments.

C. Descent and Stationarity

We analyze the fixed-dictionary objective $\mathcal{F}_Q$ in (8). For nested dictionaries, the continuation schedule can be viewed as a sequence of such fixed-dictionary problems, each initialized from the previous dictionary size.

Assumption 1. *For a fixed active dictionary size $Q+1$, the iterates remain in a bounded normalized parameter set. The regularization weights are nonnegative and $\lambda_\beta > 0$. The Tucker step either uses a line search or an equivalent descent rule such that, for constants $c > 0$ and $b > 0$,*

$$\mathcal{F}_Q(\theta^k) - \mathcal{F}_Q(\theta^{k+1}) \geq c \|\theta^{k+1} - \theta^k\|_2^2, \quad (21)$$

$$\text{dist}(0, \partial \mathcal{F}_Q(\theta^{k+1})) \leq b \|\theta^{k+1} - \theta^k\|_2. \quad (22)$$

The boundedness condition is enforced in practice by Tucker gauge normalization and regularization. The two descent conditions are standard for block coordinate and line-search gradient schemes on smooth objectives.

Proposition 3 (Monotone objective decrease). *Under Assumption 1, every fixed-dictionary outer iteration of the alternating estimator satisfies*

$$\mathcal{F}_Q(\theta^{k+1}) \leq \mathcal{F}_Q(\theta^k). \quad (23)$$

Moreover, if (21) holds for the full outer update, then

$$\sum_{k=0}^{\infty} \|\theta^{k+1} - \theta^k\|_2^2 < \infty. \quad (24)$$

Theorem 3 (Convergence to a stationary point). *Suppose Assumption 1 holds for a fixed dictionary size $Q + 1$. Then the sequence generated by the alternating estimator has finite length and converges to a critical point θ^* of $\mathcal{F}_Q$:*

$$0 \in \partial \mathcal{F}_Q(\theta^*). \quad (25)$$

The theorem does not claim global optimality; the problem remains nonconvex because of the Tucker factorization. Its role is to rule out pathological cycling for the proposed alternating estimator and to place the algorithm on the same theoretical footing as standard nonconvex block-descent methods [24], [25].

D. Diagnostic Interpretation

The diagnostic D_{link} in (17) can be interpreted as a projection test for residual structure that depends on the current Tucker estimate. Let $\Phi_{\Omega, Q}$ be the response feature matrix built from a fixed Tucker reconstruction $\mathcal{S}_T$ and let $\mathbf{P}_Q$ denote the orthogonal projector onto its column space. When $\lambda_\beta = 0$, the refreshed-response residual satisfies

$$E_{\text{link}} = \|(\mathbf{I} - \mathbf{P}_Q)y_\Omega\|_2^2. \quad (26)$$

Consequently, D_{link} is the fraction of identity-response residual energy removed by projecting the observations onto the response-feature subspace generated by the fixed Tucker fit. With ridge regularization, the same interpretation holds with $\mathbf{P}_Q$ replaced by the corresponding shrinkage operator. Large D_{link} therefore means that the residual is aligned with scalar functions of the current Tucker estimate, whereas small D_{link} suggests that the residual is better addressed by increasing multilinear rank, changing the tensor prior, or revisiting the observation model.

E. Masked Generalization

The masked objective also has the usual sampling interpretation. Let $\Theta \subset \mathbb{R}^d$ be a bounded parameter set for a degree- P polynomial-linked Tucker model, where

$$d = \prod_{n=1}^N R_n + \sum_{n=1}^N I_n R_n + (P + 1). \quad (27)$$

For $\theta \in \Theta$, define the full-entry risk $\mathcal{L}(\theta) = M^{-1} \sum_{\mathbf{i}} \ell_{\mathbf{i}}(\theta)$, $M = \prod_n I_n$, and the masked empirical risk $\hat{\mathcal{L}}_m(\theta) = m^{-1} \sum_{t=1}^m \ell_{\mathbf{i}_t}(\theta)$, where the entries $\mathbf{i}_t$ are sampled uniformly.

Theorem 4 (Masked-entry uniform convergence). *Assume $0 \leq \ell_{\mathbf{i}}(\theta) \leq B_\ell$ and $\ell_{\mathbf{i}}(\theta)$ is G -Lipschitz in θ , uniformly over $\mathbf{i}$. If Θ lies in a Euclidean ball of radius B_Θ , then, with probability at least $1 - \delta$,*

$$\sup_{\theta \in \Theta} \left| \mathcal{L}(\theta) - \hat{\mathcal{L}}_m(\theta) \right| \leq 2G\varepsilon + B_\ell \sqrt{\frac{d \log(3B_\Theta/\varepsilon) + \log(2/\delta)}{2m}} \quad (28)$$

for every $0 < \varepsilon \leq B_\Theta$.

This bound is not intended as a sharp sample-complexity result. Its role is to make clear which complexity enters the masked recovery problem: the covering dimension is the tied dictionary-linked parameter count, rather than the parameter count of an untied Tucker model with Veronese ranks.

F. Computational Scaling

Let $M = \prod_{n=1}^N I_n$, $R_\Pi = \prod_{n=1}^N R_n$, and $d_{\text{Tucker}} = R_\Pi + \sum_{n=1}^N I_n R_n$. Storing the dictionary-linked model requires

$$d_{\text{NTD}} = d_{\text{Tucker}} + (Q + 1) \quad (29)$$

parameters. A higher-rank Tucker model that explicitly matches the degree- Q Veronese lift may require mode ranks $D_n = \binom{R_n+Q-1}{Q}$ for the degree- Q channel, with core size $\prod_n D_n$ and factor size $\sum_n I_n D_n$. The proposed model keeps the storage cost tied to the original Tucker core and factors and avoids explicit storage of the lifted tensor.

For one outer iteration, Response Refresh forms $\Phi_{\Omega, Q}^T \Phi_{\Omega, Q}$ and solves a small system, costing

$$\mathcal{O}(mQ^2 + Q^3). \quad (30)$$

The dominant Tucker cost is evaluating the Tucker reconstruction and reverse contractions on the observed set or on the full tensor, depending on the implementation. Dense reconstruction scales as

$$\mathcal{O}\left(MR_\Pi + \sum_{n=1}^N I_n R_n R_\Pi / R_n\right) \quad (31)$$

up to constants determined by contraction order, followed by an additional $\mathcal{O}(MQ)$ cost to evaluate the entrywise response for dense dictionaries. In sparse masked settings, the same operations can be restricted to observed entries, replacing M by m in the response and residual computations. Since Q is small, the response-update cost is typically negligible compared with the Tucker contractions.

VI. EXPERIMENTS

The experiments evaluate the proposed model through three connected questions. First, does the learned response behave as predicted on controlled nonlinear measurements? Second, under the same Tucker rank, does the learned response improve real tensor reconstruction and completion consistently across scenes, with separate stress tests for missing entries and observation noise? Third, can a low-cost diagnostic indicate when the response model is appropriate? These questions match the modeling claim: NTD-PL is expected to be most beneficial when a matched-rank Tucker model leaves residual structure consistent with a shared entrywise response. Unless stated otherwise, percentage gains denote relative error reduction from Tucker to NTD-PL. The main real-data evidence is based on all CAVE scenes under matched Tucker rank; noisy reconstruction and cross-domain diagnostics are used to test robustness and operating conditions.

A. Controlled Nonlinear Tensor Recovery

The controlled experiment isolates the measurement-response mechanism without missing entries or observation noise. We start from an underlying Tucker signal $\mathcal{S}^*$ of known rank and generate the measured tensor by adding a nonlinear response component that is orthogonal to the constant and linear channels. Let $\bar{\mathcal{S}}$ denote the normalized underlying signal

and let $\bar{\mathcal{R}}$ denote the normalized nonlinear residual after this projection. The clean measurement is generated as

$$\mathcal{Y}_\alpha = \text{scale}_{[0,1]}(\sqrt{1-\alpha}\bar{\mathcal{S}} + \sqrt{\alpha}\bar{\mathcal{R}}), \quad (32)$$

where α controls nonlinear residual energy after the constant and linear response components have been removed. This construction keeps the intended Tucker rank fixed while making the measured tensor increasingly mismatched to a linear Tucker observation model.

Figure 1 shows the expected transition. When α is small, the measured tensor is close to the underlying Tucker signal and NTD-PL stays close to matched-rank Tucker. As the nonlinear residual energy increases, the learned response explains a growing part of the residual and the gap widens across all four target responses: s^2 , $s^2 + s^3$, $\tanh(\kappa s)$, and $(e^{\kappa s} - 1)/\kappa$. The degree-dependent behavior is reported in Appendix C-A.

All experiments in the main text use the power dictionary because it is compact, interpretable, and directly supports the rank-inflation analysis in Section V. Appendix C-A adds controlled details, including a linear-response case in which high-order channels remain inactive. These full-observation experiments serve as mechanism checks; masked recovery and real tensor recovery are evaluated in the following subsections.

B. Hyperspectral Tensor Recovery

We next evaluate whether the same linked Tucker representation improves real spectral-spatial tensors. Unless otherwise stated, hyperspectral cubes are globally normalized, the Tucker rank is fixed across Tucker and NTD-PL, and completion metrics are computed only on held-out entries. We report root mean square error (RMSE) and spectral angle mapper (SAM); lower values are better for both.

1) *CAVE Reconstruction and Completion*: We first use the CAVE multispectral image database [26], [27], which provides full spectral-spatial scenes and therefore supports both full-observation reconstruction and masked completion. Full reconstruction tests whether the learned response adds a compact correction at a fixed Tucker rank. Across the tested ranks, NTD-PL consistently improves matched-rank Tucker in both RMSE and SAM when averaged over all 15 scenes, with larger gains at lower ranks and smaller gains as the Tucker rank grows; the full capacity curves are reported in Appendix C-C. This is the expected behavior of a fixed-rank nonlinear correction: if ordinary multilinear rank is already high enough to absorb the nonlinear response, the marginal value of the response should decrease.

Table I asks a complementary capacity question: how much ordinary Tucker rank is needed to match a lower-rank NTD-PL model? Across the tested ranks, Tucker needs roughly 51–60% more parameters to match the NTD-PL RMSE, while NTD-PL still has lower SAM. The result supports the rank-inflation interpretation in Section V: the power response reuses the same low-rank Tucker signal to express structured nonlinear residuals with a small number of response coefficients.

We then add observation noise only to the full-reconstruction setting. The model is fitted to clean or noisy full CAVE cubes and evaluated against the corresponding clean

cube. This separates radiometric robustness from missing entry recovery. Across clean observations and Gaussian noise levels down to 10 dB, NTD-PL maintains positive RMSE and SAM gains over matched-rank Tucker; the full table is reported in Appendix C-C.

Held-out completion tests whether the fixed-rank benefit transfers to unobserved entries without introducing observation noise. Table II reports CAVE completion under random masks using all 15 scenes and three mask seeds per scene. At the same Tucker rank, NTD-PL improves missing-entry RMSE by 12.5–17.7% and missing-entry SAM by 13.4–29.3% across the tested missing rates. The gains are stable across mask density, and paired scene-level summaries are reported in Appendix C-C.

C. Rank-Inflation Evidence

The reconstruction and completion tables show that the linked model improves standard accuracy metrics. We next ask whether this improvement is accompanied by the rank-inflation behavior predicted by the analysis. For this mechanism view, we use the full CAVE scenes at their native 512×512 spatial resolution and fix the Tucker rank to $(12, 12, 4)$ for both Tucker and NTD-PL. We sweep the maximum power degree P and compute the spatial-mode spectra of the learned Tucker signal, the final prediction, and the response component $f_\beta(\mathcal{S}) - \mathcal{S}$.

Figure 2 shows two complementary summaries. Increasing P raises both the reconstruction gain and the effective rank of the response component, while the learned Tucker signal itself remains at the prescribed rank. Moreover, the fraction of the measured tensor’s spatial tail energy captured by the NTD-PL prediction is strongly aligned with the RMSE gain across scenes and degrees. The Spearman correlation between spatial tail capture and RMSE gain is 0.94 in this sweep. Table III reports the corresponding degree-wise summary.

D. Learned Response Curves

The previous experiment measures rank inflation in the tensor domain. We now look directly at the learned scalar response. For each full-size CAVE scene, we fit NTD-PL at rank $(12, 12, 4)$ and evaluate the learned map $s \mapsto f_\beta(s)$ on percentiles of the learned Tucker signal. Because Tucker factors have scale and sign gauges, each curve is plotted after endpoint normalization on the central percentile range; the curve shape, rather than the absolute latent scale, is the object of interest.

Figure 3 shows that the learned response is not an arbitrary high-degree oscillation. The mean curve is consistently nonlinear and mostly monotone, and the shape changes little after moderate degree. The numerical summary in Table IV matches this visual pattern: moving from $P = 1$ to $P = 4$ increases both the nonlinear deviation from the best affine fit and the RMSE gain, while the additional change from $P = 4$ to $P = 6$ is smaller. This provides a direct response-level view of the same mechanism: the model does not replace the low-rank signal by a separate high-rank tensor, but learns a shared scalar calibration of that signal.




Fig. 1. Controlled nonlinear residual sweep for four target responses. The parameter α controls nonlinear residual energy after removing constant and linear components. NTD-PL remains close to Tucker in the nearly linear regime and separates as the nonlinear component grows.

TABLE I

HIGHER-RANK TUCKER COMPARISON ON CAVE RECONSTRUCTION. FOR EACH NTD-PL TUCKER RANK, WE REPORT THE FIRST HIGHER-RANK TUCKER MODEL MATCHING NTD-PL RMSE AND THE ADDITIONAL PARAMETER COST.

NTD-PL reference				First Tucker matching RMSE				Extra params
Rank	Params	RMSE↓	SAM↓	Rank	Params	RMSE↓	SAM↓	
(18, 18, 3)	20k	0.0319	19.23	(28, 28, 3)	31k	0.0316	21.41	+59.5%
(24, 24, 4)	27k	0.0256	15.15	(35, 35, 4)	41k	0.0254	16.26	+51.3%
(33, 33, 4)	38k	0.0223	14.41	(49, 49, 4)	60k	0.0222	15.28	+56.5%

TABLE II

CAVE RANDOM-MISSING COMPLETION AT MATCHED TUCKER RANK. METRICS ARE EVALUATED ONLY ON HELD-OUT ENTRIES; GAINS ARE RELATIVE IMPROVEMENTS OVER MATCHED-RANK TUCKER.

Missing rate	RMSE*↓			SAM*↓		
	Tucker	NTD-PL	Rel. gain	Tucker	NTD-PL	Rel. gain
0.1	0.0263 ± 0.0179	0.0229 ± 0.0158	12.9%	11.82 ± 5.33	8.36 ± 3.38	29.3%
0.3	0.0263 ± 0.0179	0.0230 ± 0.0158	12.5%	14.75 ± 6.33	12.47 ± 5.06	15.5%
0.5	0.0264 ± 0.0179	0.0230 ± 0.0158	12.9%	15.56 ± 6.85	13.48 ± 5.75	13.4%
0.7	0.0266 ± 0.0180	0.0219 ± 0.0128	17.7%	16.02 ± 7.21	13.79 ± 5.43	13.9%




Fig. 2. Rank-inflation evidence on full-size CAVE reconstruction at Tucker rank (12, 12, 4). Left: increasing the polynomial degree raises the response-component rank and reconstruction gain. Right: scene-level tail capture by the NTD-PL prediction tracks RMSE gain.

E. Spectral-Regime Error Analysis

Finally, we ask where the reconstruction error is reduced in the observed tensor. Entries are grouped by their normalized observed value and by spectral band. This is an observed-domain analysis: the intensity bins are bins of entries of the normalized measured tensor, not bins of RGB brightness or spatial pixel averages. The analysis connects the learned response curve to observable error behavior by showing which

response ranges receive the largest correction.

Figure 4 shows that the gain is not spread uniformly over the cube. The largest relative improvements occur in low-to-mid intensity deciles, whereas one high-intensity decile is nearly neutral. At the same time, the highest-intensity decile contributes a large share of the total MSE reduction because its baseline error energy is large. The spectral view shows stronger gains in the green-to-red regimes than at the shortest wavelengths. Table V gives the regime-level summary: the




Fig. 3. Learned response curves on full-size CAVE reconstruction at Tucker rank (12, 12, 4). Left: endpoint-normalized response curves for $P = 6$; thin lines are scenes and the thick line is their mean. Right: nonlinear deviation and gain over the linear response increase with degree and then stabilize.

TABLE III

RANK-INFLATION SUMMARY ON FULL-SIZE CAVE SCENES AT RANK (12, 12, 4). GAINS ARE RELATIVE TO MATCHED-RANK TUCKER.

P	RMSE gain (%)	Resp. rank	Pred. lift	Tail cap.
1	0.10	12.0	0.0	0.000
2	2.82	28.0	0.0	0.009
4	7.65	42.5	3.5	0.104
6	8.79	47.5	5.0	0.127

TABLE IV

LEARNED-RESPONSE SUMMARY ON FULL-SIZE CAVE SCENES AT RANK (12, 12, 4).

P	Gain vs. $P = 1$ (%)	Nonlin.	Mono.	Curv.
1	0.00	0.000	1.000	0.01
2	3.19	0.116	0.993	1.21
4	6.85	0.203	0.976	3.79
6	7.74	0.209	0.962	4.48

TABLE V

SPECTRAL-REGIME ERROR ANALYSIS ON FULL-SIZE CAVE SCENES AT RANK (12, 12, 4).

Regime	Pooled gain (%)	Median gain (%)	MSE contrib. (%)
400–500 nm	3.54	8.16	14.5
510–600 nm	7.60	10.50	43.5
610–700 nm	7.68	9.11	42.0

510–600 nm and 610–700 nm regimes account for most of the MSE reduction in this full-size CAVE setting. These patterns support the response interpretation from another angle: NTD-PL changes the effective calibration of specific intensity and spectral regimes rather than applying a uniform improvement everywhere.

F. Diagnostics and Operating Regimes

The diagnostic D_{link} is computed by fitting a Tucker model and measuring the residual reduction obtained by refitting only the response coefficients. It is therefore substantially cheaper than full NTD-PL training and is used here as a model-selection signal for the proposed recovery model. Table VI

TABLE VI

CAVE COMPLETION DIAGNOSTIC. SCENES ARE GROUPED BY D_{link} TERCILE; ENTRIES ARE MEAN NTD-PL RMSE GAINS OVER MATCHED-RANK TUCKER.

Missing rate	Low D_{link}	Mid D_{link}	High D_{link}	Spearman ρ_s
0.1	0.75%	7.92%	12.63%	0.829
0.3	-1.13%	9.10%	13.12%	0.871
0.5	-1.09%	9.16%	13.18%	0.889
0.7	-0.62%	9.22%	13.07%	0.868

groups CAVE completion scenes by diagnostic tercile. Higher D_{link} corresponds to larger held-out RMSE gains at every missing rate, while the low-diagnostic group has little or negative average gain. The Spearman correlations, 0.829–0.889, show that the diagnostic is strongly aligned with the eventual benefit of the full model.

Figure 5 evaluates the same diagnostic beyond hyperspectral data. Each point is an independently fitted tensor unit, and the horizontal axis is D_{link} measured before full NTD-PL training. Object-view tensors show the clearest separation, action-video tensors follow the same trend, and natural images occupy a lower-yield but still positive regime. These results reinforce the main interpretation: NTD-PL is most effective when the Tucker residual contains structure that is shared across entries through the same response function.

VII. CONCLUSION

This paper developed dictionary-linked Tucker models for nonlinear low-rank tensor signal recovery. The central idea is to keep the Tucker rank budget tied to the underlying signal while modeling the measured tensor through a learned entrywise response. The power-dictionary instance produces a compact rank inflation: polynomial channels can increase the effective rank of the observation while reusing the same Tucker core and factors. Other finite scalar dictionaries can be used in the same formulation when application-specific response priors are available, but the paper uses the power dictionary as the default response basis.




Fig. 4. Observed-domain error regimes on full-size CAVE reconstruction at Tucker rank $(12, 12, 4)$ and $P = 6$. Left: RMSE gain by normalized-intensity decile. Middle: gain by spectral band. Right: joint intensity–spectral regimes. Positive values indicate lower error for NTD-PL than Tucker.




Fig. 5. Cross-domain response diagnostic. Each point plots the diagnostic D_{link} against the RMSE gain of full NTD-PL over Tucker. Larger diagnostic yield predicts larger benefit from the learned response.

The analysis formalized this mechanism through rank inflation, response generation, response approximation, and non-convex descent results. The estimator combines a small ridge update for the response coefficients with gradient-based Tucker updates through the current nonlinear response. The response diagnostic checks whether a signal is likely to benefit from the response before full nonlinear training.

Experiments support the operating regime of the model. Controlled tensors verify that gains appear when the residual is nonlinear and the active polynomial degree is expressive enough. Hyperspectral reconstruction and completion show that, at matched Tucker rank, the learned response can improve both RMSE and spectral angle metrics, and higher-rank Tucker comparisons show that similar RMSE can require substantially larger ordinary Tucker models. Noisy-observation experiments show that the improvement remains positive under perturbed measurements. Diagnostic stratification further clarifies the operating regime: when the residual is not well explained by a shared entrywise response, increasing rank or changing the tensor prior is more appropriate.

Several directions remain open. Monotone splines or

physics-informed response families may be preferable when calibration constraints are known. Sharper sample-complexity or identifiability results, richer noise models, and domain-specific response priors are also natural extensions of the framework.

APPENDIX A ADDITIONAL MODEL DETAILS

A. Explicit p -Fold Tucker Motivation

The power-dictionary channel in (6) can be viewed as a tied version of an explicit high-order Tucker interaction. This construction is useful because it separates two ideas: the high-order interaction pattern that a polynomial channel represents, and the parameter sharing that makes the proposed model compact.

For a Tucker signal, a formal p -fold Tucker channel may be written as

$$\tilde{Y}_i^{(p)} = \sum_{\mathbf{r}^{(1)}, \dots, \mathbf{r}^{(p)}} \left(\prod_{q=1}^p X_{r_1^{(q)}, \dots, r_N^{(q)}} \right) \prod_{n=1}^N \mathcal{A}_{i_n, r_n^{(1)}, \dots, r_n^{(p)}}^{(n,p)}. \quad (33)$$

For $p = 1$, this is the ordinary Tucker entry formula. For $p = 2$, each mode projection couples two Tucker index tuples, giving a quadratic Tucker-style channel. More generally, (33) is the tensor analogue of augmenting a linear map with polynomial interactions among Tucker coordinates.

The explicit construction is not the parameterization used by the estimator: each mode projection $\mathcal{A}^{(n,p)}$ has $I_n R_n^p$ entries and a separate set of such projections would be needed for every active degree. Dictionary-linked Tucker keeps the same interaction pattern but ties each p -fold projection to the original Tucker factor,

$$\mathcal{A}_{i_n, r_n^{(1)}, \dots, r_n^{(p)}}^{(n,p)} = \prod_{q=1}^p A_{i_n, r_n^{(q)}}^{(n)}. \quad (34)$$

Substituting (34) into (33) gives exactly $\mathcal{S}^{\odot p}$, where $\mathcal{S} = [\mathcal{X}; \mathbf{A}^{(1)}, \dots, \mathbf{A}^{(N)}]$. Thus a degree- p power-response channel is a shared-projection p -fold Tucker interaction: it couples p Tucker index tuples while reusing the same mode factors.

This viewpoint explains why the model can increase effective rank without storing an untied high-order tensor block. The additional trainable variables are the response coefficients, while the high-order features are induced by repeated use of the same Tucker factors.

APPENDIX B PROOFS

A. Proof of Proposition 1

Let

$$S_i = \sum_{r_1, \dots, r_N} X_{r_1, \dots, r_N} \prod_{n=1}^N A_{i_n, r_n}^{(n)}.$$

Expanding S_i^p gives products of p Tucker index tuples. For each mode n , the resulting mode-dependent terms are all degree- p monomials in the row $A_{i_n, \cdot}^{(n)}$. These monomials are indexed by multi-indices $\alpha \in \mathbb{N}^{R_n}$ with $|\alpha| = p$, of which there are $\binom{R_n+p-1}{p}$. Collecting identical monomials yields a Tucker representation of $\mathcal{S}^{\odot p}$ whose mode- n factor is the degree- p Veronese lift of $\mathbf{A}^{(n)}$. This proves the stated rank bound.

B. Proof of Theorem 1

We first record a standard fact about Veronese evaluations. Let $D = \binom{R+p-1}{p}$. There exists a matrix $\mathbf{A} \in \mathbb{R}^{D \times R}$ whose degree- p Veronese feature matrix has rank D . Indeed, the homogeneous degree- p monomials in R variables are linearly independent polynomials, so one can choose D evaluation points for which the square evaluation matrix is nonsingular. Consequently, for $I \geq D$, a generic $\mathbf{A} \in \mathbb{R}^{I \times R}$ has a degree- p Veronese lift with full column rank D .

Choose generic factor matrices $\mathbf{A}^{(n)} = [a_1^{(n)}, \dots, a_R^{(n)}] \in \mathbb{R}^{I_n \times R}$ and define the CP-diagonal Tucker signal

$$\mathcal{S} = \sum_{r=1}^R a_r^{(1)} \otimes \dots \otimes a_r^{(N)}.$$

For generic factors, each mode unfolding of $\mathcal{S}$ has rank R , so the signal has exact Tucker rank $(R, \dots, R)$.

Now set the degree- p response coefficient to one and all other coefficients to zero, so $\mathcal{Y} = \mathcal{S}^{\odot p}$. Expanding entrywise gives

$$Y_i = \left(\sum_{r=1}^R \prod_{n=1}^N A_{i_n, r}^{(n)} \right)^p = \sum_{|\alpha|=p} c_\alpha \prod_{n=1}^N \prod_{r=1}^R \left(A_{i_n, r}^{(n)} \right)^{\alpha_r},$$

where $c_\alpha = p! / (\alpha_1! \dots \alpha_R!) > 0$. Let $\mathbf{B}^{(n)} = \mathbf{A}^{(n)\langle p \rangle} \in \mathbb{R}^{I_n \times D}$ be the Veronese lift. By the preceding full-rank fact, each $\mathbf{B}^{(n)}$ has column rank D for generic $\mathbf{A}^{(n)}$. The mode- n unfolding of $\mathcal{Y}$ factors as $\mathbf{B}^{(n)}$ times a full-column-rank right factor built from the remaining lifted factors, with a nonsingular diagonal scaling by the positive coefficients c_α . Therefore every mode unfolding has rank D . The exact standard Tucker rank of $\mathcal{Y}$ is $(D, \dots, D)$, while $\mathcal{Y}$ is represented by the degree- p linked model using the rank- $(R, \dots, R)$ Tucker signal.

C. Proof of Corollary 1

By Proposition 1, the degree- p channel can be written as a Tucker tensor with ranks at most $D_n(p) = \binom{R_n+p-1}{p}$. A standalone Tucker parameterization has $\prod_n D_n(p)$ core entries and $\sum_n I_n D_n(p)$ factor entries. In contrast, the power-dictionary linked model reuses the core and factors of $\mathcal{S}$ and adds only one response coefficient per channel. Summing degrees up to Q adds $Q + 1$ response coefficients, proving the stated comparison.

D. Proof of Proposition 2

If $g(s) = \sum_{p=0}^P \gamma_p s^p$ and $\mathcal{Y} = g(\mathcal{S}^*)$, choose the Tucker component of the power-dictionary linked model to be the Tucker representation of $\mathcal{S}^*$ and set $\beta_p = \gamma_p$ for all $p = 0, \dots, P$. Then $f_\beta(\mathcal{S}^*) = g(\mathcal{S}^*) = \mathcal{Y}$, giving exact representation with the same Tucker rank.

E. Proof of Theorem 2

By the Weierstrass approximation theorem, for every $\epsilon > 0$ there exists a polynomial $q(s) = \sum_{p=0}^P \beta_p s^p$ such that $\sup_{s \in [a, b]} |f(s) - q(s)| \leq \epsilon$. Every entry of $\mathcal{S}^*$ lies in $[a, b]$, so each entrywise approximation error is bounded by ϵ . Therefore

$$\|f(\mathcal{S}^*) - q(\mathcal{S}^*)\|_F^2 = \sum_i |f(S_i^*) - q(S_i^*)|^2 \leq M \epsilon^2,$$

which gives the stated bound after taking square roots.

F. Proof of Proposition 3

For a fixed Tucker signal, Response Refresh solves the β -subproblem exactly, hence it cannot increase $\mathcal{F}_Q$. The subsequent Tucker update is assumed to satisfy the sufficient decrease condition (21), and therefore also cannot increase $\mathcal{F}_Q$. Combining the two stages gives monotone decrease over the outer iteration. Since $\mathcal{F}_Q$ is bounded below by zero when the regularization weights are nonnegative, summing (21) over k gives

$$c \sum_{k=0}^K \|\theta^{k+1} - \theta^k\|_2^2 \leq \mathcal{F}_Q(\theta^0) - \mathcal{F}_Q(\theta^{K+1}) \leq \mathcal{F}_Q(\theta^0).$$

Letting $K \rightarrow \infty$ proves square summability of the steps.

G. Proof of Theorem 3

For fixed Q , the objective is a polynomial function of the Tucker parameters and response coefficients when the monomial basis is used. For any other fixed polynomial basis spanning the same space, it remains polynomial after a linear change of coefficients. The linear-spline dictionary is piecewise polynomial and semi-algebraic; standard RBF dictionaries are analytic and belong to the standard tame classes covered by KL convergence theory. With standard box or normalization constraints used to fix the Tucker gauge, the corresponding constrained objective satisfies the Kurdyka–Lojasiewicz property. Assumption 1 gives sufficient decrease and a relative error bound. The standard KL convergence theorem for nonconvex descent sequences then implies finite length of the iterates and convergence to a critical point of $\mathcal{F}_Q$.

H. Proof of Theorem 4

Fix $0 < \varepsilon \leq B_\Theta$, and let $\mathcal{N}_\varepsilon$ be an ε -net of Θ in Euclidean norm. Since Θ lies in a d -dimensional ball of radius B_Θ , it admits a net with

$$|\mathcal{N}_\varepsilon| \leq \left(\frac{3B_\Theta}{\varepsilon} \right)^d.$$

For any fixed $\theta_0 \in \mathcal{N}_\varepsilon$, the random variables $\ell_{i_\ell}(\theta_0)$ are independent and bounded in $[0, B_\ell]$. Hoeffding's inequality gives

$$\Pr \left(\left| \mathcal{L}(\theta_0) - \widehat{\mathcal{L}}_m(\theta_0) \right| \geq u \right) \leq 2 \exp \left(-\frac{2mu^2}{B_\ell^2} \right).$$

A union bound over $\mathcal{N}_\varepsilon$ shows that, with probability at least $1 - \delta$,

$$\max_{\theta_0 \in \mathcal{N}_\varepsilon} \left| \mathcal{L}(\theta_0) - \widehat{\mathcal{L}}_m(\theta_0) \right| \leq B_\ell \sqrt{\frac{d \log(3B_\Theta/\varepsilon) + \log(2/\delta)}{2m}}.$$

For an arbitrary $\theta \in \Theta$, choose $\theta_0 \in \mathcal{N}_\varepsilon$ with $\|\theta - \theta_0\|_2 \leq \varepsilon$. The uniform Lipschitz assumption implies

$$|\mathcal{L}(\theta) - \mathcal{L}(\theta_0)| \leq G\varepsilon, \quad |\widehat{\mathcal{L}}_m(\theta) - \widehat{\mathcal{L}}_m(\theta_0)| \leq G\varepsilon.$$

Combining these two approximation errors with the net bound and taking the supremum over Θ gives the stated inequality.

APPENDIX C

ADDITIONAL EXPERIMENTAL DETAILS

A. Controlled Nonlinear Tensors

The controlled experiment uses synthetic tensors whose underlying signal is exactly generated by a Tucker model. For each random seed, factor matrices are sampled with normalized columns, a small Tucker core is sampled independently, and the underlying signal is rescaled to a fixed numerical range before the nonlinear response is applied. The target response family is

$$h \in \{s^2, s^2 + s^3, \tanh(\kappa s), (e^{\kappa s} - 1)/\kappa\}. \quad (35)$$

For each response, the nonlinear component is projected away from the constant and linear span on the sampled signal

TABLE VII
CONTROLLED LINEAR-RESPONSE CONSISTENCY. THE COLUMN $p > 1$
REPORTS THE TOTAL CONTRIBUTION OF NONLINEAR POLYNOMIAL
CHANNELS.

Method	bias = 0		bias = 0.5	
	RMSE↓	$p > 1$	RMSE↓	$p > 1$
Tucker	0.0114	–	0.0647	–
NTD-PL, $\beta_0 = 0$	0.0116	4.19%	0.0522	39.17%
NTD-PL, β_0 free	0.0117	3.19%	0.0118	2.98%

entries. The resulting residual is then normalized so that α controls nonlinear residual energy after removing the constant and linear response components:

$$\mathcal{Y}_\alpha = \text{scale}_{[0,1]} \left(\sqrt{1 - \alpha} \bar{S} + \sqrt{\alpha} \bar{R} \right). \quad (36)$$

This construction makes the mechanism visible: a linear Tucker model receives the same Tucker rank budget for every α , while the observation becomes increasingly incompatible with a purely linear measurement map.

All controlled runs are full-observation reconstructions. Tucker, CP, tensor-train, tensor-ring, and NTD-PL are evaluated by RMSE on the generated measurement tensor. NTD-PL uses the same Tucker rank as the matched Tucker baseline and varies only the active polynomial degree. The α sweep is visualized in Fig. 1; the supplementary checks below focus on degree-dependent behavior and linear-response consistency under the same controlled generator. Table VII gives a linear-response control: when the measurement is linear and unbiased, the high-order channels remain nearly inactive and NTD-PL reduces to Tucker-level behavior; when a constant offset is present, freeing β_0 absorbs the bias without requiring high-order terms.

B. Metrics and Normalization

For a test index set Γ , the RMSE is

$$\text{RMSE} = \left(|\Gamma|^{-1} \sum_{i \in \Gamma} (Y_i - \widehat{Y}_i)^2 \right)^{1/2}. \quad (37)$$

For completion, Γ contains held-out entries only. For full reconstruction, Γ is the full tensor support. Spectral angle mapper is computed pixelwise on spectral vectors,

$$\text{SAM}(\mathbf{y}, \widehat{\mathbf{y}}) = \arccos \frac{\langle \mathbf{y}, \widehat{\mathbf{y}} \rangle}{\|\mathbf{y}\|_2 \|\widehat{\mathbf{y}}\|_2 + \epsilon}, \quad (38)$$

and then averaged over valid pixels. The small ϵ is used only to avoid division by zero for empty or near-empty spectra. Hyperspectral cubes are globally normalized before fitting, and all compared methods use the same normalization and the same train/test masks.

C. Hyperspectral Protocol

The CAVE multispectral image database [26], [27] is used for the main hyperspectral experiments because it provides fully observed spectral-spatial scenes. Full reconstruction evaluates the fitted tensor on all entries; in the noisy reconstruction




Fig. 6. Degree-dependent controlled gains. Entries report the RMSE reduction of NTD-PL relative to matched-rank Tucker; gains appear when the active polynomial degree can approximate the response.




Fig. 7. Representative optimization traces in the controlled experiment. Degree continuation activates nonlinear channels gradually, so the model starts from the Tucker-compatible linear response and then refines the shared response coefficients.

TABLE VIII

SCENE-LEVEL CONSISTENCY ON CAVE FULL RECONSTRUCTION UNDER CLEAN AND NOISY OBSERVATIONS. W/L COUNTS SCENES WHERE NTD-PL IMPROVES OR WORSENS MATCHED-RANK TUCKER; GAINS ARE ABSOLUTE METRIC REDUCTIONS AGAINST THE CLEAN TENSOR.

SNR	RMSE gain			SAM gain		
	W/L	Median	Wilcoxon p	W/L	Median	Wilcoxon p
Clean	15/0	0.0011	3.05×10^{-5}	12/3	0.3399	0.008
40 dB	15/0	0.0011	3.05×10^{-5}	12/3	0.3415	0.008
30 dB	15/0	0.0011	3.05×10^{-5}	12/3	0.3416	0.008
20 dB	15/0	0.0011	3.05×10^{-5}	12/3	0.3598	0.006
10 dB	15/0	0.0012	3.05×10^{-5}	14/1	0.9096	7.63×10^{-4}

TABLE IX

SCENE-LEVEL CONSISTENCY ON CLEAN CAVE RANDOM-MISSING COMPLETION. W/L COUNTS SCENES WHERE NTD-PL IMPROVES OR WORSENS MATCHED-RANK TUCKER AFTER AVERAGING MASK SEEDS; GAINS ARE ABSOLUTE METRIC REDUCTIONS ON HELD-OUT ENTRIES.

Missing rate	RMSE gain			SAM gain		
	W/L	Median	Wilcoxon p	W/L	Median	Wilcoxon p
0.1	12/3	0.0014	0.003	13/2	1.7338	8.54×10^{-4}
0.3	12/3	0.0014	0.003	11/4	0.6953	0.041
0.5	12/3	0.0014	0.003	10/5	0.4144	0.151
0.7	12/3	0.0014	0.002	11/4	0.3021	0.151

check, Gaussian noise is added to the full observation used for fitting and metrics are computed against the clean cube. Random-missing completion samples clean observation masks at the requested missing rates and evaluates RMSE and SAM only on the held-out entries. Tucker and NTD-PL use identical Tucker ranks; the default NTD-PL instance adds only the polynomial response coefficients and the associated training schedule. Table VIII reports scene-level consistency for full reconstruction at the main matched rank, and Table IX reports the corresponding paired consistency summary for clean random-missing completion.

The noisy-observation robustness evaluation uses the full-

reconstruction protocol on all 15 CAVE scenes resized to 128×128 spatial resolution with rank $(12, 12, 6)$. Gaussian noise is added after global max normalization at SNR levels 40, 30, 20, and 10 dB; the clean case uses the same protocol without added noise. The clean evaluation tensor is shared by Tucker and NTD-PL for each scene. Both methods use a shorter iteration budget than the main CAVE experiments, making this a controlled noise-robustness check complementary to the full-resolution clean reconstruction and completion benchmarks.

External hyperspectral experiments use Jasper Ridge, Samson, and Urban cubes after the same global normalization. Table XI reports the shapes and Tucker ranks used in the external reconstruction check, and Table XII reports the corresponding full-reconstruction results. The same fixed-rank NTD-PL model improves RMSE by 3.6–7.9% and also improves SAM on all three cubes. The NMSE changes are positive as well, with gains of 0.32–0.72 dB. These results serve as a supplementary cross-dataset check beyond the CAVE acquisition setting. Additional low-diagnostic checks verify the expected operating regime: when the response diagnostic is near zero, the learned response is not expected to provide a reliable improvement.

The higher-rank Tucker comparison in Table I searches over higher Tucker ranks and reports the first Tucker model whose RMSE matches the lower-rank NTD-PL reference. Its purpose is to quantify the parameter increase required by increasing ordinary Tucker rank when the residual can be explained by the learned response.

D. Response Diagnostic Details

The diagnostic D_{link} is computed from a fixed Tucker fit. First, a matched-rank Tucker model is trained under the same mask and normalization as the main experiment. Second, the Tucker reconstruction is held fixed and the response coefficients are refit by the ridge update in (11). The diagnostic




Fig. 8. CAVE full reconstruction at matched Tucker rank. NTD-PL improves both RMSE and SAM, and the gap decreases as Tucker rank increases.

TABLE X

CAVE NOISY-OBSERVATION RECONSTRUCTION. MODELS ARE FITTED TO CLEAN OR NOISY FULL OBSERVATIONS, AND METRICS ARE COMPUTED AGAINST THE CLEAN TENSOR; GAINS ARE RELATIVE ERROR REDUCTIONS OVER MATCHED-RANK TUCKER.

SNR	Clean RMSE↓			Clean SAM↓		
	Tucker	NTD-PL	Rel. gain	Tucker	NTD-PL	Rel. gain
Clean	0.0320 ± 0.0274	0.0305 ± 0.0266	4.86%	14.71 ± 7.57	13.87 ± 6.73	5.74%
40 dB	0.0320 ± 0.0274	0.0305 ± 0.0266	4.86%	14.72 ± 7.58	13.87 ± 6.73	5.77%
30 dB	0.0320 ± 0.0274	0.0305 ± 0.0266	4.86%	14.71 ± 7.55	13.87 ± 6.73	5.73%
20 dB	0.0321 ± 0.0274	0.0305 ± 0.0266	4.87%	14.95 ± 7.72	13.99 ± 6.80	6.41%
10 dB	0.0326 ± 0.0273	0.0310 ± 0.0265	4.88%	16.01 ± 7.93	14.63 ± 6.86	8.58%

TABLE XI

EXTERNAL HYPERSPECTRAL RECONSTRUCTION PROTOCOL. TUCKER AND NTD-PL USE THE SAME TUCKER RANK ON EACH CUBE.

Dataset	Shape	Normalization	Backbone rank
Jasper Ridge	100 × 100 × 198	global max	(11, 11, 43)
Samson	95 × 95 × 156	global max	(10, 10, 35)
Urban	307 × 307 × 162	global max	(43, 43, 47)

TABLE XII

EXTERNAL HYPERSPECTRAL RECONSTRUCTION. NTD-PL USES THE SAME TUCKER RANK AS THE TUCKER BASELINE AND ADDS ONLY THE POLYNOMIAL RESPONSE COEFFICIENTS.

Dataset	RMSE↓			SAM↓		
	Tucker	NTD-PL	Rel. gain	Tucker	NTD-PL	Rel. gain
Jasper Ridge	0.0462	0.0430	6.9%	9.14	8.16	10.7%
Samson	0.0263	0.0242	7.9%	5.70	5.65	1.0%
Urban	0.0444	0.0428	3.6%	8.23	7.94	3.5%

is the fractional residual energy removed by this one response update, as defined in (17). This procedure is cheaper than full NTD-PL training because it does not update the Tucker core and factors through the nonlinear response.

The CAVE diagnostic table in the main text groups scenes into terciles by D_{link} and compares the eventual held-out RMSE gain of full NTD-PL. We use the same fixed-Tucker computation for external checks and low-diagnostic regimes where the learned response is not expected to help.

E. Cross-Domain Diagnostic Data

The cross-domain diagnostic in Fig. 5 treats each tensor unit as an independent recovery problem. Natural image units include CBSD68 and CIFAR-10 images [28], [29]; object-view units include COIL-100 and smallNORB tensors [30], [31]. Each unit is normalized independently, fitted by matched-rank

Tucker and NTD-PL, and summarized by response diagnostic yield versus the eventual RMSE gain. These experiments are diagnostic tensor-recovery checks. Their purpose is to test whether this response-based residual pattern appears beyond hyperspectral cubes.

CBSD68 units are resized natural images with rank (32, 32, 3). CIFAR-10 uses 1000 class-balanced test images of shape $32 \times 32 \times 3$ with rank (20, 20, 3). COIL-100 uses one object at a time as a view tensor with rank (8, 12, 12, 2). smallNORB uses one object instance at a time as a stereo view tensor with rank (12, 24, 24, 2). KTH-Action [32] and UCF101 [33] provide the action-video tensor units. All units are normalized independently, and diagnostic labels are assigned from D_{link} , not from the eventual NTD-PL gain.

F. Compute Resources

The reported runs used a CPU Python implementation based on NumPy, SciPy, and TensorLy on a workstation with a 13th Gen Intel Core i5-13400F CPU, 32 GB system RAM, Windows 11, and Python 3.13. An NVIDIA RTX 5060 GPU with 8 GB memory was present, but the reported scripts do not use GPU kernels. Sweep scripts cap BLAS/OpenMP threads per worker and parallelize independent jobs, using up to 12 workers for the main low-rank sweeps and 6 workers for the diagnostic batches.

From the recorded runtime fields, controlled synthetic residual jobs were sub-second to about 1 s per configuration. Full-resolution CAVE reconstruction took about 10–17 s per scene for matched-rank Tucker and 172–217 s per scene for NTD-PL in the main low-rank reconstruction sweep. CAVE random completion runs took about 0.9–18.6 s per scene/mask/seed, while structured completion runs were about 1.6–5.0 s. The CAVE response diagnostic itself took about 0.12–0.59 s per fitted refresh. Non-HSI diagnostic batches with 6 CPU workers took about 77 s for CBSD68/CIFAR-10/COIL-100/smallNORB and about 5–6 minutes for each action-video

TABLE XIII
DATASET-LEVEL STATISTICS BEHIND THE CROSS-DOMAIN DIAGNOSTIC SCATTER IN FIG. 5. GAIN COLUMNS REPORT MEAN NTD-PL RMSE IMPROVEMENT OVER INDEPENDENTLY FITTED MATCHED-RANK TUCKER WITHIN D_{link} TERCILES.

Domain	Dataset	Units	Med. D_{link}	ρ_s	Low	Mid	High
Natural images	CBSD68	68	0.091	0.85	4.17%	8.40%	14.60%
Natural images	CIFAR-10	1000	0.008	0.93	-0.01%	2.59%	7.85%
Object-view	COIL-100	100	0.208	0.93	2.77%	5.46%	9.14%
Object-view	smallNORB	25	0.264	0.95	6.92%	9.91%	14.16%
Action video	KTH-Action	240	0.107	0.88	2.74%	4.27%	10.87%
Action video	UCF101	240	0.065	0.89	1.44%	2.97%	5.89%

batch. Peak memory was not instrumented per run; empirically the largest reported CPU runs completed within the 32 GB RAM budget.

G. Dataset Asset and Usage Information

All datasets used in the experiments are public research benchmarks. They are used only for reconstruction, completion, or diagnostic evaluation; processed summaries and scripts are redistributed rather than raw dataset files.

ACKNOWLEDGMENT

The authors would like to thank ...

REFERENCES

- [1] T. G. Kolda and B. W. Bader, "Tensor decompositions and applications," *SIAM Review*, vol. 51, no. 3, pp. 455–500, 2009.
- [2] N. D. Sidiropoulos, L. D. Lathauwer, X. Fu, K. Huang, E. E. Papalexakis, and C. Faloutsos, "Tensor decomposition for signal processing and machine learning," *IEEE Transactions on Signal Processing*, vol. 65, no. 13, pp. 3551–3582, 2017.
- [3] J. Liu, P. Musialski, P. Wonka, and J. Ye, "Tensor completion for estimating missing values in visual data," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 35, no. 1, pp. 208–220, 2013.
- [4] A. Liu and A. Moitra, "Tensor completion made practical," in *Advances in Neural Information Processing Systems*, vol. 33, 2020, pp. 18905–18916.
- [5] P. E. Debevec and J. Malik, "Recovering high dynamic range radiance maps from photographs," in *Proceedings of the 24th Annual Conference on Computer Graphics and Interactive Techniques*, ser. SIGGRAPH '97, 1997, pp. 369–378.
- [6] M. D. Grossberg and S. K. Nayar, "Modeling the space of camera response functions," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 26, no. 10, pp. 1272–1282, 2004.
- [7] J. M. Bioucas-Dias, A. Plaza, N. Dobigeon, M. Parente, Q. Du, P. Gader, and J. Chanussot, "Hyperspectral unmixing overview: Geometrical, statistical, and sparse regression-based approaches," *IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing*, vol. 5, no. 2, pp. 354–379, 2012.
- [8] Y. Altmann, N. Dobigeon, and J.-Y. Tourneret, "Supervised nonlinear spectral unmixing using a postnonlinear mixing model for hyperspectral imagery," *IEEE Transactions on Image Processing*, vol. 21, no. 6, pp. 3017–3025, 2012.
- [9] R. Heylen, M. Parente, and P. Gader, "A review of nonlinear hyperspectral unmixing methods," *IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing*, vol. 7, no. 6, pp. 1844–1868, 2014.
- [10] M. Udell, C. Horn, R. Zadeh, and S. Boyd, "Generalized low rank models," *Foundations and Trends in Machine Learning*, vol. 9, no. 1, pp. 1–118, 2016.
- [11] R. Ganti, L. Balzano, and R. Willett, "Matrix completion under monotonic single index models," in *Advances in Neural Information Processing Systems*, vol. 28, 2015.
- [12] Y. Luo, X. Zhao, Z. Li, M. K. Ng, and D. Meng, "Low-rank tensor function representation for multi-dimensional data recovery," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 46, no. 5, pp. 3351–3369, 2024.
- [13] V. Saragadam, R. Balestrieri, A. Veeraraghavan, and R. G. Baraniuk, "DeepTensor: Low-rank tensor decomposition with deep network priors," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 46, no. 12, pp. 10337–10348, 2024.
- [14] D. Ahn, U. S. Saini, E. E. Papalexakis, and A. Payani, "Neural additive tensor decomposition for sparse tensors," in *Proceedings of the 33rd ACM International Conference on Information and Knowledge Management*, 2024, pp. 14–23.
- [15] A. Cichocki, D. P. Mandic, A.-H. Phan, C. F. Caiafa, G. Zhou, Q. Zhao, and L. D. Lathauwer, "Tensor decompositions for signal processing applications: From two-way to multiway component analysis," *IEEE Signal Processing Magazine*, vol. 32, no. 2, pp. 145–163, 2015.
- [16] L. D. Lathauwer, B. D. Moor, and J. Vandewalle, "A multilinear singular value decomposition," *SIAM Journal on Matrix Analysis and Applications*, vol. 21, no. 4, pp. 1253–1278, 2000.
- [17] I. V. Oseledets, "Tensor-train decomposition," *SIAM Journal on Scientific Computing*, vol. 33, no. 5, pp. 2295–2317, 2011.
- [18] Q. Zhao, G. Zhou, S. Xie, L. Zhang, and A. Cichocki, "Tensor ring decomposition," *arXiv preprint arXiv:1606.05535*, 2016.
- [19] L. Yang, J. Fang, H. Li, and B. Zeng, "An iterative reweighted method for tucker decomposition of incomplete tensors," *IEEE Transactions on Signal Processing*, vol. 64, no. 18, pp. 4817–4829, 2016.
- [20] D. Hong, T. G. Kolda, and J. A. Dueresch, "Generalized canonical polyadic tensor decomposition," *SIAM Review*, vol. 62, no. 1, pp. 133–163, 2020.
- [21] A. Schein, M. Zhou, D. Blei, and H. Wallach, "Bayesian poisson tucker decomposition for learning the structure of international relations," in *Proceedings of the International Conference on Machine Learning*, 2016, pp. 2810–2819.
- [22] Z. Xu and F. Yan, "Infinite tucker decomposition: Nonparametric bayesian models for multiway data analysis," *arXiv preprint arXiv:1108.6296*, 2011.
- [23] B. Liu, L. He, S. Zhe, Y. Li, and Z. Xu, "DeepCP: Flexible nonlinear tensor decomposition," in *NeurIPS Workshops*, 2017.
- [24] H. Attouch, J. Bolte, and B. F. Svaiter, "Convergence of descent methods for semi-algebraic and tame problems: Proximal algorithms, forward-backward splitting, and regularized gauss-seidel methods," *Mathematical Programming*, vol. 137, no. 1–2, pp. 91–129, 2013.
- [25] J. Bolte, S. Sabach, and M. Teboulle, "Proximal alternating linearized minimization for nonconvex and nonsmooth problems," *Mathematical Programming*, vol. 146, no. 1–2, pp. 459–494, 2014.
- [26] Columbia CAVE Laboratory, "Multispectral image database," <https://cave.cs.columbia.edu/repository/Multispectral>, n.d., accessed: 2026-04-13.
- [27] F. Yasuma, T. Mitsunaga, D. Iso, and S. K. Nayar, "Generalized assorted pixel camera: Postcapture control of resolution, dynamic range, and spectrum," *IEEE Transactions on Image Processing*, vol. 19, no. 9, pp. 2241–2253, 2010.
- [28] D. Martin, C. Fowlkes, D. Tal, and J. Malik, "A database of human segmented natural images and its application to evaluating segmentation algorithms and measuring ecological statistics," in *Proceedings of the IEEE International Conference on Computer Vision*, 2001, pp. 416–423.
- [29] A. Krizhevsky, "Learning multiple layers of features from tiny images," University of Toronto, Tech. Rep., 2009. [Online]. Available: <https://www.cs.toronto.edu/~kriz/cifar.html>
- [30] S. A. Nene, S. K. Nayar, and H. Murase, "Columbia Object Image Library (COIL-100)," Columbia University, Tech. Rep. CUCS-006-96, 1996. [Online]. Available: <https://cave.cs.columbia.edu/repository/COIL-100>
- [31] Y. LeCun, F. J. Huang, and L. Bottou, "Learning methods for generic object recognition with invariance to pose and lighting," in *Proceedings*

TABLE XIV
DATASET SOURCES AND USAGE NOTES. “NO” IN THE REDISTRIBUTION COLUMN MEANS THAT RAW DATA FILES ARE NOT INCLUDED.

Dataset	Source / citation	License or usage note	Redistribution / use
CAVE Multispectral Image Database	Columbia CAVE Laboratory; [26], [27]	Public download and citation information are provided by the official page; used only for research evaluation.	No; reconstruction, completion, diagnostics.
CBSD68 / Berkeley natural images	Berkeley Segmentation Dataset; [28]	Official usage statement permits research and non-commercial use and requests citation.	No; non-HSI diagnostic evaluation.
CIFAR-10	University of Toronto CIFAR page; [29]	Public research benchmark download and citation information are provided by the official page.	No; non-HSI diagnostic evaluation.
COIL-100	Columbia Object Image Library; [30]	Public download and citation information are provided by the official page; used only for research evaluation.	No; non-HSI diagnostic evaluation.
smallNORB	NYU/Yann LeCun dataset page; [31]	Official page provides the database for research purposes and requests citation.	No; non-HSI diagnostic evaluation.
KTH-Action	KTH action-recognition dataset page; [32]	Official page states that the database is publicly available for non-commercial use and requests citation.	No; non-HSI diagnostic evaluation.
UCF101	UCF CRCV dataset page; [33]	Public benchmark download and citation information are provided by the official page; used only for research evaluation.	No; non-HSI diagnostic evaluation.

of the *IEEE Computer Society Conference on Computer Vision and Pattern Recognition*, vol. 2, 2004, pp. 97–104.

- [32] C. Schuldt, I. Laptev, and B. Caputo, “Recognizing human actions: A local SVM approach,” in *Proceedings of the 17th International Conference on Pattern Recognition*, vol. 3, 2004, pp. 32–36. [Online]. Available: <https://www.csc.kth.se/cvap/actions/>
- [33] K. Soomro, A. R. Zamir, and M. Shah, “UCF101: A dataset of 101 human action classes from videos in the wild,” University of Central Florida, Tech. Rep. CRCV-TR-12-01, 2012. [Online]. Available: <https://www.crcv.ucf.edu/data/UCF101.php>